

---

University of Connecticut  
Computer Science and Engineering  
CSE 4402/5095: Network Security

# Security of the Domain Name System (DNS)

©Amir Herzberg  
Version of Fall '24

# Domain Name System (DNS) Security

---

- **Introducing DNS (or a quick recap)**
- **DNS poisoning**
  - Motivation: the hacker's swiss knife
  - Method 1 (historical): by Gratuitous `glue' RR
  - Method 2: send from corrupt NS
  - Method 3: send spoofed DNS response
  - Method 4: dangling DNS records
- **DNSSEC: Cryptographic security for DNS**
- **DNS Privacy issues and defenses**

# Domain Names and IP Addresses

---

- IP packets contain source, dest IP addresses
  - 32 bits, e.g. 128.33.44.223
- Routers use IP Addresses
  - To deliver packets to their destinations
- Users use Domain Names, e.g. `www.foo.edu`
- Domain Names are hierarchical, and:
  - Meaningful: `*.edu`: university, `www.*`: web server
  - Easier to manage, remember and use
  - The Domain Name System (DNS) maps domain names to IP addresses

# Benefits of using domain names, not IPs

---

- Usability: easier to remember, meaningful
- Abstraction
  - Fixed domain-name mapped to sometimes changing IP
  - IP changes for mobility, optimization (close to customer), business (CDN/cloud, change provider)
  - Multiple IPs for the same domain: redundancy, IPv4 & IPv6, localization (for efficiency)
  - Organization: hierarchy, wildcards
- Map domain names to different **record types**:
  - IPv4 and IPv6 addresses
  - And many more types of **resource record (RR)** types

# Some Domain Name System Record Types

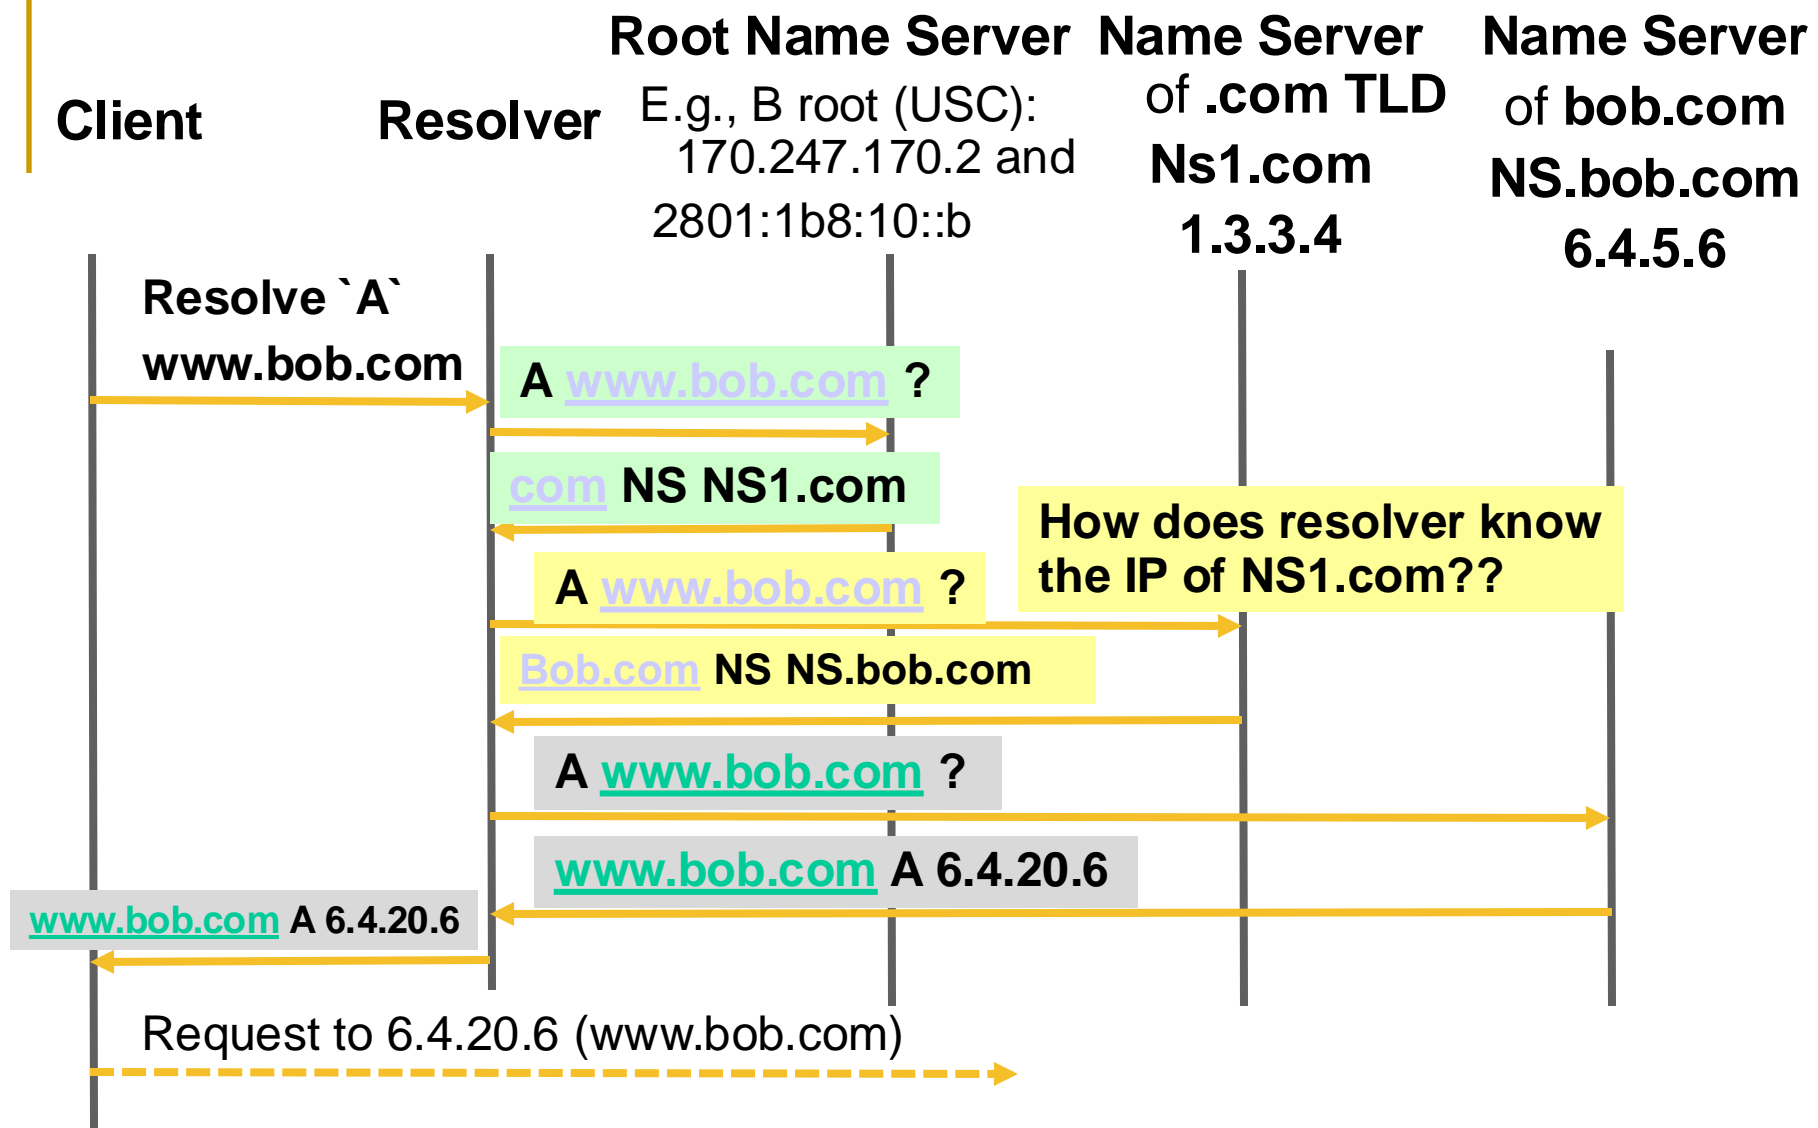
- A: IPv4 address, AAAA: IPv6 **IP of domain**
- 'Alias' name: CNAME **Name of email (SMTP) server of domain**
- Mail server: MX **Name of name server of domain**
- Name server: NS **Servers authorized to send email for domain**
- Server for service: SRV **Indication if IP is spammer**
- Policies and PKs: CAA, SPF, DKIM, ... **CAs used by BoA**
- Blacklists
- Any text: TXT
- Reverse-DNS ('pointer') records, map an IP to a domain: PTR **Domain name of owner of IP**

# DNS Entities

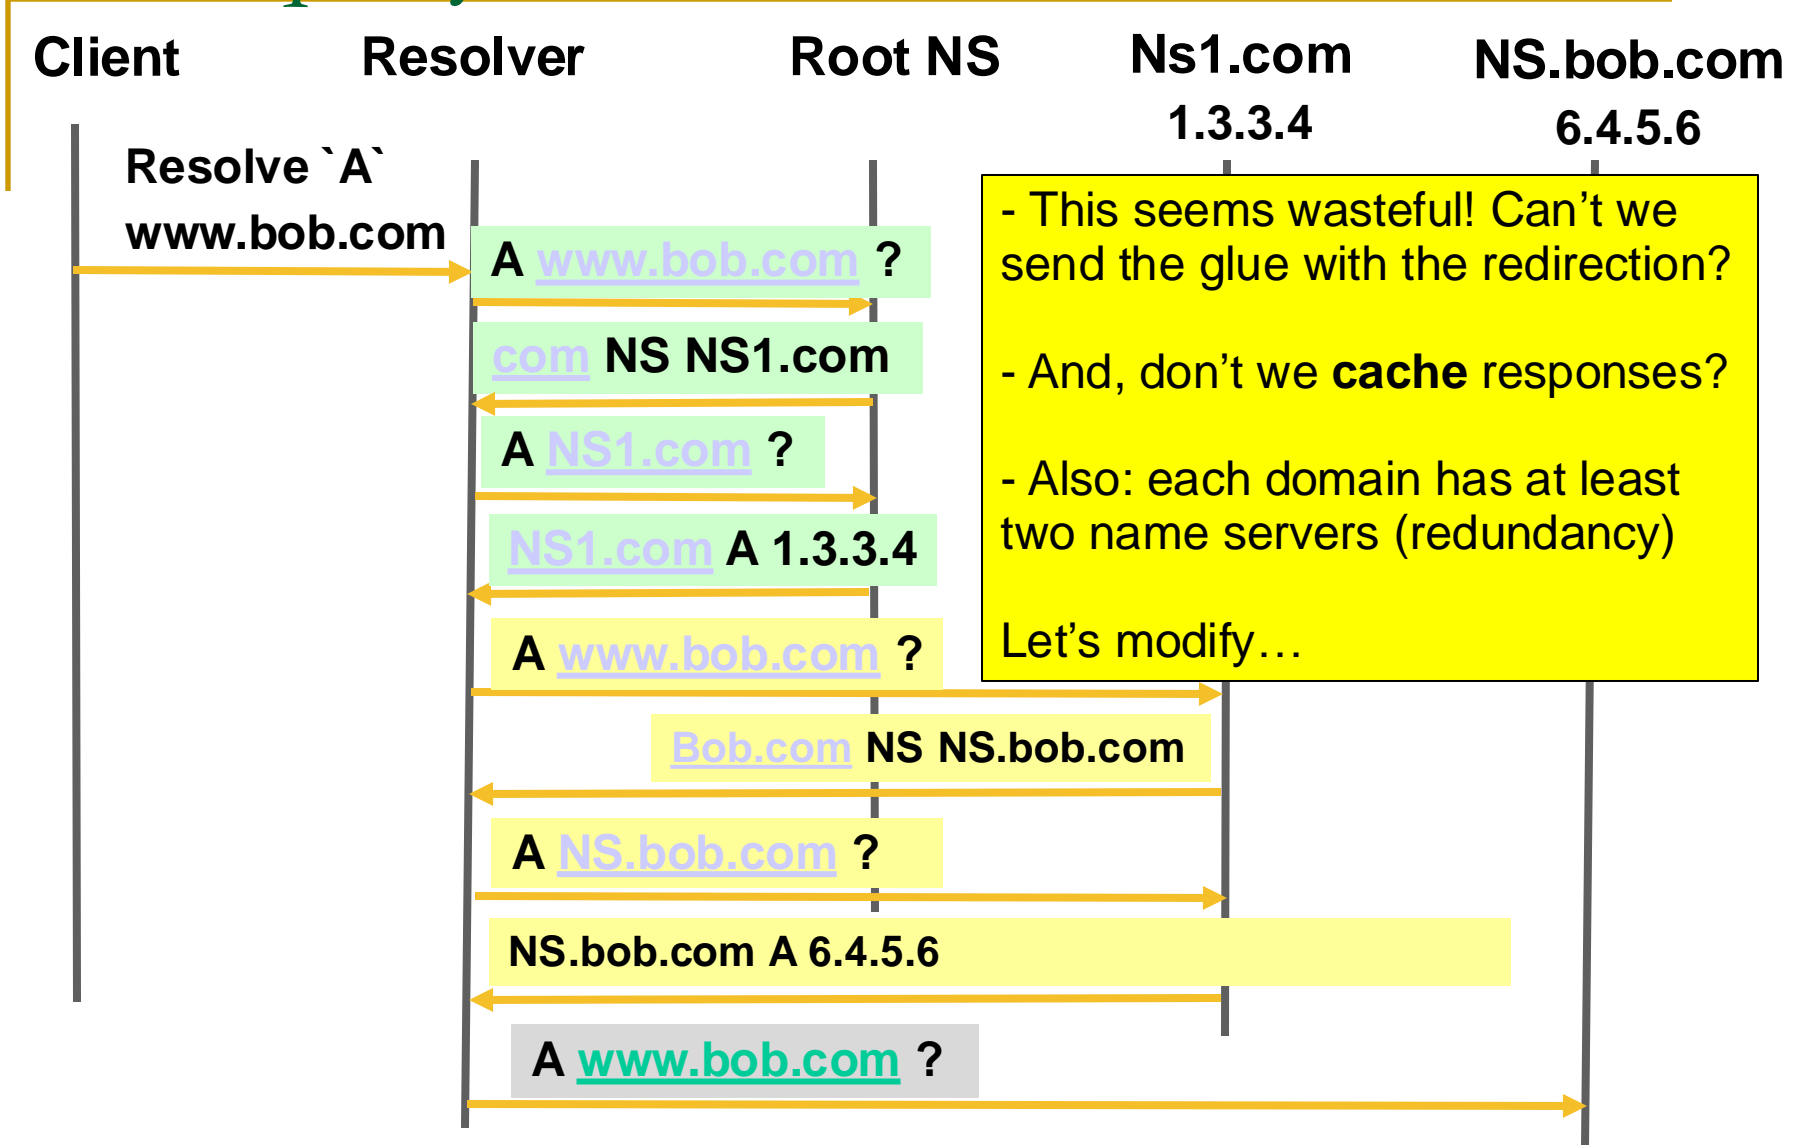
---

- DNS Client (aka DNS stub resolver)
  - Provides DNS resource records (RRs) to applications
  - If requested RR (type, domain) not in cache, asks resolver
  - Contact resolver by IP address (set manually or by DHCP)
- Resolver (aka recursive resolver, DNS cache, or local NS)
  - Returns resource records (RRs) requested by clients
  - Requests RRs from Name Servers, if not in resolver's cache
    - Selects order of name servers to try based on past performance: failures and delays
- Name servers (aka authoritative NS)
  - Provides the RRs of one or more domains
  - Hierarchical, e.g., NS of .com identifies NS of bob.com

# DNS Resolution Process (simplified)



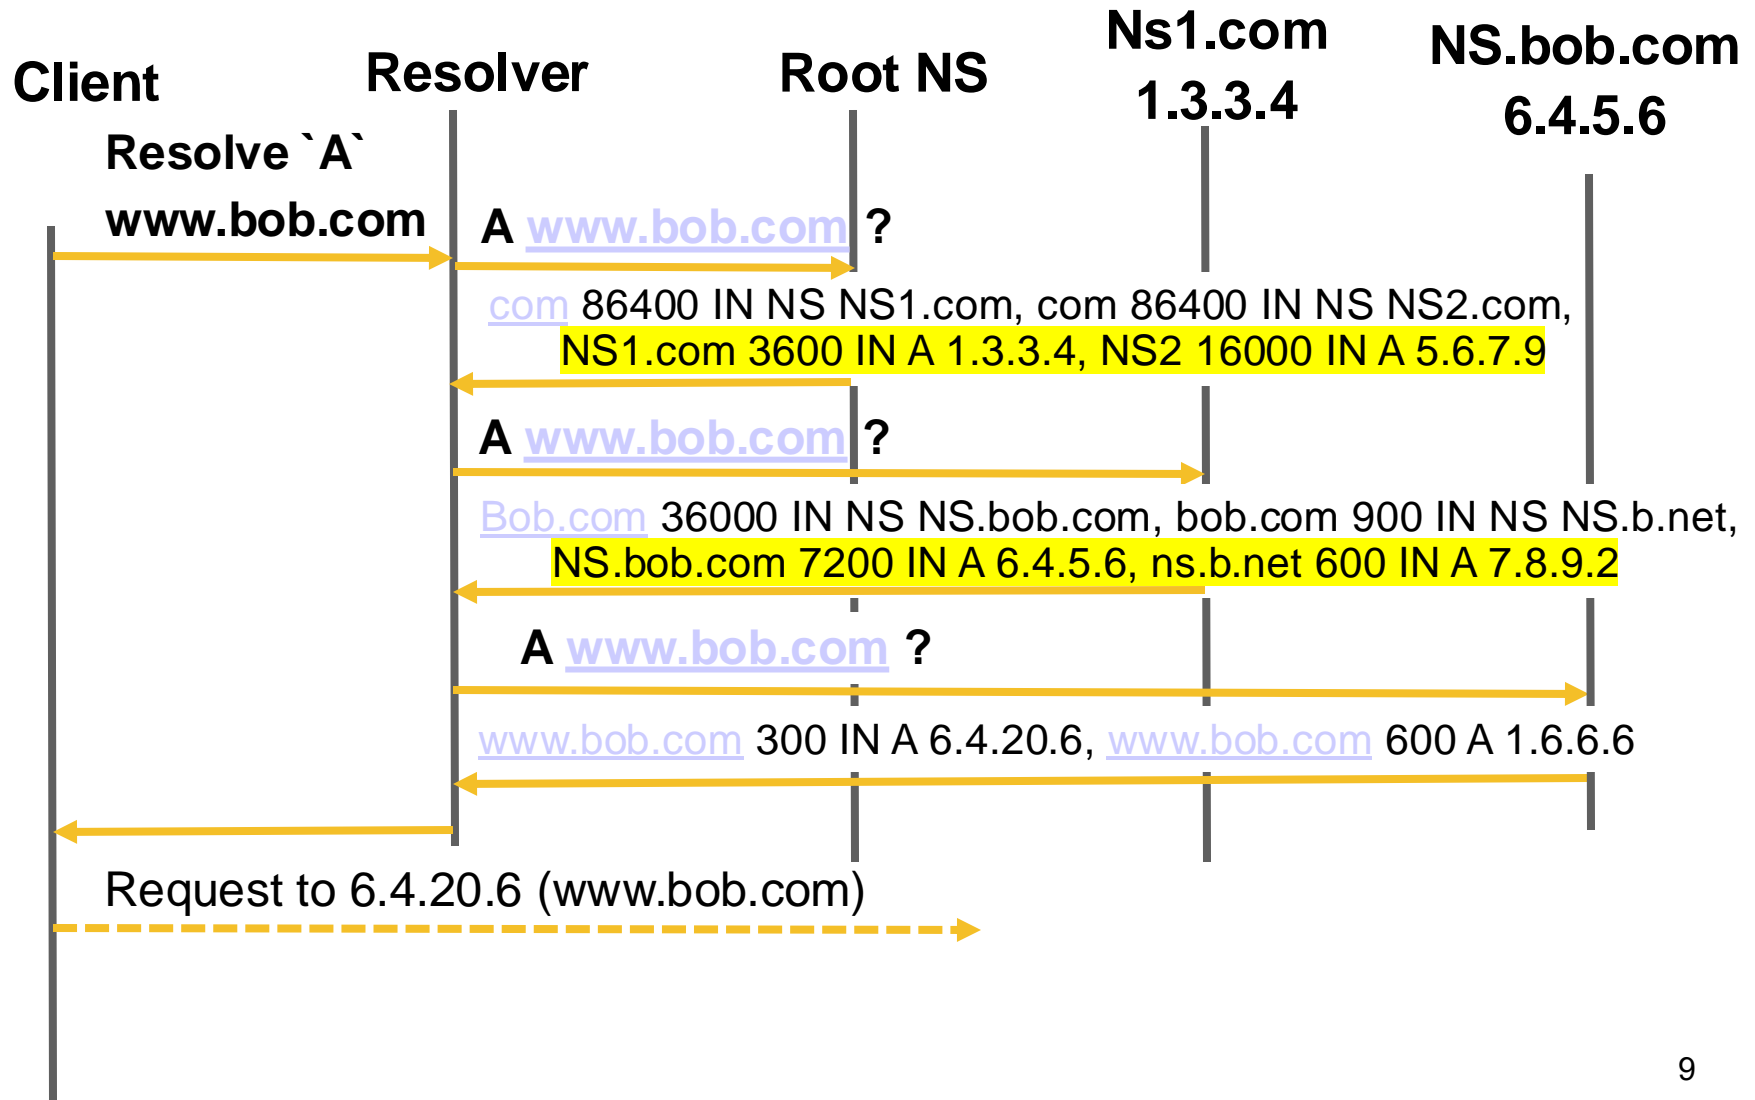
# Glue query to find IPs of name servers





# DNS Resolution Process: with Glue,

multiple name servers and TTL (caching)



# DNS Caching

---

- Caching of RRs is critical for performance
  - An RR is cached for TTL value (specified in RR)
  - Non-existing domain (Nxdomain) responses cached for period set by resolver
- Each DNS client maintains a cache
  - And sends requests to resolver (which keeps a cache too)
- Resolver cache
  - Sends query to NS of most specific domain in cache
    - E.g., send query to NS of Bob.com if known; if not, to NS.com (or root)
  - Often nested (univ. resolver → ISP's resolver)
  - 'Open resolver': provides service to any client

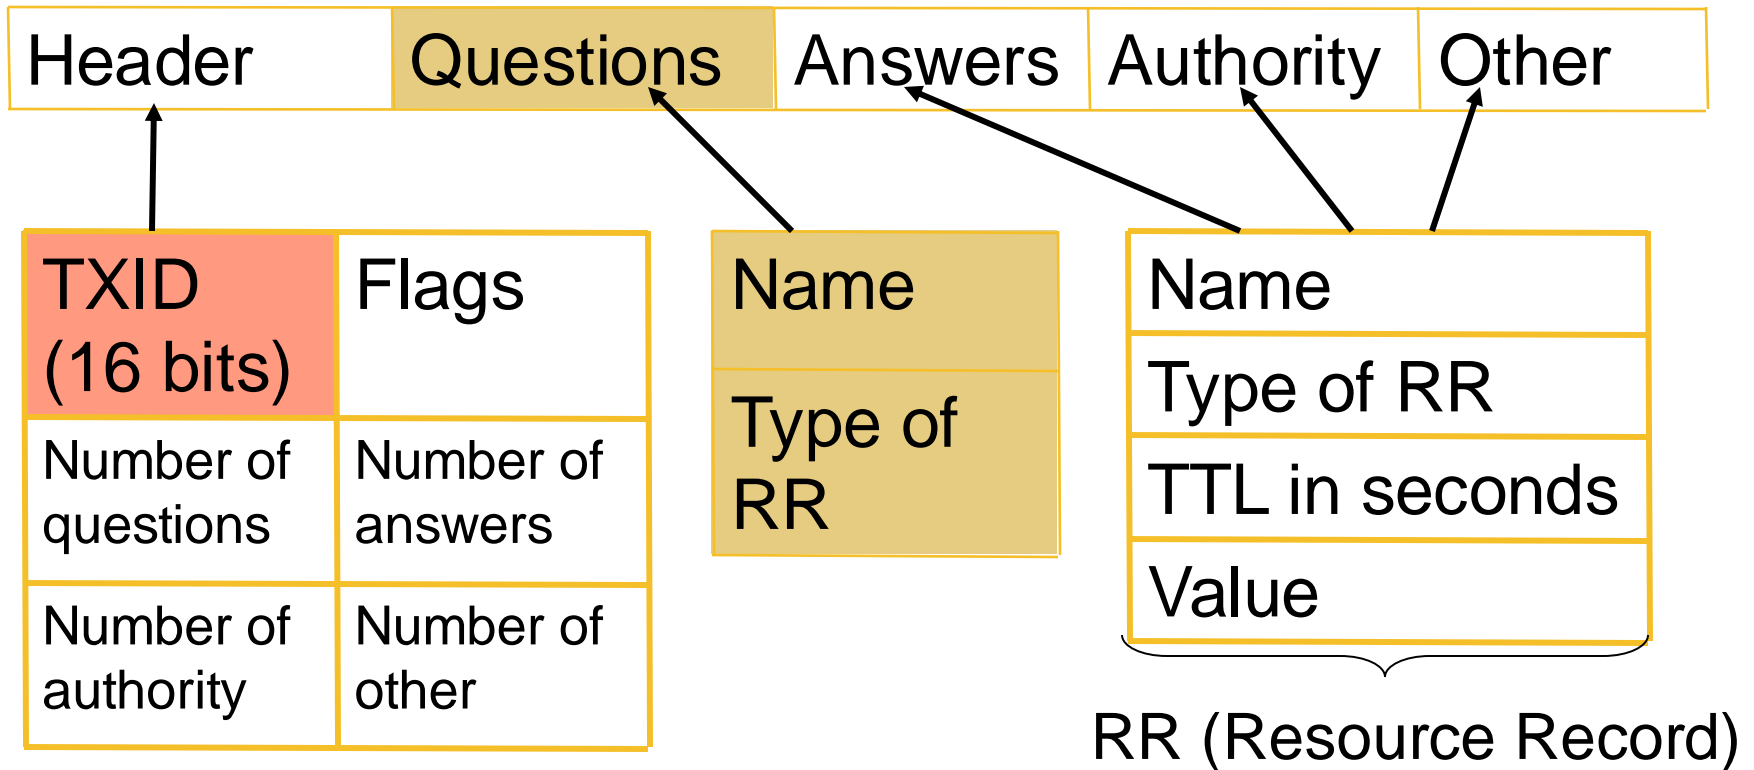
# Examples of Domain Name System Records

---

- A: IPv4 address, AAAA: IPv6 address
  - Bob.com A 1.2.3.4, bob.com AAAA 2024:1111:2222:333:4444:FFFF
- 'Alias' name: CNAME
  - Bob.com CNAME bob.com.cdn.net
- Mail server: MX
  - Bob.com MX 1 smtp.bob.com, bob.com MX 5 backup.mail.bob.com
- Name server: NS
  - Bob.bom NS ns1.bob.com
- Server for service (IMAP, SIP, XMPP, LDAP,...): SRV
  - Service, protocol, priority, weight, port, server-domain
  - \_imap.\_tcp.bob.com. 86400 IN SRV 10 4 143 imap.bob.com
- Policies and PKs: CAA, SPF, DKIM,...
  - CAA (CA authorization) defines allowed issuing CA, policy, etc.
- Reverse-DNS ('pointer', rDNS): PTR
  - 6.6.6.6.in-addr.arpa PTR 666.org, 2024:::666.ip6.arpa PTR 666.org

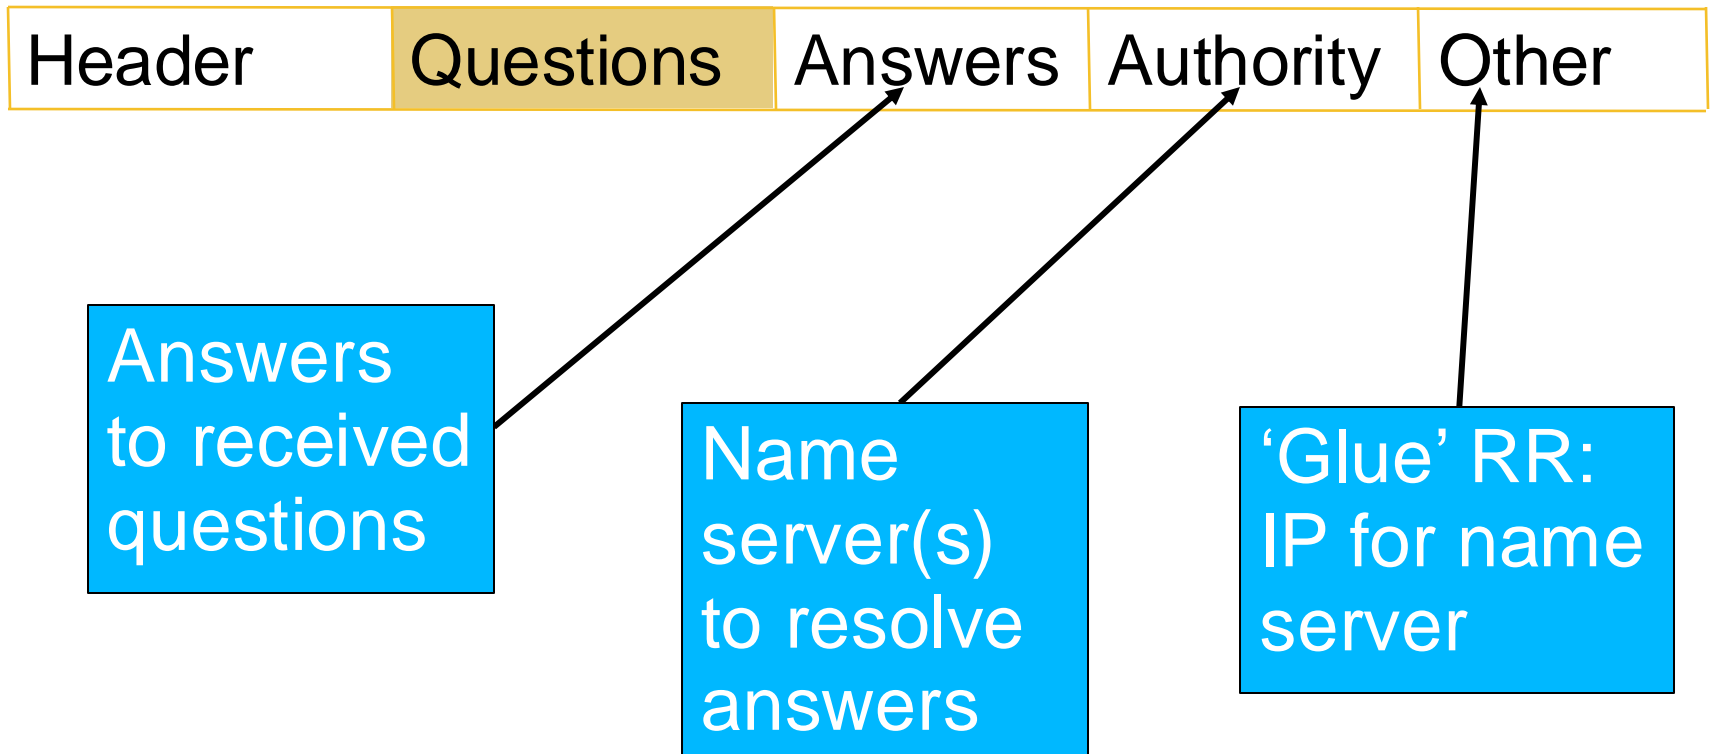
# DNS Messages

- DNS protocol: send request, receive reply
- Single format for requests & replies



# DNS Messages

- DNS protocol: send request, receive reply
- Single format for requests & replies
  - Typically, requests contain only header & questions



# DNS Messages

- DNS protocol: send request, receive reply
- Single format for requests & replies

Header	Questions	Answers	Authority	Other
--------	-----------	---------	-----------	-------

Answers to  
received questions:  
Bob.com A 1.2.3.4

Name server(s) to  
resolve answers:  
b.com NS NS.b.net

'Glue' RR:  
IP for name server  
ns.b.net A 7.8.9.2

# DNS Security: Goals

---

## ■ Authenticity

- ❑ Owners should control mappings (name → IP)
- ❑ Challenge-response defence against **off-path attacker**
- ❑ Or: DNSSEC to protect against **MitM attacker**
  - DNSSEC uses signatures and cryptographic hashing, see later

## ■ Availability

- ❑ Prevent Denial of Service (DoS) attacks

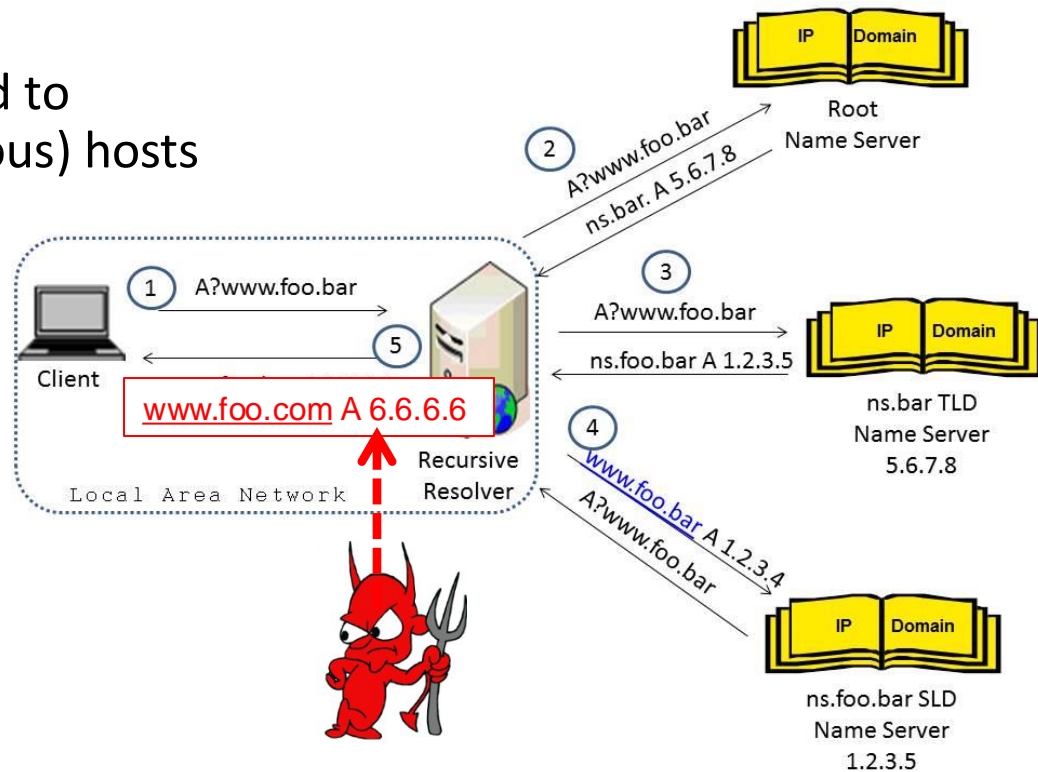
## ■ Privacy / Confidentiality ??

- Privacy for server/data from user: not a goal
  - Protocol allows any client to query any server
  - Undesirable: `what's there?` query
- Privacy for user's queries ?
  - Your resolutions tell a lot about you! Motivation to offer open resolvers ☺
  - Defense from MitM/ISP: DNS Over HTTPS / TLS (DOH, DOT)
  - Defense from resolver: NAT ?

# Threat: Cache Poisoning

Problem: no authentication of DNS responses

- Attacker can provide spoofed records
- Resolver caches
- Clients redirected to incorrect (malicious) hosts





# DNS Poisoning: the Hacker's Knife

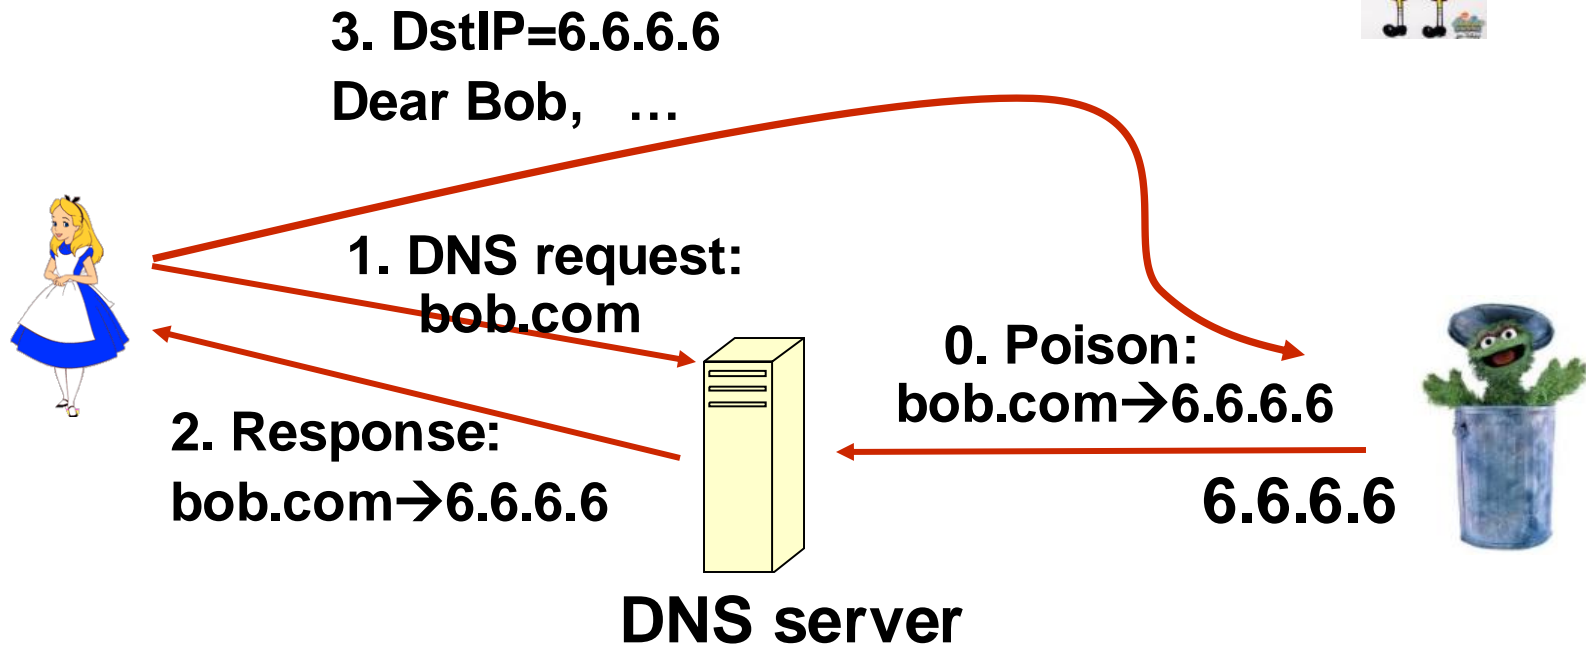


# MITM via DNS Poisoning

## ■ Allows offpath attacker to become MITM

- ❑ Web spoofing / phishing attacks
- ❑ Poison IP of NS once  
→ poison domain forever!

**Bob.com**  
**129.4.4.5**



# Steal Password by Poisoning

- Password recovery procedure exploit
  - Attacker poisons MX record used by webserver

**PayPal**

## Can't log on?

Just provide us with the following details and we'll help you access your account.

What's the problem?

☒ I've forgotten my password

Email address

☐ I've forgotten my email address

☐ I've forgotten my password and email address

✖ Enter the email address account.

**ebay**.co.uk

## Forgot your user ID?

### Here's how to fix this

Enter your email address below. Click the "Continue" button

Enter email address:

Re-enter your email address:

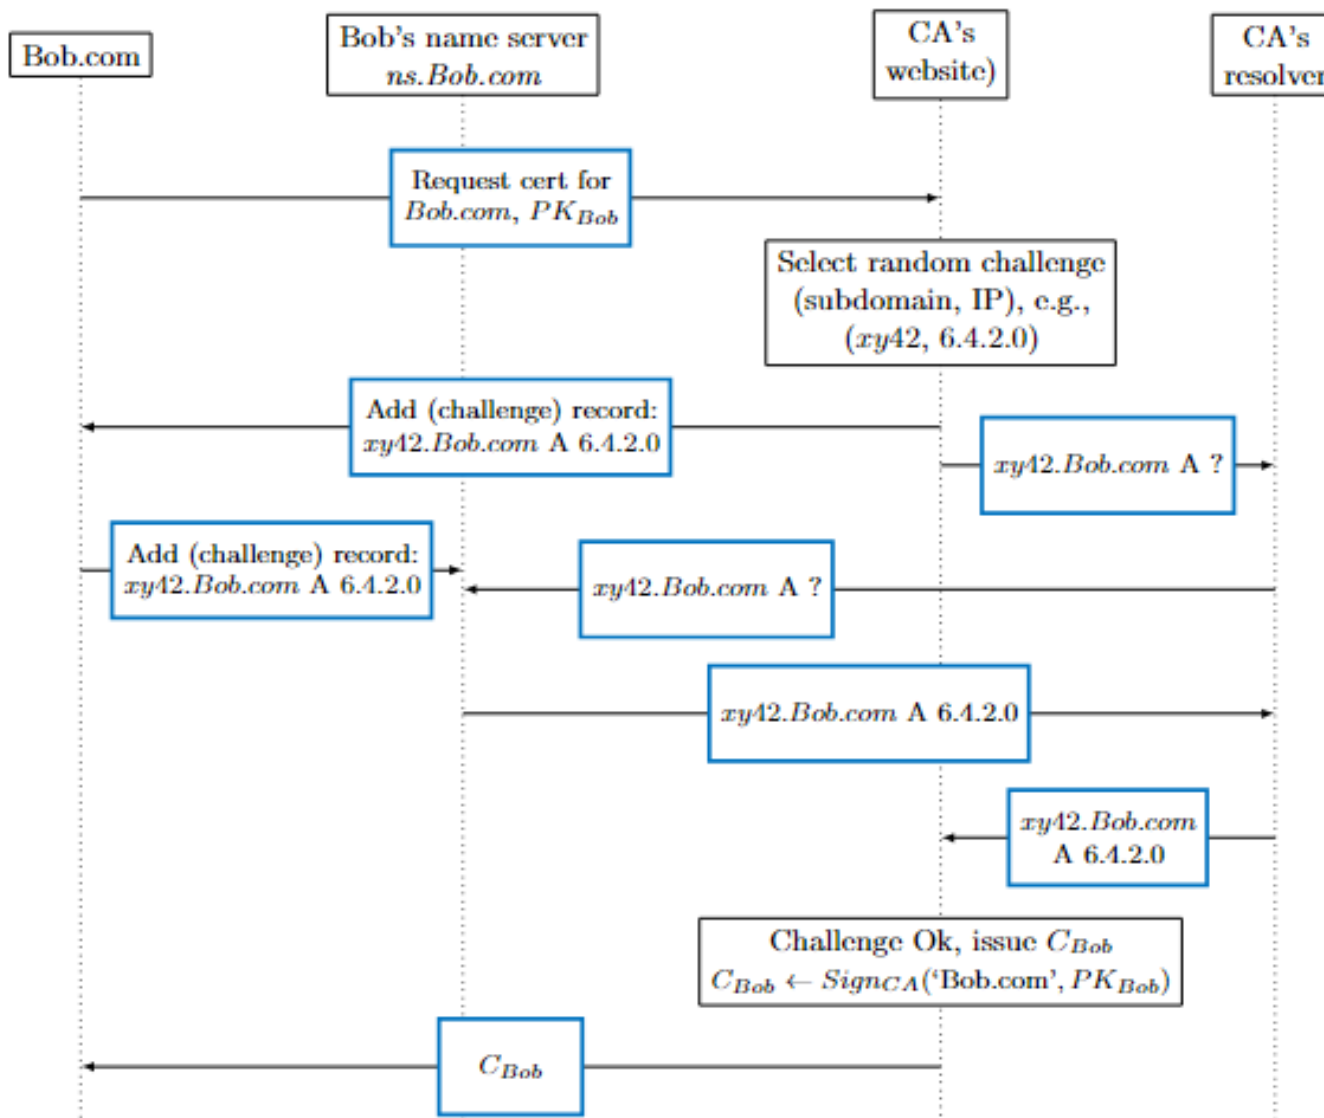
**Continue**

# Poisoning to get a rogue certificate certified

---

- A CA signs a certificate authenticating pk for domain
- The CA should verify that the request is valid
- Domain validation: validate control over the domain
- Poison the CA's resolver circumvents domain validation
- Let's first see how domain validation works...

# Domain Validation

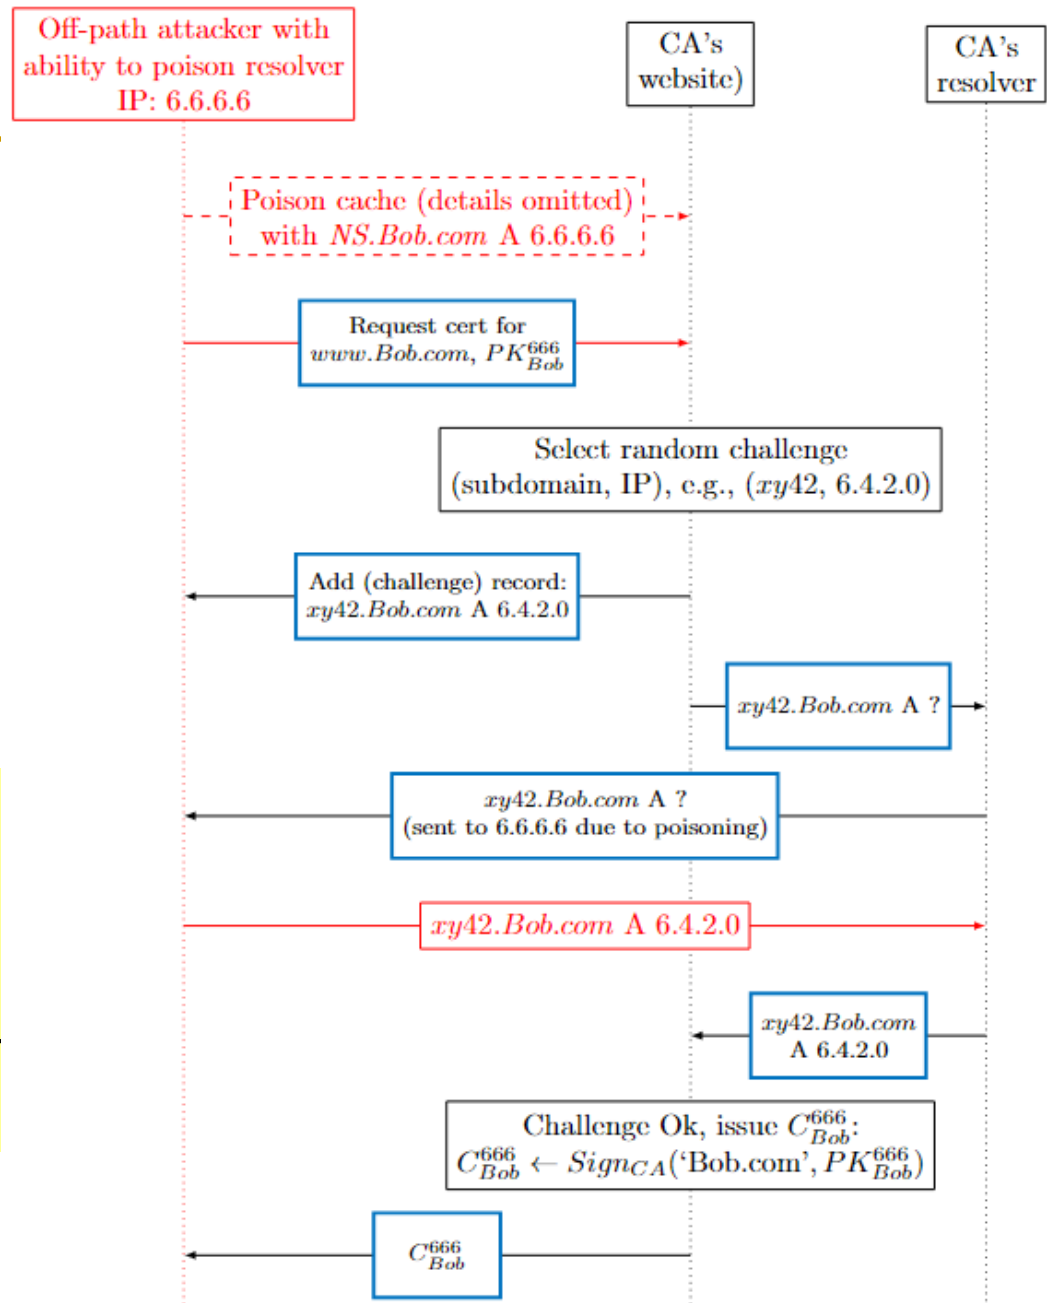


# Attack

Tricking a domain-validating CA to issue a fake certificate for [www.Bob.com](http://www.Bob.com) by DNS poisoning, against domain-validating CA.

We show later **how** the poisoning is done.

Exercise: draw a sequence diagram showing how the off-path attacker can exploit  $C_{Bob}^{666}$ . You may assume that the attacker can DNS-poison additional resolvers.



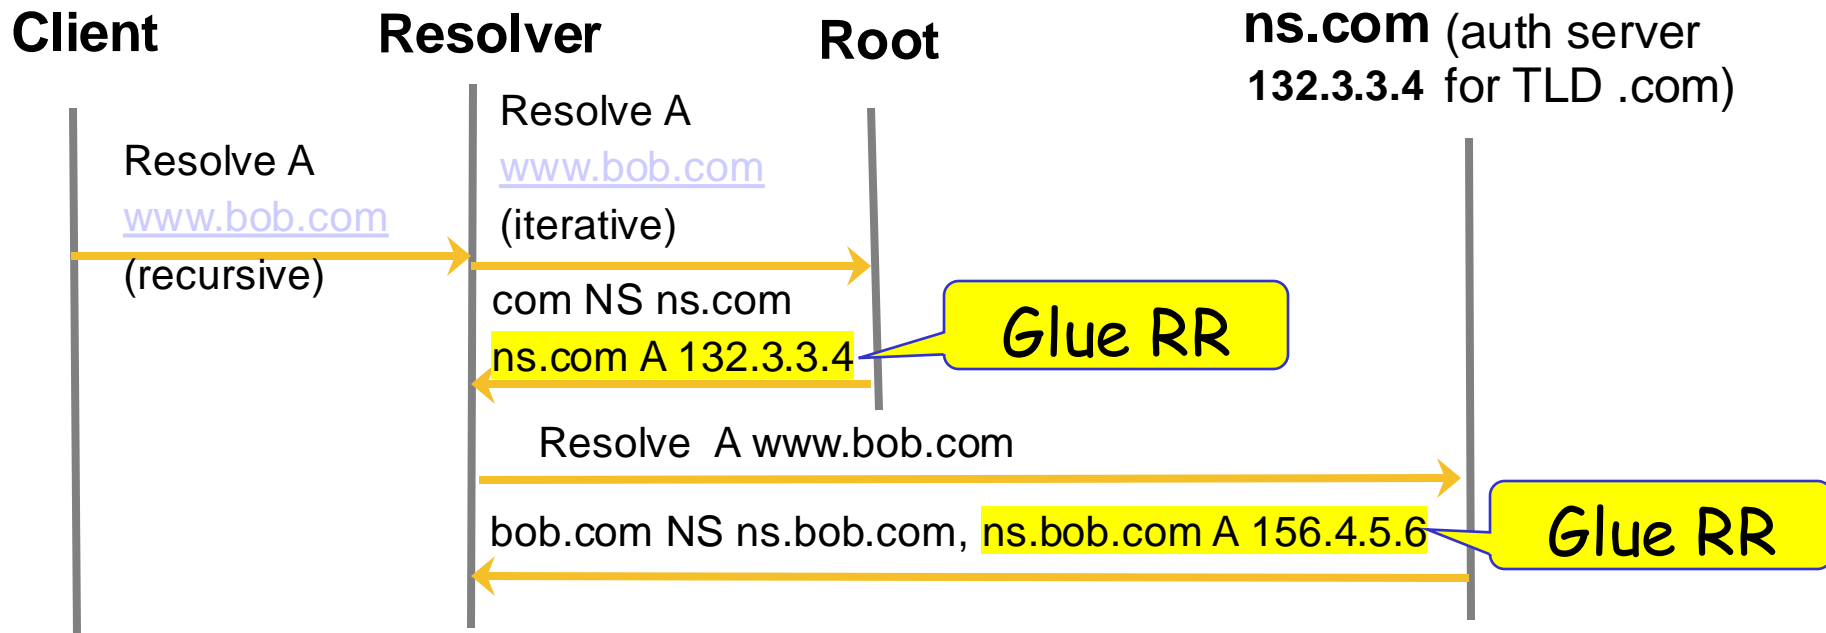
# DNS Poisoning: How?

---

- ❑ DNS Poisoning: `inject' fake DNS RR
- ❑ **Method 1 (historical):** by Gratuitous `glue' RR
  - e.g. query A [www.eve.com](http://www.eve.com), response: eve.com NS download.com and **download.com A 6.6.6.6**
  - **Bailiwick Rule:** allow answers only for subdomains
    - ns.eve.com can't answer for download.com
    - Some resolvers cache combo: eve.com NS-IP 6.6.6.6
- ❑ **Method 2:** send from corrupt NS
  - Some referral-chains are very long
    - Transitive trust relationships
  - Some (many?) NSs run vulnerable versions!
- ❑ **Method 3:** send spoofed DNS response
- ❑ **Method 4:** dangling DNS records

# Gratuitous glue RR in Responses

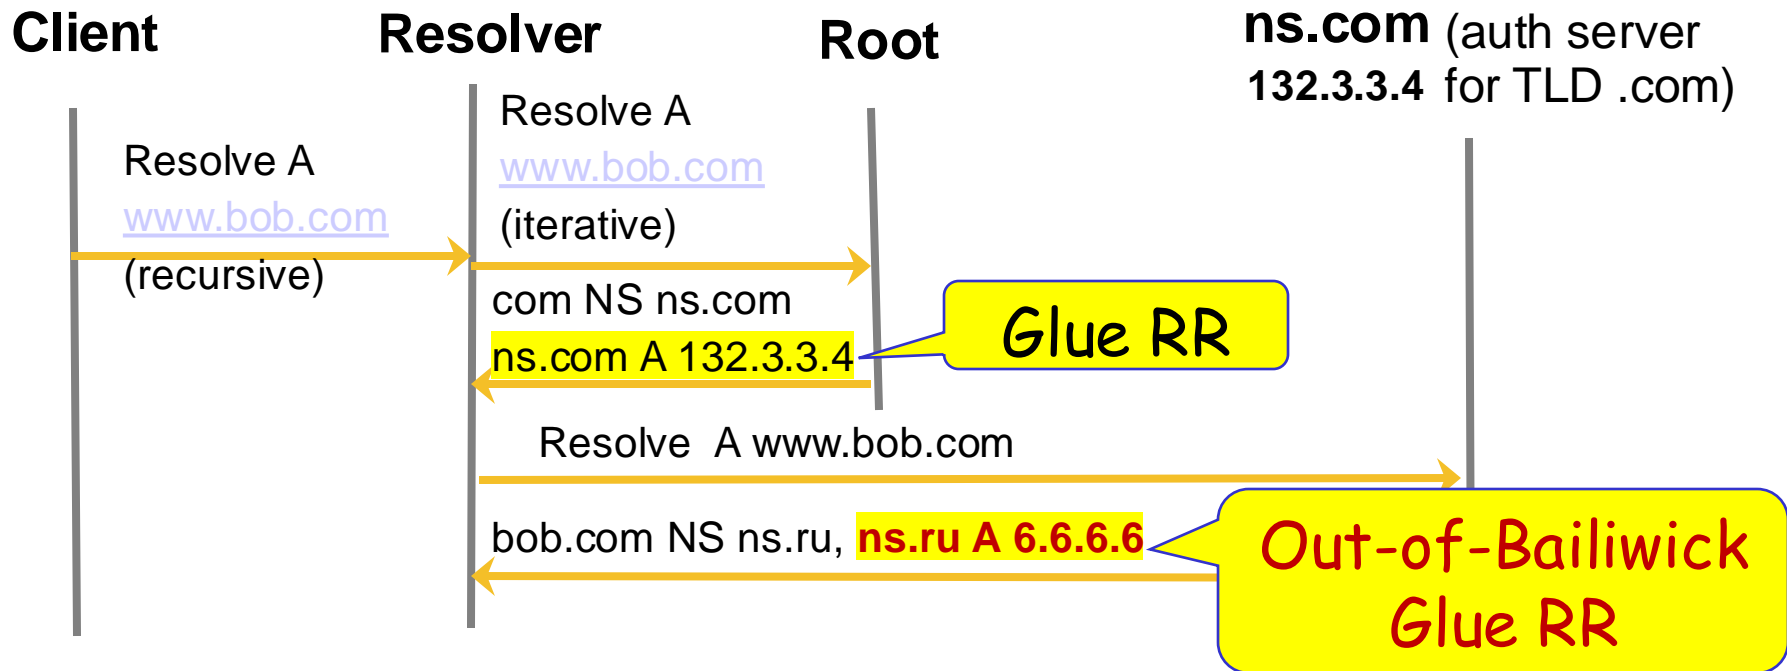
- Normally: RR is received to fulfil request
- **Gratuitous RR**: received **without** request
  - In response to different request, or appended to a request
- Main use: glue (A) RR , sent with a referral (NS)





# Out-of-Bailiwick Gratuitous Glue RR

- **Out-of-Bailiwick Gratuitous RR**: a RR for one domain, received from a name server for another
- **Can be abused for DNS cache poisoning!**
  - Out-of-bailiwick RRs ignored since ~1997



# Out of Bailiwick Attack [historical]


Recursive  
Resolver



Attacker  
6.6.6.6



Name server  
ns.ebay.com



A ? attacker.com

attacker.com A 6.6.6.6  
attacker.com NS ns.ebay.com  
ns.ebay.com A 6.6.6.6

A?www.ebay.com

# Bailiwick Rule:

## Accept Records Only within Bailiwick

- Prevents out-of-Bailiwick attack
- A Name Server can respond for records within its domain
  - Root servers can return any record
  - TLD server can return anything within that TLD
  - ns.bank.com can return anything for bank.com
- Resolvers ignore out-of-bailiwick responses
  - Some resolvers use out-of-bailiwick for glue (resolution of NS), and do not cache it (no impact on other queries)
  - Or, cache combo, e.g.: eve.com NS-IP 6.6.6.6


# DNS Poisoning: How?

---

- ❑ DNS Poisoning: `inject' fake DNS RR
- ❑ Method 1 (historical): by Gratuitous `glue' RR
  - e.g. query A www.eve.com, response: eve.com NS download.com and download.com A 6.6.6.6
  - **Bailiwick Rule**: allow answers only for subdomains
    - ns.eve.com can't answer for download.com
    - Some resolvers cache combo: eve.com NS-IP 6.6.6.6
- ❑ **Method 2: send from corrupt NS**
  - Some referral-chains are very long
    - **Transitive trust** relationships
  - Some (many?) NSs are vulnerable (e.g., unpatched)
- ❑ Method 3: send spoofed DNS response
- ❑ Method 4: dangling DNS records

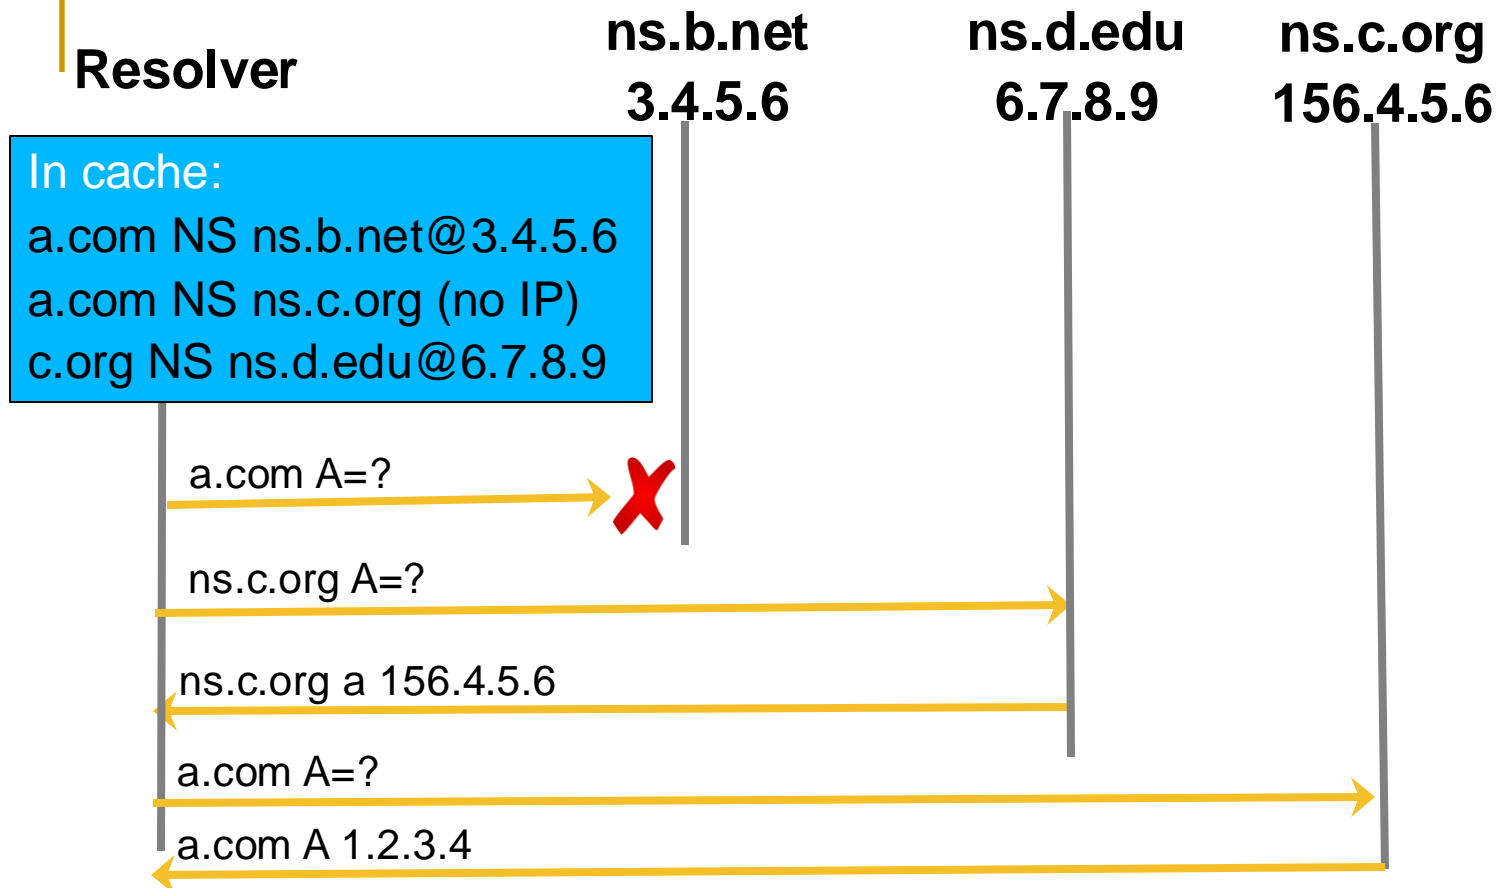
# DNS Trust Relationships

- Often, some name servers of a domain are from other domains (for backup, performance)
  - a.com has name server (also) in b.net
  - Resolvers identify & use the 'best performing' name server



# DNS: Transitive Trust

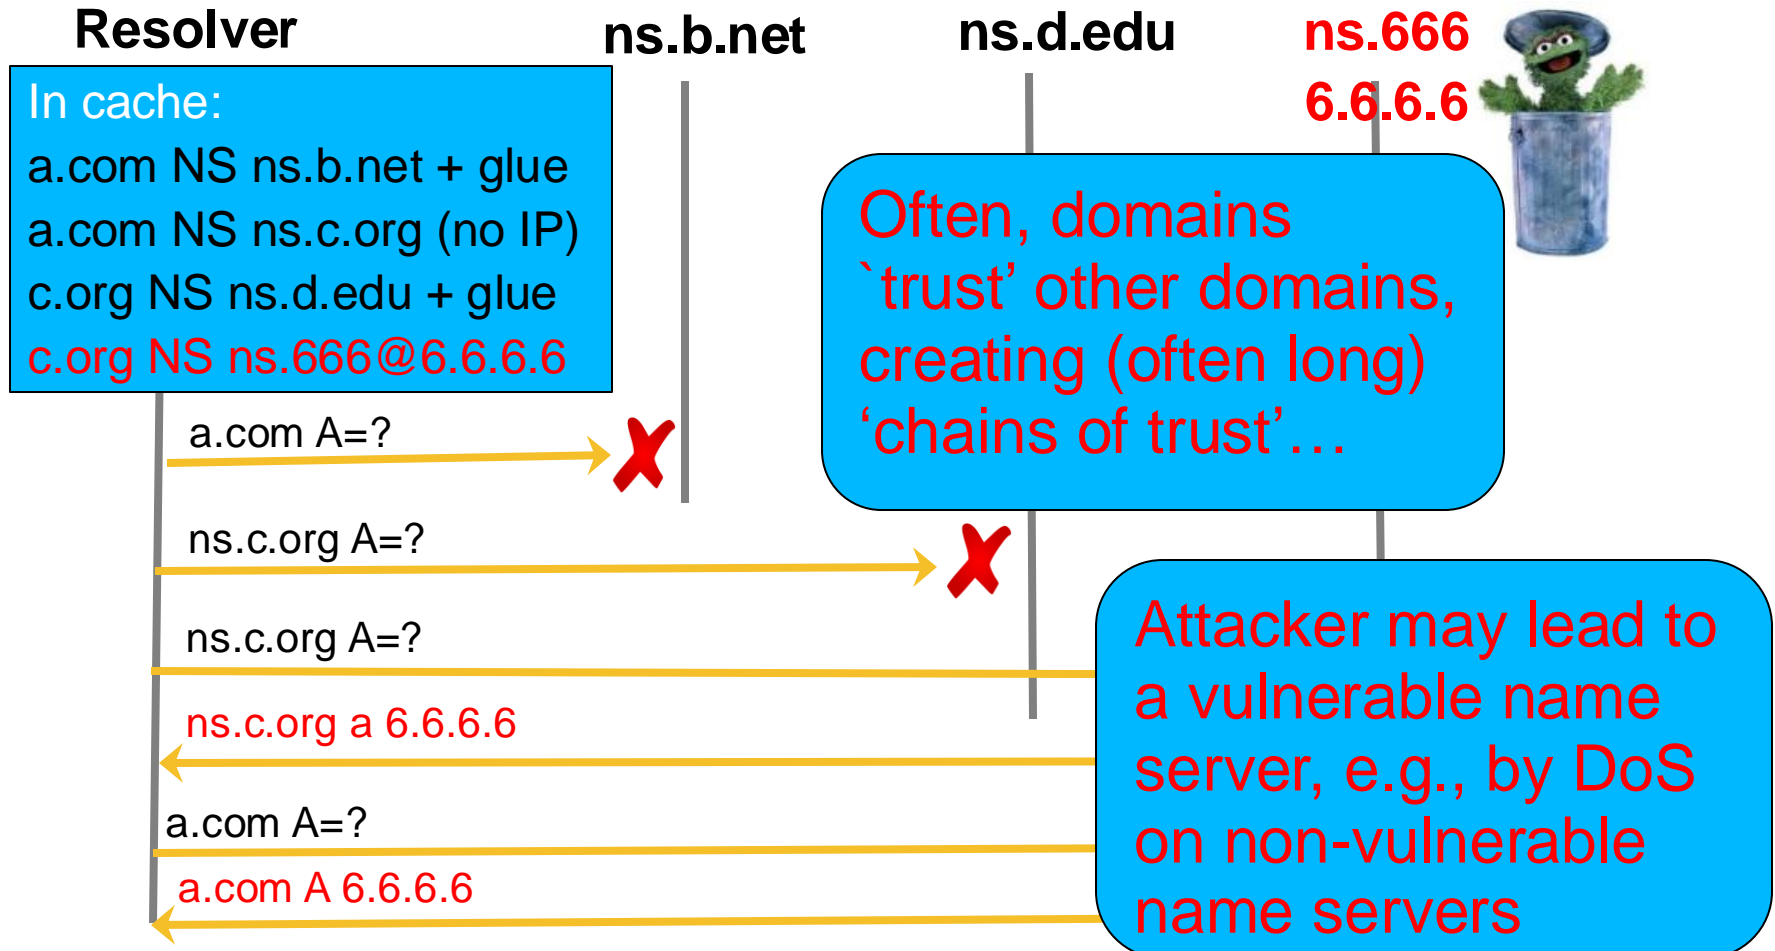
- Resolver may need to find IP of name servers, too



a.com trusts b.net, c.org and c.org trusts d.edu  
→ a.com also trusts d.edu

# DNS: Transitive Trust Attack

- Off-path attacker, can drop packets [how?]
- Controls 'transitively-trusted' NS, e.g., ns.666



# Domain Name System (DNS) Security

---

- DNS: quick recap
- DNS poisoning
  - Motivation: the hacker's swiss knife
  - Method 1 (historical): by Gratuitous `glue' RR
  - Method 2: send from corrupt NS
  - **Method 3: send spoofed DNS response**
  - Method 4: dangling DNS records
- DNSSEC: Cryptographic security for DNS
- DNS Privacy issues and defenses




# Off-path DNS poisoning attacks

---


- DNS is clearly vulnerable to MitM attacks
  - Against MitM, use DNSSEC (crypto-defense)
  - Or... **assume attacker is only off-path**
- But... DNS vulnerable to off-path attacks too!
  - Motivating DNSSEC: use signatures, hashing (against MitM!)
- We present **Spoofed-response attacks**:
  - Concept and challenges
  - Kaminisky's attack
  - Improved DNS security ('post-Kaminsky' defenses)
  - Source port de-randomization attacks ('post-Kaminsky' attacks)

# DNS Poisoning by Spoofed Response




Challenge 1: **could** attacker send a query to the resolver? Often, yes:


- Open resolver (anyone can ask)
- Incoming email validation
- Hyperlink (e.g., <IMG>) in visited website
- Compromised machine (Zombie/bot)



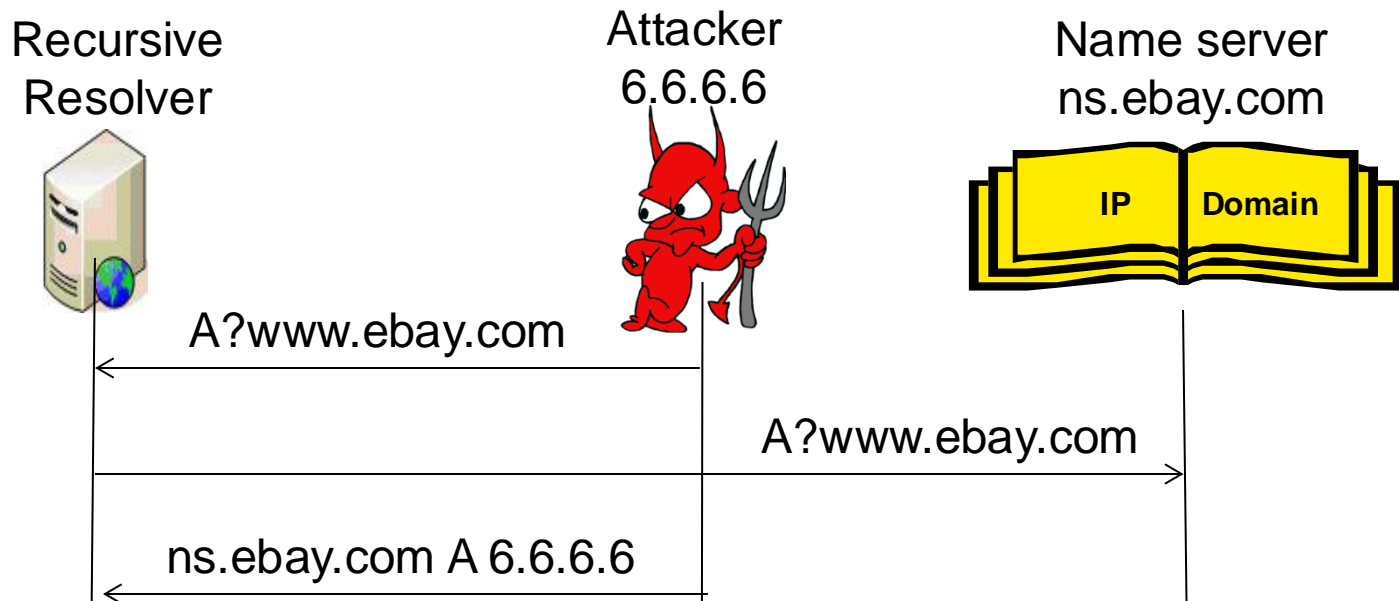
# DNS Poisoning by Spoofed Response



Challenge 2: resolver would not send a query if record is already cached!  
- solutions: wait, clear cache, or ask for name not in cache (more later)



# DNS Poisoning by Spoofed Response



Challenge 3: would resolver accept the spoofed response?

Challenge 3A: source IP address should be of NS to which request was sent!

**Solution:** attacker sends packet with the source IP address of NS (**spoofed**)

Challenge 3B: response header should contain the same (16-bit) TXID field as in the header of the request

**Solution:** send many spoofed responses, with different TXID values...



# DNS Poisoning by Spoofed Response

- Attacker sends many responses with different TXID values
- But, **must win the race:**
  - Legit response surely has correct TXID
  - Would the legit response arrive before the 'lucky' spoofed response (with correct TXID)?
  - First correct response is accepted and cached
    - For TTL time (minute, hour, day,...)
      - TTL is the duration of validity of a given mapping
  - Subsequent, incorrect D are ignored

➔ **Winning the race considered impractical...  
Until Kaminsky's attack [2008]**

RIP Dan Kaminsky: 1979-2021

# Kaminsky's attack: four ideas

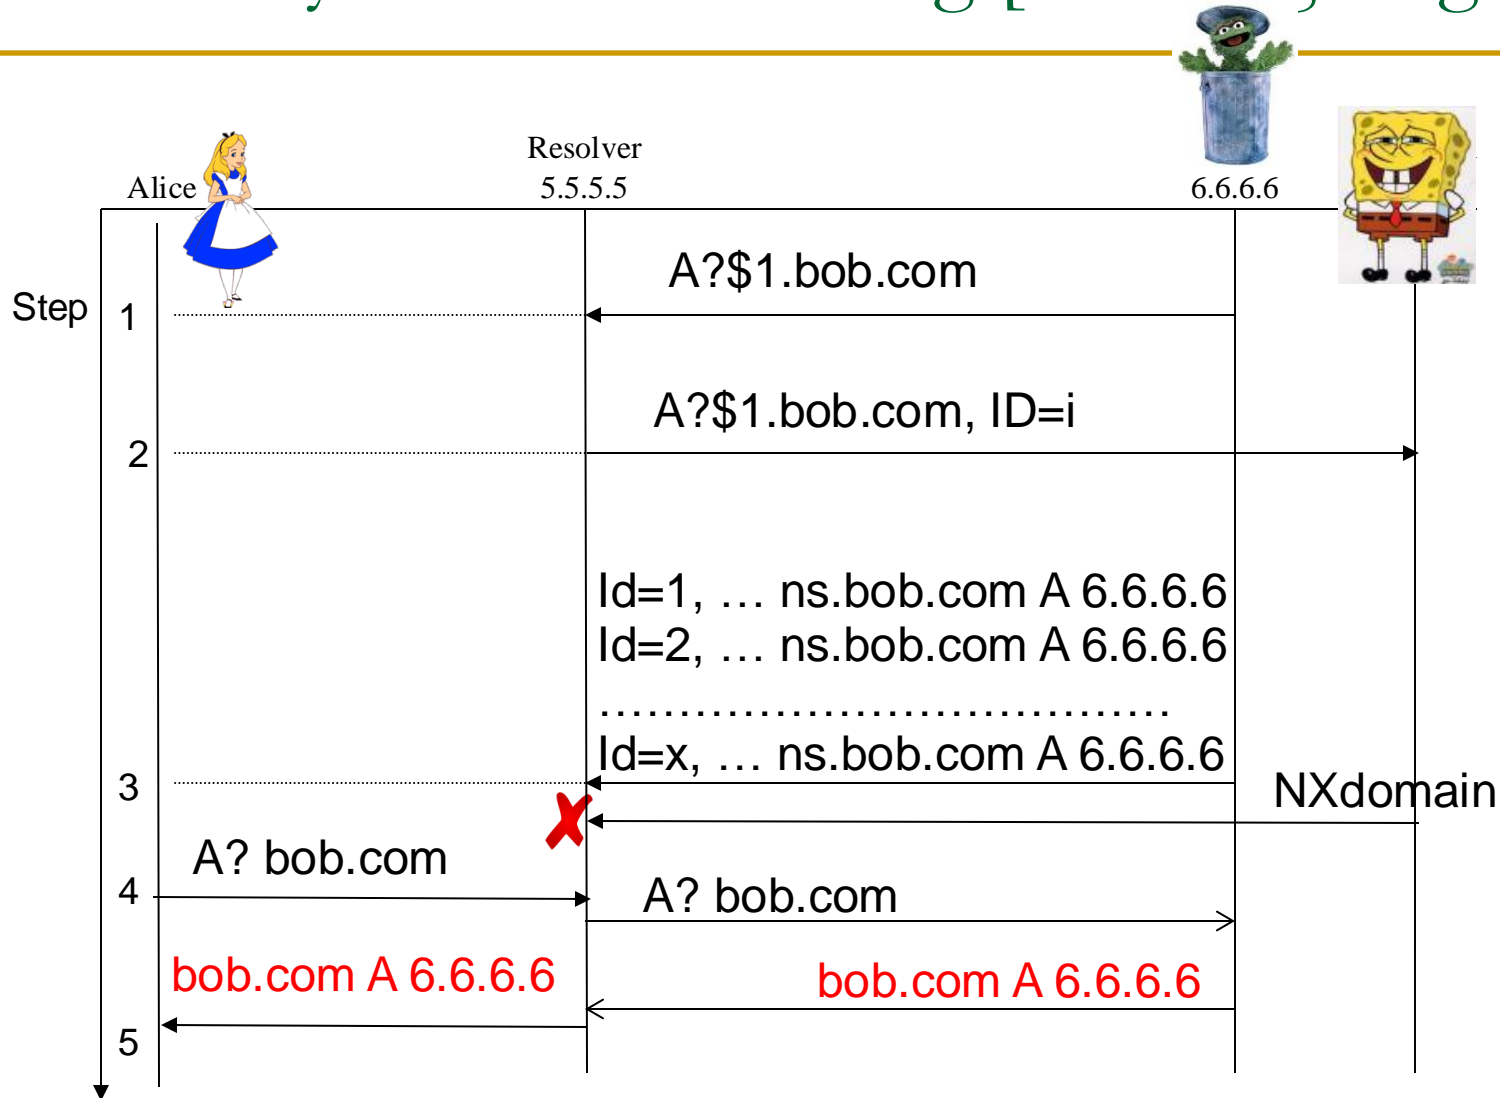
---

- Challenge 1: ensuring response not in cache
- Idea 1: query is for a non-existing domain → surely not in cache!
- Challenge 2: only a limited number of spoofed response can arrive before the legit response
- Idea 2: no problem, just try again with a different non-existing domain!
- Challenge 3: how to exploit poisoning response to a non-existing domain?
  - Exercise: steal cookies, even with 'secure' attribute (get a cert)
  - Kaminsky's (3<sup>rd</sup>) idea: **poisoned glue**

# Kaminsky's 3<sup>rd</sup> Idea: poisoned glue (or NS)

- ❑ Query: 'A' record for (say) 585.bank.com
- ❑ Spoofed response:
  - Option 1, **poisoned glue**:  
Answer section: empty (or an answer)  
Authority section: bank.com NS ns.bank.com  
Glue (additional section): **ns.bank.com A 6.6.6.6**
  - Option 2, **poisoned NS**:  
Answer section: empty (or an answer)  
Authority section: **bank.com NS ns.666.com**  
Additional section: (none or a glue)
- ❑ Caching **ns.bank.com A 6.6.6.6** or **bank.com NS ns.666.com** → entire bank.com domain poisoned
- ❑ Would they cache? Depends on resolver
  - Cautious resolvers don't cache referral glue and NS RRs
  - Other spoofing payloads, e.g., use CNAME [KSW'17]

# Kaminsky's DNS Poisoning [showing single query]



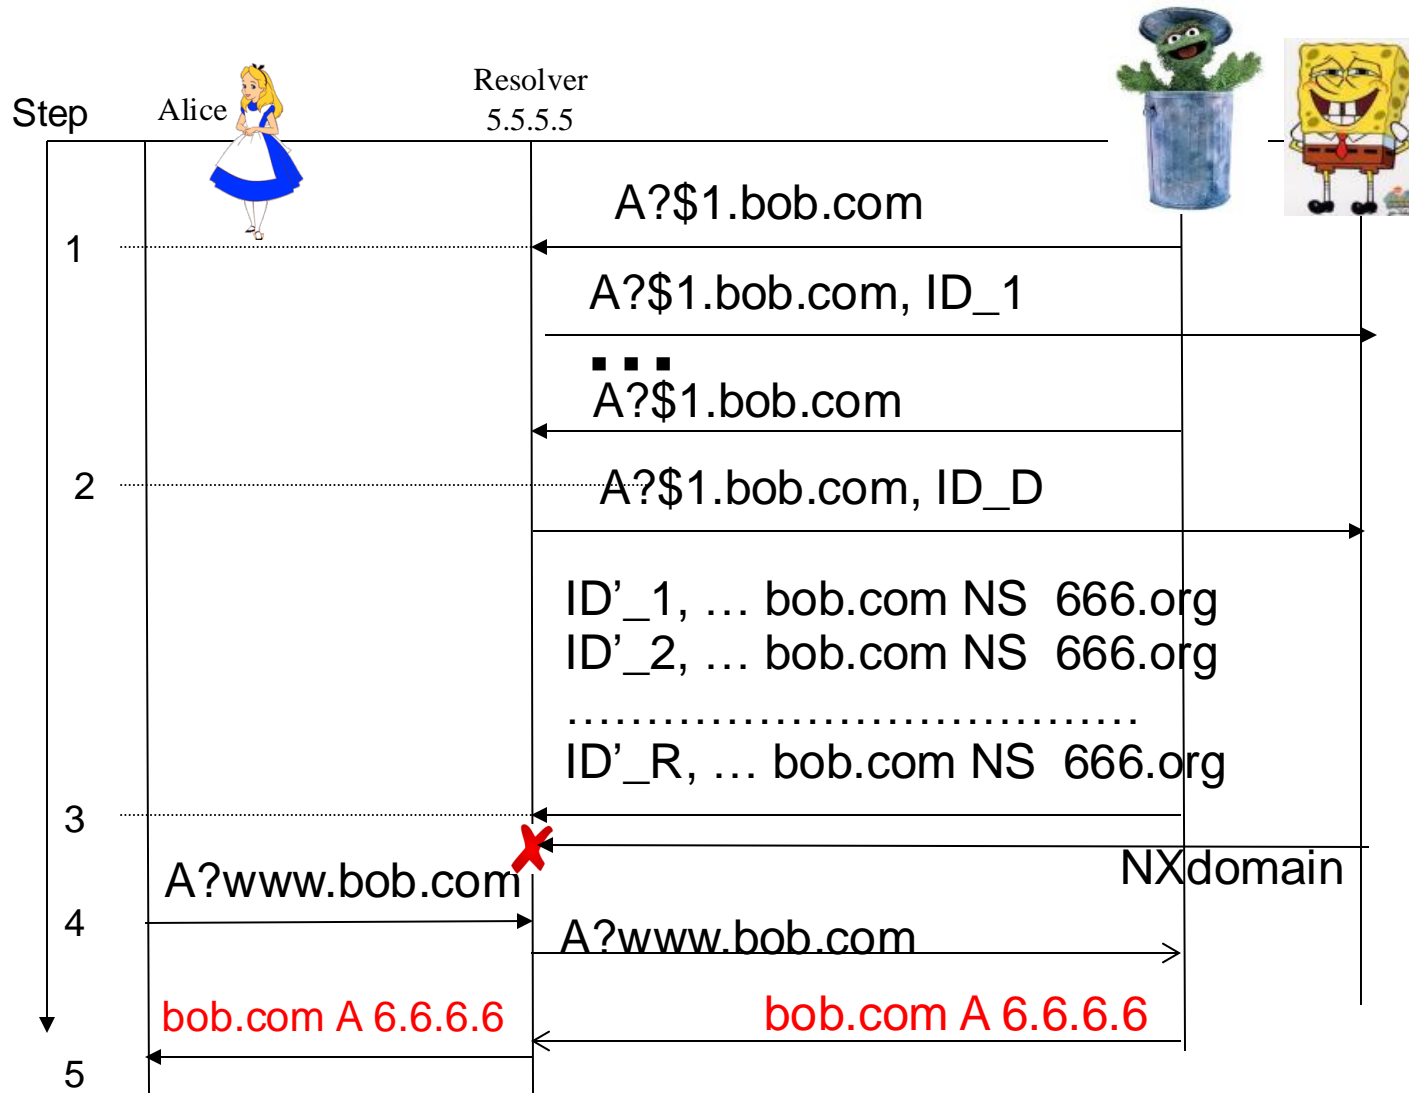
If none of the x responses match... repeat with \$2 subdomain!  
Actually, Kaminsky had one more idea!



## 4<sup>th</sup> idea: duplicate requests → birthday!

- Attacker sends  $D$  duplicate requests
- All for exactly the same query (domain)
- Some resolvers send  $D$  identical requests
- Probability of a response to match:  $D/I$ 
  - Match: 16-bit DNS-ID, ports, IP as in a request
  - If only DNS-ID is random:  $I = 2^{16}$
- Attacker sends  $R$  responses
- Probability of one of them to match?
- Birthday paradox:  $p \lesssim \min\left(1, \frac{D \cdot R}{I}\right)$

# Kaminsky's Birthday Attack (duplicate requests)



# Defenses against Kaminsky's Attack

---

- Don't cache referral glue and NS RRs
- RFC 5452: Local server must validate as follows...
  - Same question section as in request
  - Response received within reasonable delay
  - Ignore if already received valid response for query
  - Same (16-bit) ID field
    - Local server **must choose ID randomly**
- $\text{Resp}(\text{Src.IP}) = \text{Req}(\text{Dst.IP})$ 
  - Most domains have 1 to 3 likely-to-be-used name servers from a given resolver
- $\text{Resp}(\text{Dest.IP}, \text{Dest.port}) = \text{Req}(\text{Src.IP}, \text{Src.port})$ 
  - **Main Defense: Source Port Randomization (SPR), i.e., resolver selects random source port**
  - **Preferably, also randomize source IP**

# Source Port Randomization (SPR)

---

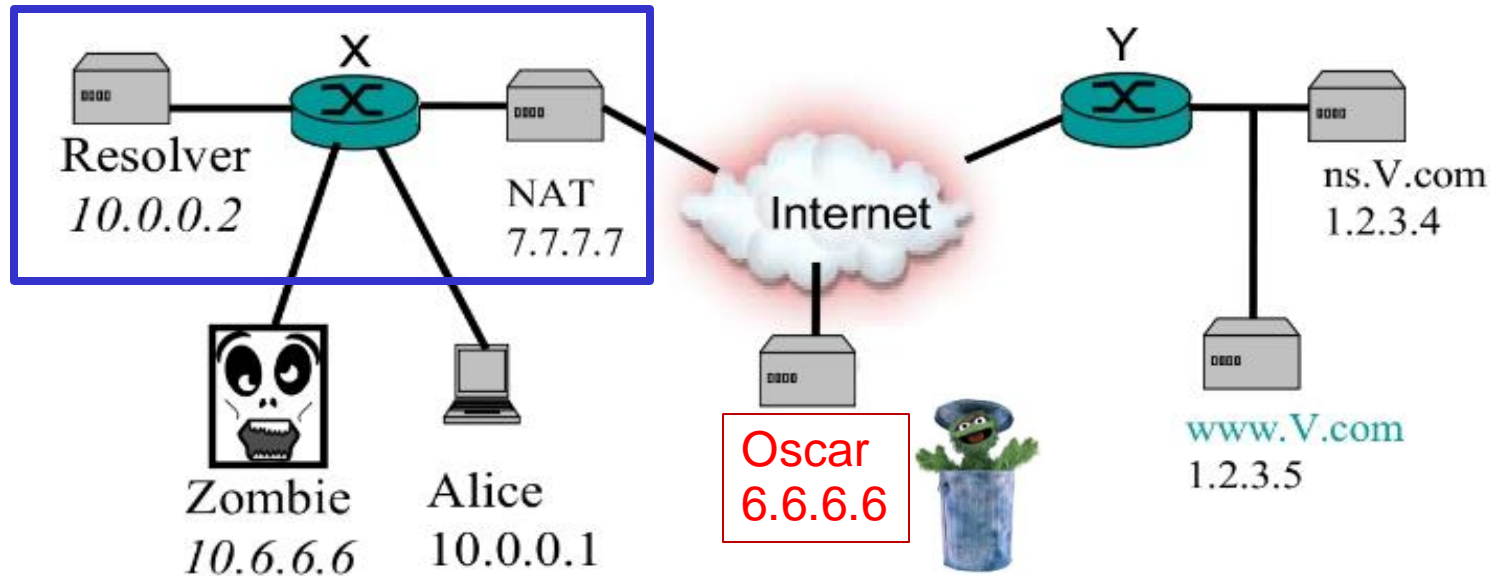
- Send requests from random/unpredictable ports [Bernstein2002, RFC5452]
  - Port field - 16 bits
  - Increases the search space: 16 bit ID, 16 bit port  $\rightarrow \sim 2^{32}$
- Makes Kaminsky's attack less practical
  - Birthday attack:  
 $2^{16}$  requests +  $2^{16}$  responses  $\rightarrow$  poison with prob.  $\sim 1/2$
  - Many (most?) resolvers also prevent/limit birthday attack
    - Detect new query identical to pending one  $\rightarrow$  don't resend

# Source Port Randomization (SPR)

---

- Send requests from random/unpredictable ports [Bernstein2002, RFC5452]
  - Port field - 16 bits
  - Increases the search space: 16 bit ID, 16 bit port  $\rightarrow \sim 2^{32}$
- Makes Kaminsky's attack less practical
  - $2^{16}$  requests +  $2^{16}$  responses  $\rightarrow$  poison with prob.  $\sim 1/2$
  - Many (most?) resolvers also prevent/limit birthday attack
  - Reduced motivation to deploy DNSSEC ? ☹
- Several **Source Port De-Randomization** attacks
  - Next: resolver behind NAT source port de-randomization
  - In IP security: fragmentation-based de-randomization
  - In papers: other source port de-randomization attacks
    - Recent, very effective: SadDNS attack(s); and others

# Resolver behind NAT




Typical home/office network:

- Includes a resolver and different hosts, some possibly 'zombies' (controlled by attacker)
- Connected via a NAT device (IP: 7.7.7.7)
- Just in case, let us explain what's a NAT

# NAT: Network Address Translation

Goal: share IP addresses among multiple hosts in Net  
Net uses private ranges IPs (10.0/8, 172.16/12, 192.168/16)  
NAT uses one 'external' IP; maps internal (IP:port) pairs to an external port



How are external ports allocated ?

# NAT: External Port Allocation Methods

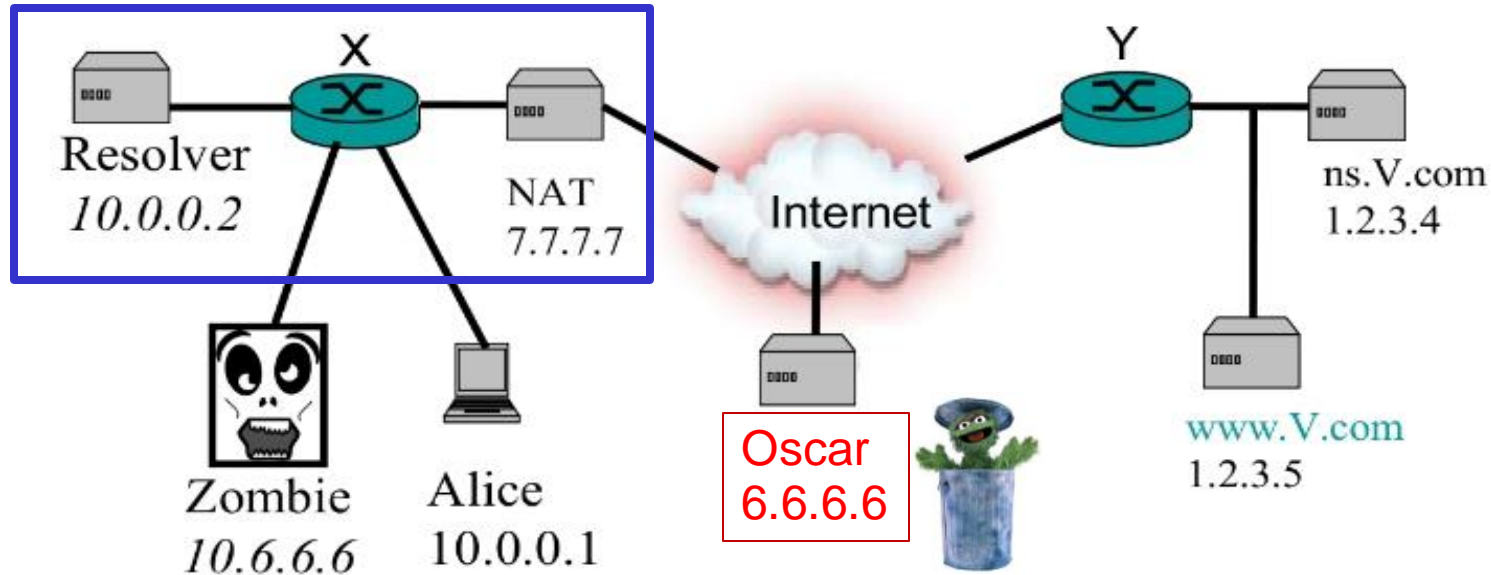
- Goal: different external ports per connection
  - Assume up to  $2^{16} = 65536$  concurrent connections
- Naïve solution 1: on the  $p_x$  connection to  $x$ , use port  $p_x$ 
  - Where  $x=(IP:port)$  is a destination
- Disadvantages?
  - Need to maintain  $p_x$  forever: restart  $p_x$  if not used/saved
  - Security concern?
  - Attacker can guess value of  $p_x$
  - This foils source port randomization (SPR)
    - ➔ attacker can use Kaminiski's attack (against SPR resolver)
- Naïve solution 2: select a random port for each connection
  - Disadvantage?
  - Collisions (birthday paradox: after \_\_\_ concurrent connections)



# Per-Dest-Incrementing Port Allocation

- Goal 1: different external ports per connection
- Goal 2: off-path attacker cannot predict next port
- Naïve solution 3: select random ports, remember ports in use (avoid collisions)
  - Inefficient (remember all ports)
- Alternative: Per-Dest-Incrementing port allocation:
  - On connection to  $x$ , increment port  $p_x \pmod{2^{16} = 65536}$
  - If no current  $p_x$ , select a random (or pseudo-random) value
    - One way to select pseudo-random  $p_x \leftarrow PRF_k(x||time)$
    - A common simplification using hash:  $p_x \leftarrow h(k||x||time)$
  - Very common port-allocation method (e.g., Linux)

# Derandomizing SPR: Resolver behind NAT

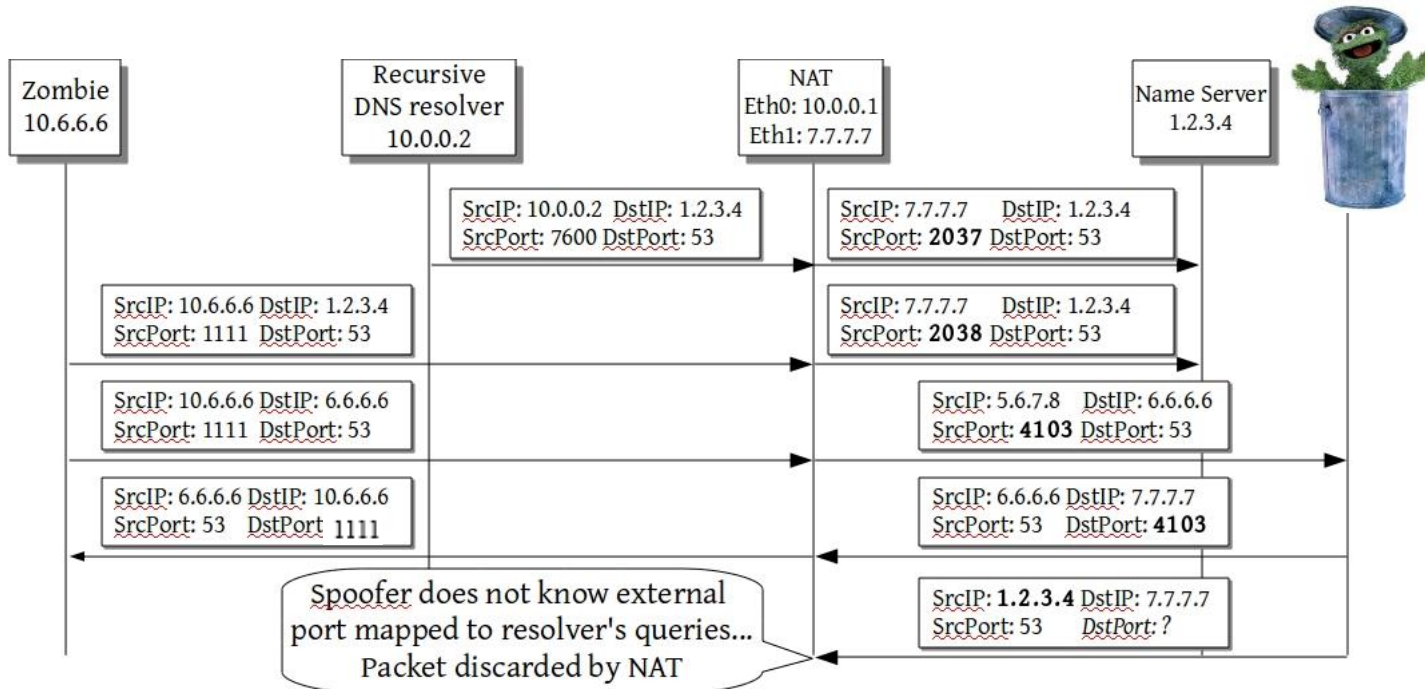



Typical home/office network:

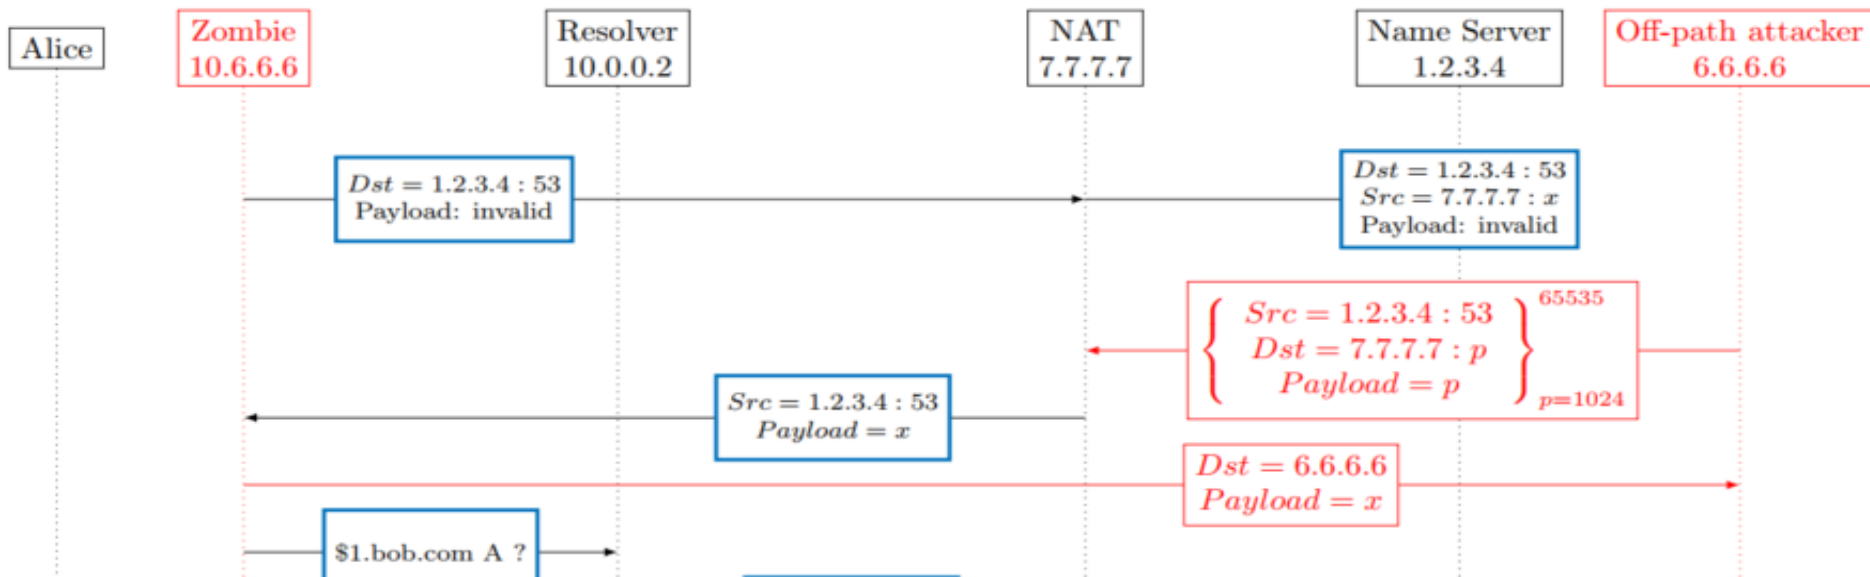
- Connected via a NAT device (IP: 7.7.7.7)
- Attack assumptions:
  - Network includes a Zombie - compromised host
  - NAT uses **per-dest-incrementing port allocation**
    - A (less efficient) variant works without this assumption
    - Even more efficient variant for globally-incrementing ports

# Ports for queries via NAT

- For per-dest incrementing ports NAT
  - For each destination: select first port at random, then choose next available port
- Can attacker predict port assigned by NAT?



# Source Port deRandomization



- Zombie sends a packet to port 53 of the name server
- NAT allocates port  $x$  to port 53 of the name server
- Spoofer sends a packet with payload  $p$  to each port  $p \in \{1, \dots, 65536\}$
- Only the one sent to port  $x$  gets through to zombie (with payload  $x$ )
- Zombie echoes current port ( $x$ ) to attacker
- Attacker continues with Kaminsky's attack (and known source port!)



# Domain Name System (DNS) Security

---

- DNS: quick recap
- DNS poisoning
  - Motivation: the hacker's swiss knife
  - Method 1 (historical): by Gratuitous `glue' RR
  - Method 2: send from corrupt NS
  - Method 3: send spoofed DNS response
  - **Method 4: dangling DNS records**
- DNSSEC: Cryptographic security for DNS
- DNS Privacy issues and defenses

# Dangling DNS records

- Recall: several DNS records map resources to domain names, including CNAME, NS and MX, e.g.:
  - foo.com NS ns1.CDN.net
  - [www.fee.com](http://www.fee.com) CNAME server765.cloud.com
- Dangling DNS record: domain freed, record remains
  - E.g., using a new NS: foo.com NS ns.foo.com
  - But forgot to remove old NS record to ns1.CDN.net
    - Or old NS records is still in cache of resolvers
- Attacker may get control of old domain (ns1.CDN.net)
  - E.g., when domain name allocated by CDN/cloud to customers
  - Allows to map resource to attacker-controlled IP !

# Domain Name System (DNS) Security

---

- DNS: quick recap
- DNS poisoning
  - Motivation: the hacker's swiss knife
  - Method 1 (historical): by Gratuitous `glue' RR
  - Method 2: send from corrupt NS
  - Method 3: send spoofed DNS response
  - Method 4: dangling DNS records
- **DNSSEC: Cryptographic security for DNS**
- DNS Privacy issues and defenses



# DNS Security (DNSSEC) [RFCs 4033 to 4035]

---

- Goal: prevent MitM (and off-path) attacks
  - Also defends against exploits of vulnerable name servers
- Main idea: **sign DNS RRsets**
  - **RRset: set of Resource Records with same name and type**
    - E.g., all 'A' records of foo.com
  - Signature in a separate DNS RR, of type RRSIG
  - Private signing key of domain owner [offline?]
    - authoritative DNS server
  - Hierarchical certification of domains' public key
- ➔ prevents cache poisoning attacks
- ➔ secures use of DNS-based policies, public keys
- Includes a (very limited) 'DNS-specific' PKI

# Domain Name System with DNSSEC

Client

DNSSEC validation at **client**?

- Essential if resolver isn't trusted
- But 'breaks' recursive service
- Requires support from OS, resolver
- Supported by some clients, incl. popular browsers
- Some resolvers 'break' DNSSEC
  - Esp. : NXDomain → Ad

www.ebay.com



1.2.3.4

A?www.ebay.com


Is valid?



Recursive  
Resolver

A?www.ebay.com

www.ebay.com A 1.2.3.4



Name Server  
NS.ebay.com

# DNSSEC: Public Keys and Record Types

---

- Typical domains have two public keys: **KSK** and **ZSK**
- **Key Signing Keys (KSK)**
  - Authenticated by parent domain, or by operator/software ('trust anchor', e.g., for root domain)
- **Zone Signing Keys (ZSK)**
  - Authenticated (signed) by KSK of domain
- Domain may use multiple ZSKs/KSKs, to phase-in a new key or for different signing algorithms
- Basic DNSSEC Record types:
  - **DNSKEY**: public keys (mainly, KSKs and ZSKs)
  - **RRSIG**: signatures (over an RRset, e.g., DNSKEY RRset)
  - **DS**: hash to authenticate KSK's DNSKEY (Hash-then-Sign)

# The RRSIG, DS and DNSKEY RRs

## Root Zone (KSK embedded in browsers, tag=3861)

. 86400 IN DNSKEY 256 3 5 (... ..) ; ZSK1 (tag=1127)  
. 86400 IN DNSKEY 256 3 5 (... ..) ; ZSK2  
. 86400 IN RRSIG DNSKEY 5 0 86400 (... 3861 ...)  
com. 86400 IN RRSIG DS 5 1 86400 (... 1127 ...)  
com. 86400 IN DS 50322 5 1 (2BB18 ... A53B0A)

Signature  
over both  
DNSKEYs  
using KSK  
3861

Signature of...  
(using ZSK 1127)

Hash of  
KSK 50322

## com Zone

com. 6400 IN DNSKEY 257 3 5 (... ..) ; KSK (50322)  
com. 6400 IN DNSKEY 256 3 5 (... ..) ; ZSK1 (2623)  
com. 6400 IN DNSKEY 256 3 5 (... ..) ; ZSK2 (14672)  
com. 6400 IN RRSIG DNSKEY 5 1 6400 (... 50322 ...)  
www.com. 3600 IN A 1.2.3.4 ; A RR  
www.com. 3600 IN RRSIG A 5 2 (... 2623 ... ) ; Signs A RRset

Signature  
over all  
DNSKEYs  
using KSK  
50322.

Signature over type-A RRset using ZSK 2623

# DNSSEC Resource Record (RRs): RRSIG

- **RRSIG**: signature over an **RRset** and validity period
  - Signature over the concatenation of all RRs of given name, type
  - Sent automatically with response to request for RRset, if request requests DNSSEC records (DNSSEC Ok (DO) bit)
  - Example: signature over set of RRs of type **A** of **www.foo.com**:  
**www.foo.com. 6400 IN RRSIG A 5 3 6400**  
**20230218173103** (20220218083059 **2624** **foo.com.**  
**oJB1W6WNGv+ldv.....DQfsS3Ap3o=**)
  - TTL of signature, RRset is 6400s (1hour) [usually identical]
  - Signing algorithm is '5' (RSA/SHA1)
  - Number of 'labels' in signed name is 3 (www.foo.com).
  - **Expiration time is 17:31:03 at 2/18/2023;**  
signing time is 08:30:59 at 2/18/2022
  - **2624** is the 'key tag' - identifies signer's key (may have few).  
For most algs: 16-bits sum of the data of the DNSKEY data
  - **Signer is foo.com** and oJB1... is the Base64-encoded signature

# DNSSEC Resource Record (RRs): RRSIG

- **RRSIG**: signature over an **RRset** and validity period
  - Signature over the concatenation of all RRs of given name, type
  - Sent automatically with response to request for RRset, if request requests DNSSEC records (DNSSEC Ok (DO) bit)
  - Example: signature over set of RRs of type **A** of www.foo.com:  
www.foo.com. 6400 IN **RRSIG A** 5 3 6400  
20230218173103 (20220218083059 2624 foo.com.  
oJB1W6WNGv+ldv.....DQfsS3Ap3o=)
  - TTL of signature, RRset is 6400s (1hour) [usually identical]
  - **Expiration time is 17:31:03 at 2/18/2023;**  
signing time is 08:30:59 at 2/18/2022
- Actually, something is weird here! What?
  - Expiration should be longer than TTL, but by so much??
  - Attacker can resend the signature anyway
  - To ensure updates, validity shouldn't be much more than TTL

# The RRSIG, DS and DNSKEY RRs

## Root Zone (KSK embedded in browsers, tag=3861)

. 86400 IN DNSKEY 256 3 5 (... ..) ; ZSK1 (tag=1127)

. 86400 IN DNSKEY 256 3 5 (... ..) ; ZSK2

. 86400 IN RRSIG DNSKEY 5 0 86400 (... 3861 ...)

com. 86400 IN RRSIG DS 5 1 86400 (... 1127 ...)

com. 86400 IN DS 50322 5 1 (2BB18 ... A53B0A)

Signature  
over both  
DNSKEYs  
using KSK  
3861

Signature of...  
(using ZSK 1127)

Hash of  
KSK 50322

Why??

## com Zone

com. 6400 IN DNSKEY 257 3 5 (... ..) ; KSK (50322)

com. 6400 IN DNSKEY 256 3 5 (... ..) ; ZSK1 (2623)

com. 6400 IN DNSKEY 256 3 5 (... ..) ; ZSK2 (14672)

com. 6400 IN RRSIG DNSKEY 5 1 6400 (... 50322 ...)

www.com. 3600 IN A 1.2.3.4 ; A RR

www.com. 3600 IN RRSIG A 5 2 (... 2623 ... ) ; Signs A RRset

Signature  
over all  
DNSKEYs  
using KSK  
50322.

Signature over type-A RRset using ZSK 2623

# DNSSEC Resource Record (RRs): DNSKEY

- RRSIG: signature over an RRset and validity period
- DNSKEY: DNS public verification key(s), e.g.:

foo.com. 86400 IN DSNSKEY 256 3 5 (Abl4...z3w==)  
(domain) (TTL) (flags) (prot) (alg) (public key )

- 16-bit flags field; only bits 7 and 15 used (other: zeros)
  - Details in next slide
- 3 is 'protocol' field (only value allowed is 3)
- 5 is 'signature algorithm' field; 5 is RSA-SHA1.  
Other options defined (by IANA), e.g., ECDSA
- Public key is BASE-64 encoded
- Note: keytag is not specified; it is computed as a (known) function of the public key
  - If multiple keys have the same tag, try all of them [KeyTrap]



# DNSSEC Resource Record (RRs): DNSKEY

- RRSIG: signature over an **RRset** and validity period
- DNSKEY: DNS public verification key(s), e.g.:  
foo.com. 86400 IN DSNSKEY 256 3 5 (Abl4...z3w==)  
(domain) (TTL) (flags) (prot) (alg) (public key )
  - 16-bit flags field; only bits 7 and 15 used (others must be zero)
    - Bit 7 signals key validates RRsigs
    - Bit 15 signals Secure Entry Point (SEP): key validated by parent zone or 'pre-trusted' by resolver; one KSK must have SEP set
    - Flag bits are numbered from left (MSb) to right (LSb: bit 15)
    - Values: 256 (bit 7 - RRSIG key), 257 (also SEP)
  - Typically, we use two type of keys:
    - Key Signing Keys (KSK), Secure Entry Point (SEP), flags: 257, used to validate DNSKEY RRSIGs;
    - Zone Signing Keys (ZSK), flags: 256, validate (other) RRSIGs, validated by KSK

# DNSSEC Resource Record (RRs): DNSKEY

- RRSIG: signature over an **RRset** and validity period
- DNSKEY: DNS public verification key(s) ), e.g.:  
foo.com. 86400 IN DSNSKEY 256 3 5 (Abl4...z3w==)
- 16-bit flags field; only bits 7 and 15 used (other: zeros)
  - Bit 7 signals key validates RRsigs
  - Bit 15 signals Secure Entry Point (SEP)
- Key Signing Keys (KSK), flags=257, for DNSKEY RRSIGs
- Zone Signing Keys (ZSK), flags=256, for other RRSIGs
- Why separate KSK from ZSK?
- ZSK used more → more exposed → change often
- ZSK roll-over: keep new ZSK in parallel to old ZSK
- Changing KSK requires re-validation by parent zone
- KSK is typically longer, kept offline, changed rarely
- Parent domain validates SEP keys (KSKs) using the DS RR

# The RRSIG, DS and DNSKEY RRs

## Root Zone (KSK embedded in browsers, tag=3861)

. 86400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK1 (tag=**1127**)

. 86400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK2

. 86400 IN **RRSIG** **DNSKEY** 5 0 86400 (... **3861** ...)

com. 86400 IN **RRSIG** **DS** 5 1 86400 (... **1127** ...)

com. 86400 IN **DS** **50322** 5 1 (2BB18 ... A53B0A)

Signature  
over both  
DNSKEYs  
using KSK  
**3861**

Signature of...  
(using ZSK 1127)

Hash of  
KSK **50322**

## com Zone

com. 6400 IN **DNSKEY** 257 3 5 (... ..) ; KSK (**50322**)

com. 6400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK1 (**2623**)

com. 6400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK2 (14672)

com. 6400 IN **RRSIG** **DNSKEY** 5 1 6400 (... **50322** ...)

www.com. 3600 IN **A** 1.2.3.4 ; A RR

www.com. 3600 IN **RRSIG** **A** 5 2 (... **2623** ...) ; Signs A Rrset

Signature  
over all  
DNSKEYs  
using KSK  
50322.

Signature over type-A RRset using ZSK 2623

# DNSSEC Resource Record (RRs): DS

- RRSIG: signature over an **RRset** and validity period
- DNSKEY: public key of zone, mainly:
  - Zone Signing Key (ZSK): validates non-key RRsets
  - Key Signing Key (KSK): validates ZSKs
- DS (Delegation Signer): hash of SEP(KSK) of subdomain
  - Essential to allow separate owner / zone file for subdomain
  - Example of DS for foo.com (in zone file of .com):  
foo.com 86400 IN DS 50322 5 1 (2BB18 ... A53B0A)
  - TTL of key is 86400s (24hours)
  - 50322 is the 'key tag' - identifies signer's key (may have few).  
For most algs: 16-bits sum of the data of the DNSKEY data
  - Signing algorithm is '5' (RSA/SHA1) and hashing is '1' (SHA1).
  - 2BB18 ... A53B0A ... is the Base64-encoded hash, over owner name concatenated to DNSKEY data
  - DNSKEY should specify Secure Entry Point (flags=257)
  - h

# The RRSIG, DS and DNSKEY RRs

## Root Zone (KSK embedded in browsers, tag=3861)

. 86400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK1 (tag=**1127**)

. 86400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK2

. 86400 IN **RRSIG** **DNSKEY** 5 0 86400 (... **3861** ...)

com. 86400 IN **RRSIG** **DS** 5 1 86400 (... **1127** ...)

com. 86400 IN **DS** **50322** 5 1 (2BB18 ... A53B0A)

Signature  
over both  
DNSKEYs  
using KSK  
**3861**

Signature of...  
(using ZSK 1127)

Hash of  
KSK **50322**

## com Zone

com. 6400 IN **DNSKEY** 257 3 5 (... ..) ; KSK (**50322**)

com. 6400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK1 (**2623**)

com. 6400 IN **DNSKEY** 256 3 5 (... ..) ; ZSK2 (14672)

com. 6400 IN **RRSIG** **DNSKEY** 5 1 6400 (... **50322** ...)

www.com. 3600 IN **A** 1.2.3.4 ; A RR

www.com. 3600 IN **RRSIG** **A** 5 2 (... **2623** ...) ; Signs A RRset

Signature  
over all  
DNSKEYs  
using KSK  
**50322**.

Signature over type-A RRset using ZSK 2623

# DNSSEC: a DNS Public Key Infrastructure?

- DNSSEC is basically a DNS-specific Public Key Infrastructure (PKI)
  - No revocation mechanism (only expiration)
  - No transparency mechanism, no extensions...
  - But: naming constraints are built-in
- **Cipher agility in DNSsec ?**
  - DS records may use different hashing algs; all are sent to the client, who can validate using only some of them
    - Attacker cannot drop any since DS RRset is signed (RRSIG)
  - Offer multiple signatures using different algorithms, keys
- **Downgrade attack : attacker causing client (resolver) to use weaker algorithm than supported by server**
  - Leaving only weaker-security records / signatures
- **Question: is DNSSEC vulnerable to downgrades??**
- **x**

# DNSSEC Downgrade Prevention

- Downgrade attack 1: send only DNSKEYs using weak algs
- Prevention:
  - Domain must have DNSKEY for each algorithm indicated by the domain's DS RRset in the parent domain
    - With valid hash as per the DS RR
- Downgrade attack 2: send only RRSIGs using weak algs
- Prevention:
  - For every signing alg  $S$  in any DNSKEY record of the domain, there must be an RRSIG using alg  $S$  (with a key from DNSKEY)
  - ➔ If given only RRSIGs using weak alg, response is invalid !
  - Concern: req' applies to servers, not validators [RFC6840 sect.5.11]  
➔ several resolver implementations do not prevent such downgrades!

# DNSSEC Downgrade Prevention

- Downgrade attack 1: send only DNSKEYs using weak algs
- Downgrade attack 2: send only RRSIGs using weak algs
- Downgrade attack 3: don't send any DNSKEYs
  - As if domain is not using DNSSEC at all

Root Zone (KSK embedded in browsers, tag=3861)	
. 86400 IN DNSKEY 256 3 5 (... ..) ; ZSK1 (tag=1127)	Signature over both DNSKEYs using KSK 3861
. 86400 IN DNSKEY 256 3 5 (... ..) ; ZSK2	
. 86400 IN RRSIG DNSKEY 5 0 86400 (... 3861 ...)	
com. 86400 IN RRSIG DS 5 1 86400 (... 1127 ...)	Signature of... (using ZSK 1127)
com. 86400 IN DS 50322 5 1 (2BB18 ... A53B0A)	

com Zone	
www.com. 3600 IN A 1.2.3.4	; A RR
... (other non-DNSSEC records)	

- Prevention: ???



# DNSSEC Downgrade Prevention

- Downgrade attack 1: send only DNSKEYs using weak algs
- Downgrade attack 2: send only RRSIGs using weak algs
- Downgrade attack 3: don't send any DNSKEYs
- Downgrade attack 4: don't send the DS from parent domain, or don't send RRSIG records for a signed RR
  - As if domain/RR is not protected using DNSSEC
- Prevention: server must send proof of non-existence of DNSSEC (DS or RRSIG) records
- Related attack: send NXDOMAIN (non-existing domain) in response to query for an existing domain
  - E.g., blacklist, policy

# Secure DNS: proof of no (signed) RR

- What if bar.com has no public key?
  - Does not yet support Secure DNS
- Can send unsigned RRs...
- But: attacker may send unsigned RRs, even if bar.com does have public key!
- Maybe "com" would sign a 'NoSec RR', indicating bar.com has no public key?
- Any concerns?
  - *Efficiency - need to sign 'NoSec RR' for every unsigned domain*
  - *Worse - we have to sign in real time!*
    - *To prevent replays, and if subdomain doesn't exist*
- **NSEC RR Solution: "Next SECure record is <name>"**
  - Use lexicographic order, sign in advance

# The NSEC (Next SEcure) RR

- NSEC record: al.com NSEC don.com A RRSIG NSEC
  - Means: no RR alphabetically between al.com and don.com
  - Also: al.com has only A, RRSIG and NSEC records
- Sent in response to NSEC query for any record from al.com to don.com (alphabetically)
- Example: COM={al.com, don.com, ed.com, jon.com}
- Request is bob.com; response=?
  - Response is: al.com. 300 IN NSEC don.com A RRSIG NSEC
  - Authenticated (signed) using RRSIG
- Request is guy.com; response=?
  - Response is ed.com. 300 IN NSEC jon.com A RRSIG NSEC
- Request is sam.com; response=?
  - Response is jon.com. 300 IN NSEC al.com A RRSIG NSEC

# The NSEC RR – and Zone Enumeration

- NSEC stands for 'Next Secure'
- NSEC record: al.com NSEC don.com A RRSIG NSEC
  - Means: no RR alphabetically between al.com and don.com
  - Also: al.com has only A, RRSIG and NSEC records
- Sent in response to NSEC query for any record from al.com to don.com (alphabetically)
- Allows zone enumeration (discovery of subdomains)
- A potential security/privacy concern:
  - Directed attacks at subdomain (e.g., website)
  - Some name may identify vulnerability
    - E.g.: proxy.x.com → maybe open proxy??
  - Names may be sensitive, e.g., new company name
- Challenge: authenticate non-existing-domain response, but prevent zone enumeration!

# NSEC3 Hashed Non-Existence Response

- Authenticated non-existing RR response:  
`<hash1> NSEC3 <alg> <opt> <n> <salt> <hash2> <RRtypes>`  
e.g.: `F3...D.com NSEC3 1 0 9 4A...F 6C...8 A RRSIG NSEC`
- Hashed: goal is to prevent zone enumeration (but...)
  - Suppose  $h(al.com) = F3...D$  and  $h(don.com) = 6C...8$
  - Means: no RR  $x.com$  such that  $F3...D < h(x.com) < 6C...8$
  - Proving non-existence of RRs (e.g., of  $x.com$ )
  - Also:  $al.com$  has only `A RRSIG` and `NSEC` records
- Parameters: `<alg> <opt> <n> <salt>`
  - `<alg>`: hashing algorithm, mostly 1 for SHA-1
  - `<opt>`: if 0, NSEC3 done for all records in domain; if 1, NSEC3 only for protected records. (see later)
  - `<n>`: number of iterations (of hash function), here 9
  - `<salt>`: salt input to hash function, here `4A...F` [base64]

# NSEC3 Hashed Non-Existence Response

- Authenticated non-existing RR response:  
`<hash1> NSEC3 <alg> <opt> <n> <salt> <hash2> <RRtypes>`
- NSEC3 hash calculation:
  - Goal: (limited) defense against dictionary attacks
  - Iterated `n` times (to make computations slower [note: DoS?])
  - Salted: against precomputation (same salt for entire domain)
    - Actually, two fields: `salt` and `salt length` (0 to 255 bytes)
  - Hash of *name* with `salt` after *i* iterations:

$$H_{name}^{i,salt} = h(H_{name}^{i-1,salt} || salt || i - 1);$$

$$H_{name}^{0,salt} = h(name || salt)$$

- Compare this salt to salt of PW file

# Zone Enumeration with NSEC3 ?

- NSEC3 hash calculation:
  - Hash of *name* with *salt* after *i* iterations:
$$H_{name}^{i,salt} = h(H_{name}^{i-1,salt} || salt || i - 1); H_{name}^{0,salt} = h(name || salt)$$
- Each NSEC3 record contains next hash of name
  - To prove non-existence of names hashed in-between
- **Bernstein's attack:** collect hashed values (about one request per record), then offline dictionary attack!
- 'White lies' defense [RFC7129]: if *name* doesn't exist, send NSEC3 for  $H_{name}^{i,salt} - 1$ , with next-hash  $H_{name}^{i,salt} + 1$ .
  - Disadvantage: online signing (with zone's private key)
  - Adding random 'white lies' records in advance (offline), doesn't help much
- NSEC5 proposal: uses special private key (still, online)

# Zone Enumeration with NSEC3 ?

- NSEC3 hash calculation:
  - Hash of *name* with *salt* after *i* iterations:
$$H_{name}^{i,salt} = h(H_{name}^{i-1,salt} || salt || i - 1); H_{name}^{0,salt} = h(name || salt)$$
- Each NSEC3 record contains next hash of name
  - To prove non-existence of names hashed in-between
- **Bernstein's attack:** collect hashed values (about one request per record), then offline dictionary attack!
- Per-name salt defense:
  - Add 'salt' record type:  
name SALT <salt>
  - NOT signed; if query to non-existing name: return random salt
  - Does this prevent attack?




# The Opt-Out Flag <opt>

- Recall the <opt> Opt-Out flag of NSEC3:  
<hash1> NSEC3 <alg> <opt> <n> <salt> <hash2> <RRtypes>
- Opt-Out=0 (no Opt-Out): NSEC3 records done for all RRs in domain
- Opt-Out=1 (Opt-Out used): NSEC3 records done only for RRs of DNSSEC-protected subdomains
  - And <hash2> is of next DNSSEC-protected subdomain
- In many zones, most subdomains are unsigned
  - E.g., Top-Level Domains (TLDs), such as com.
- ➔ such domains often use Opt-Out=1
  - Less NSEC3 records ➔ less hashing and less signing
  - In fact, a common (main?) motivation to use NSEC3, not NSEC!

# DNSSEC: Status

- DNSSEC signed domains
  - Root + (most) TLDs are signed
  - Much less signing by `regular' domains



# DNSSEC: Status

---

- DNSSEC signed domains
  - Root + (most) TLDs are signed
  - Much less signing by `regular' domains
- DNSSEC validation
  - Concern: DNSSEC interoperability failures
    - And DNS response modification (mainly NX)
  - Invalid: resolvers respond with SERVFAIL
    - Causes client to try any alternate resolver known
  - Limited incentive: not visible to customer
    - Few/no clients makes validation mandatory
    - Some clients use trusted, secure resolver
    - Some applications use trusted secure resolvers (DoH, DoT)

# Domain Name System (DNS) Security

---

- DNS: quick recap
- DNS poisoning
  - Motivation: the hacker's swiss knife
  - Method 1 (historical): by Gratuitous `glue' RR
  - Method 2: send from corrupt NS
  - Method 3: send spoofed DNS response
  - Method 4: dangling DNS records
- DNSSEC: Cryptographic security for DNS
- DNS Privacy issues and defenses

# DNS Privacy Concerns [RFC7626,9076]

---

- Privacy? Isn't DNS data public?
- DNS leaks data about activities on Internet
- Most Internet activities involve one/few DNS queries
- Queries can expose private information:
  - What website user visits
  - And more: email sender/recipient domains, sw user runs...
  - Fingerprinting → identify user / device
    - Allow even identification of device behind NAT
  - Yet... DNS queries are sent in clear text
- By default, queries sent to DHCP-defined resolver
- ISPs & CDNs often add end-user data to queries
  - Goals: parental filtering and geo-optimized responses

# DNS Privacy Concerns [RFC7626,9076]

---

- Privacy? Isn't DNS data public?
- DNS leaks data about activities on Internet
- Most Internet activities involve one/few DNS queries
- Queries can expose private information
- **DNS poisoning can expose privacy, even if using https!**
  - Allows off-path attacker to be MitM to https traffic
  - Launch attacks against TLS that require MitM capabilities
    - Most TLS attacks assume MitM capabilities!
    - E.g., downgrade attacks, cryptanalytical attacks, ...
  - Allows TLS-traffic analysis and disconnection attacks
  - And: DNS has important non-web applications! (like what?)

# DNS Privacy Defenses

---

- Focus on stub-to-resolver communication, query data
  - Resolver-to-Authoritative is harder, less critical
  - Harder: prevent exposure from traffic analysis, side-channels
- DoT: DNS over TLS (RFC 7858,8310)
  - Resolver listens on port 853
  - TCP handshake, TLS handshake...
  - Use same connection for all request; pipeline for efficiency
  - Server authentication (public key validation):
    - Opportunistic: no server authentication
    - Out-of-band key-pinned
    - Several other options in RFC 8310 [deployment?]
- DoH: DNS over HTTPS (RFC 8484)

# DNS Privacy Defenses

- Focus on stub-to-resolver communication, query data
  - Resolver-to-Authoritative is harder, less critical
  - Harder: prevent exposure from traffic analysis, side-channels
- DoT: DNS over TLS (RFC 7858)
- DoH: DNS over HTTPS (RFC 8484)
  - Performs HTTPS request to predefined URI template
    - Domain, partial path, and extension format [RFC6570]
    - E.g.: <https://doh.opendns.com/dns-query>{?dns} (for FF)
  - A single DNS query per request, using GET or POST
    - With GET, {?dns} is replaced by the formatted DNS query
  - All requests use DNS-ID of 0 (zero), to allow caching
    - Freshness lifetime (Max-Age)  $\leq$  min TTL in response RRs
    - Client can still match request to response, since... ?



# DNS Security: Summary

---

- DNS is a key part of Internet infrastructure
- Often abused and attacked
- 'Plain' DNS is vulnerable
  - Trivial MitM attack, clever off-path attacks
- DNSSEC improves authentication
  - Deployed - but (too?) slowly
  - Security challenges remain (key agility allows downgrade?)
- Do we need privacy defenses, too?
  - DoH! (DNS over HTTP)
  - Or DoT, or (few other designs)
  - Limited: only client-to-resolver and only content
  - No protection against traffic analysis, side-channels