

---

University of Connecticut  
Computer Science and Engineering  
CSE 4402/5095: Network Security

# TCP/IP Security

©Amir Herzberg

# TCP/IP Security

---

- **Internet Protocol (IP) Security**
  - Fragmentation attacks
  - IPsec
- **Secure Transport**
  - TCP injections and other attacks
  - Quic
- **Conclusion**

# IP Security: Goals, Models, Expectations

<b>Attacker models</b> → <b>Security goals</b>	<b>MitM</b>	<b>Eavesdropper</b>	<b>Off-Path</b>
Confidentiality and privacy	None (without IPsec)	None (without IPsec)	Expected
Integrity and authentication	None (without Ipsec)	Trivial	Expected: spoofing, but no modification
Availability (and efficiency)	None	Trivial	Expected (except by clogging)

# The Internet Protocol: Fragmentation

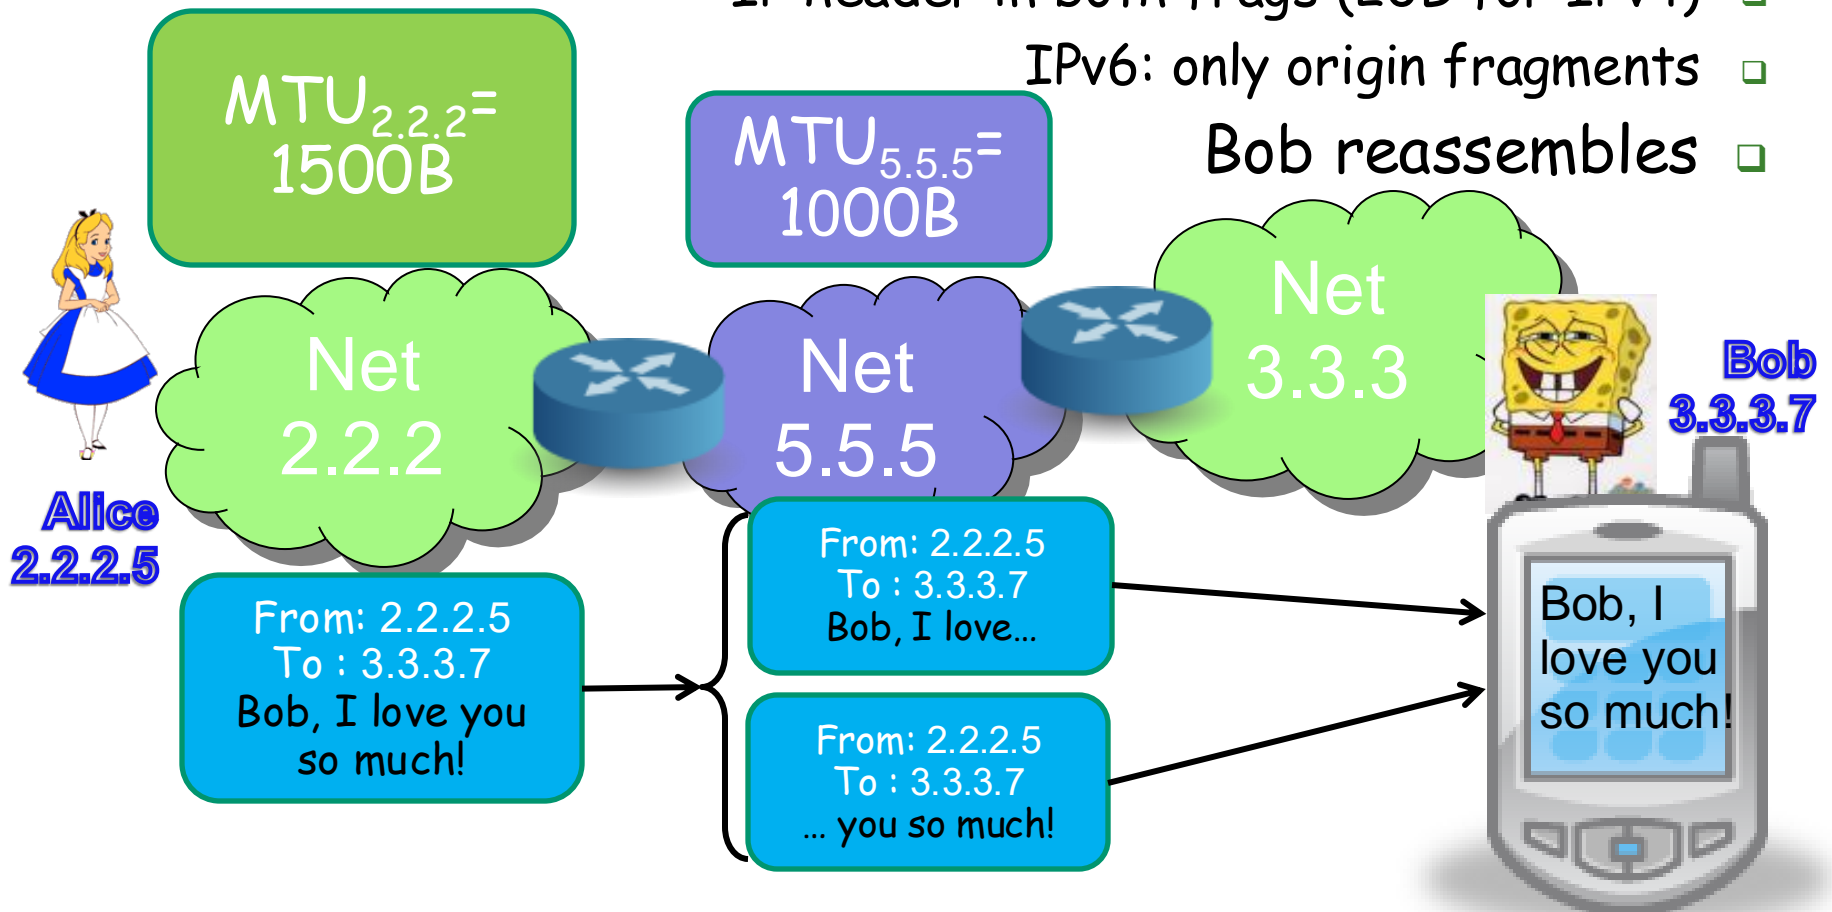
- Every network has a size-limit on packet size (MTU)
- What if we need to send more?

Solution: Fragmentation ■

IP header in both frags (20B for IPv4) □

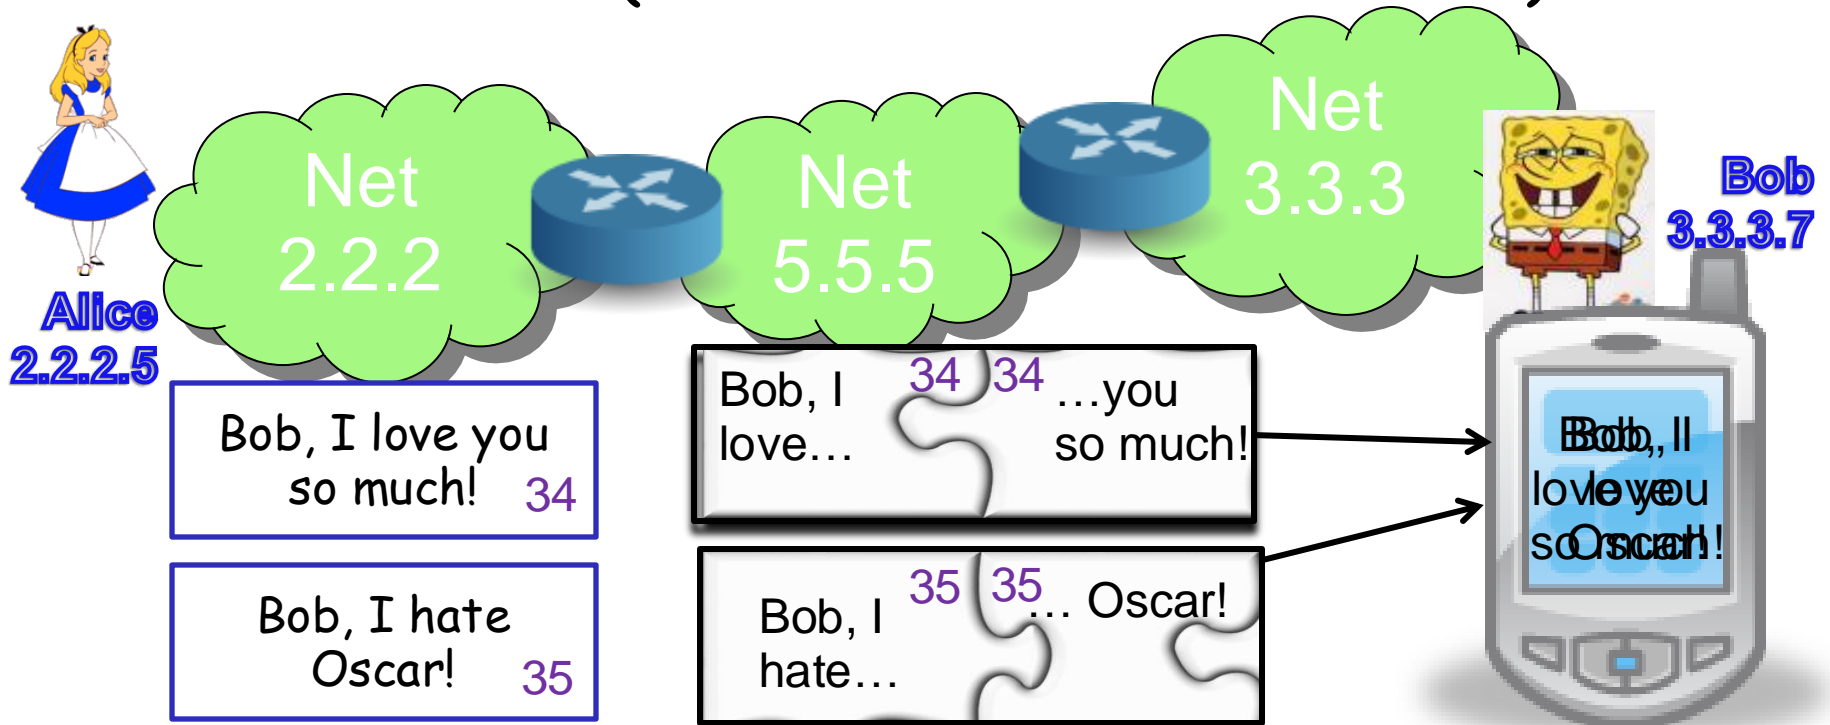
IPv6: only origin fragments □

Bob reassembles □



# Packet Reassembly: Careful(!)

- Bob receives fragments of multiple packets
- How to reassemble without mixing?
- Identify each packet
  - By Src, Dst addresses and protocol
  - And: IP-ID (16bit in IPv4; 32bit in IPv6)



# ‘Fragmentation considered Harmful’

---

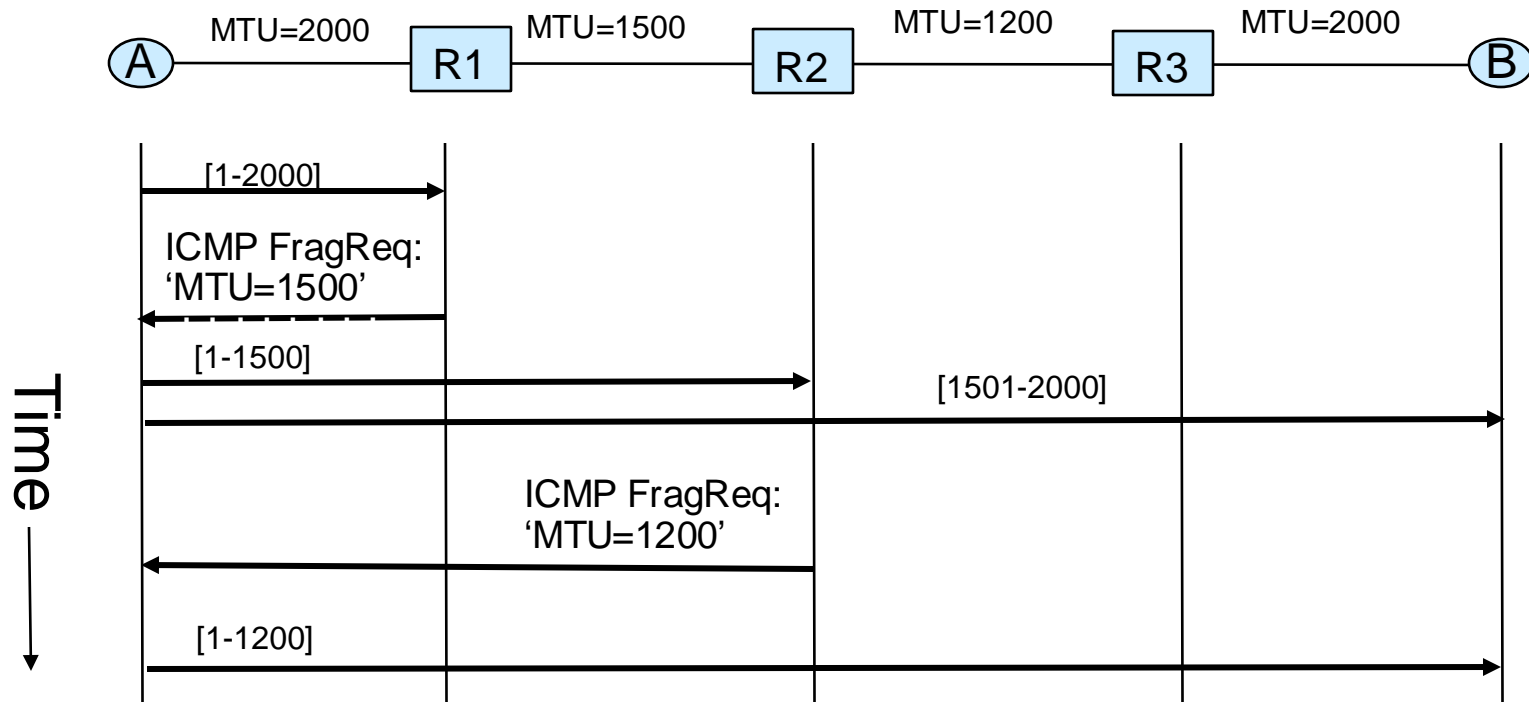
- IP fragmentation is conceptually easy, but...
  - Complexities: may arrive late or out of order
  - What if fragments overlap? How much storage?
  - Significant overhead
- ➔ Can we avoid fragmentation ? How?
- Always send only short packets
  - IPv4 allows short MTUs, but very rarely used
  - IPv6: MTU always at least 1200B
- Find minimal MTU on path
  - pMTUd: Path-MTU discovery
  - ➔ Limit packets to Path-MTU ➔ no frag!

# Avoiding Fragmentation

- Always send only short packets
  - DNS: originally limited to 512 bytes payload
  - But later, Extended-DNS allows longer packets
  - IPv6: all links have  $MTU > 1200B$ 
    - No fragmentation - except at source (UDP; why?)
- **Path MTU discovery (pMTUd)**
  - Send packets with Do not Fragment (DF) flag
  - Router sends back ICMP `Frag required` with MTU
  - Use path-MTU segments (TCP) or source-frag
- Works! few ( $< 0.5\%$ ) packets fragmented
  - Yet fragmentation is still used
    - Mainly: UDP, often with source-fragmentation

# Path MTU discovery (pMTUd)

- ❑ All packets sent with DF (don't fragment) bit set
- ❑ Periodically (rarely) try again to use larger MTU



Resending [1-1200] etc.:  
source-fragmented (UDP) or as new (shorter) packet (TCP)

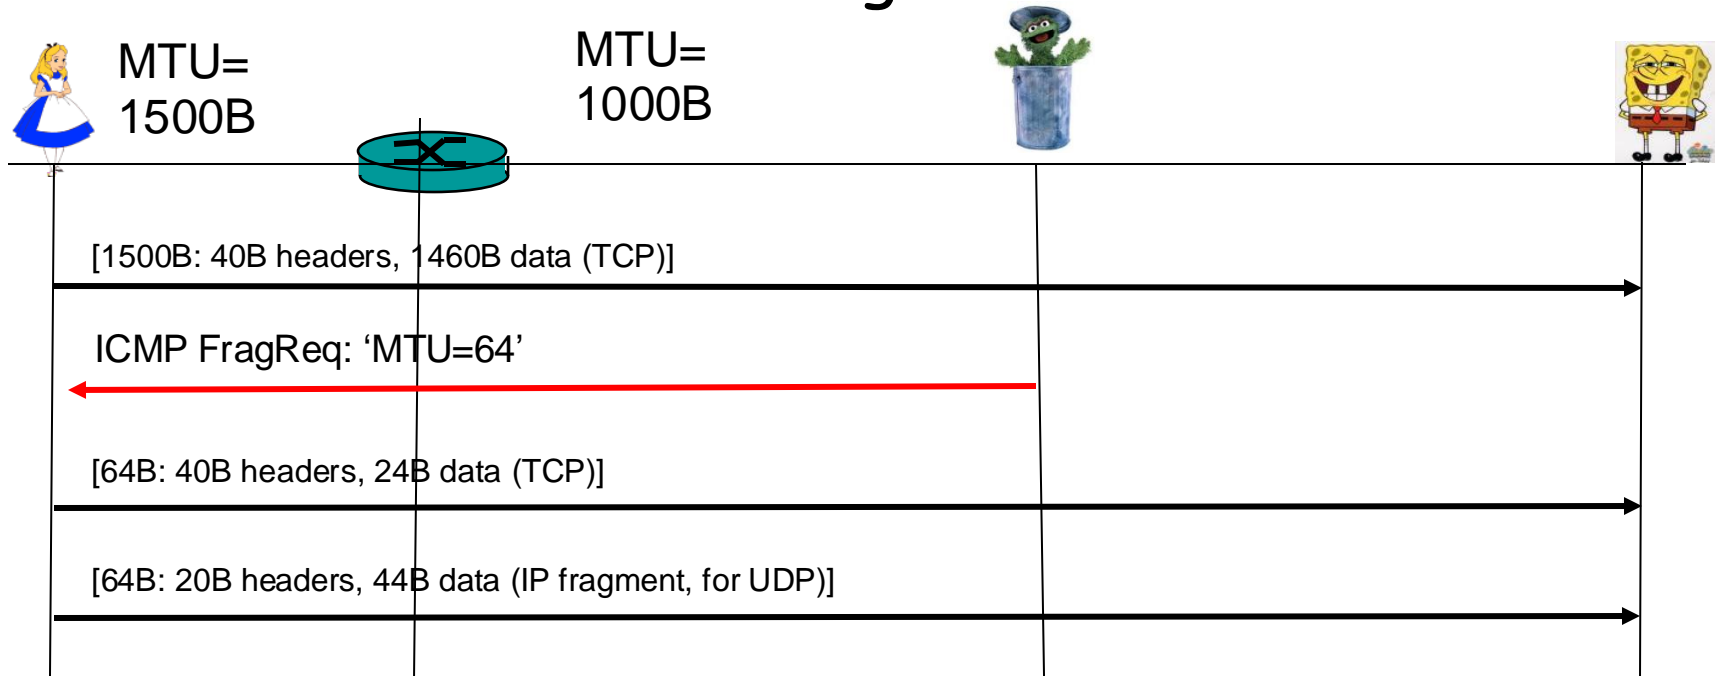


# Fragmentation attacks on IP

<b>Attacker models</b> → <b>Security goals</b>	<b>MitM</b>	<b>Eavesdropper</b>	<b>Off-Path attack Exploiting IP-ID</b>
Confidentiality and privacy	Trivially broken without crypto	Broken (without IPsec)	<b>2<sup>nd</sup> Frag interception attack</b>
Integrity and authentication	Trivially broken without crypto	Trivial	<b>2<sup>nd</sup> Frag spoofing attack</b>
Availability (and efficiency)	Trivially broken	Trivial	<b>Frag-based packet drop and overhead attacks</b>
Stealthy scan	Trivially broken	Trivial	<b>Off-path stealthy TCP scan</b>

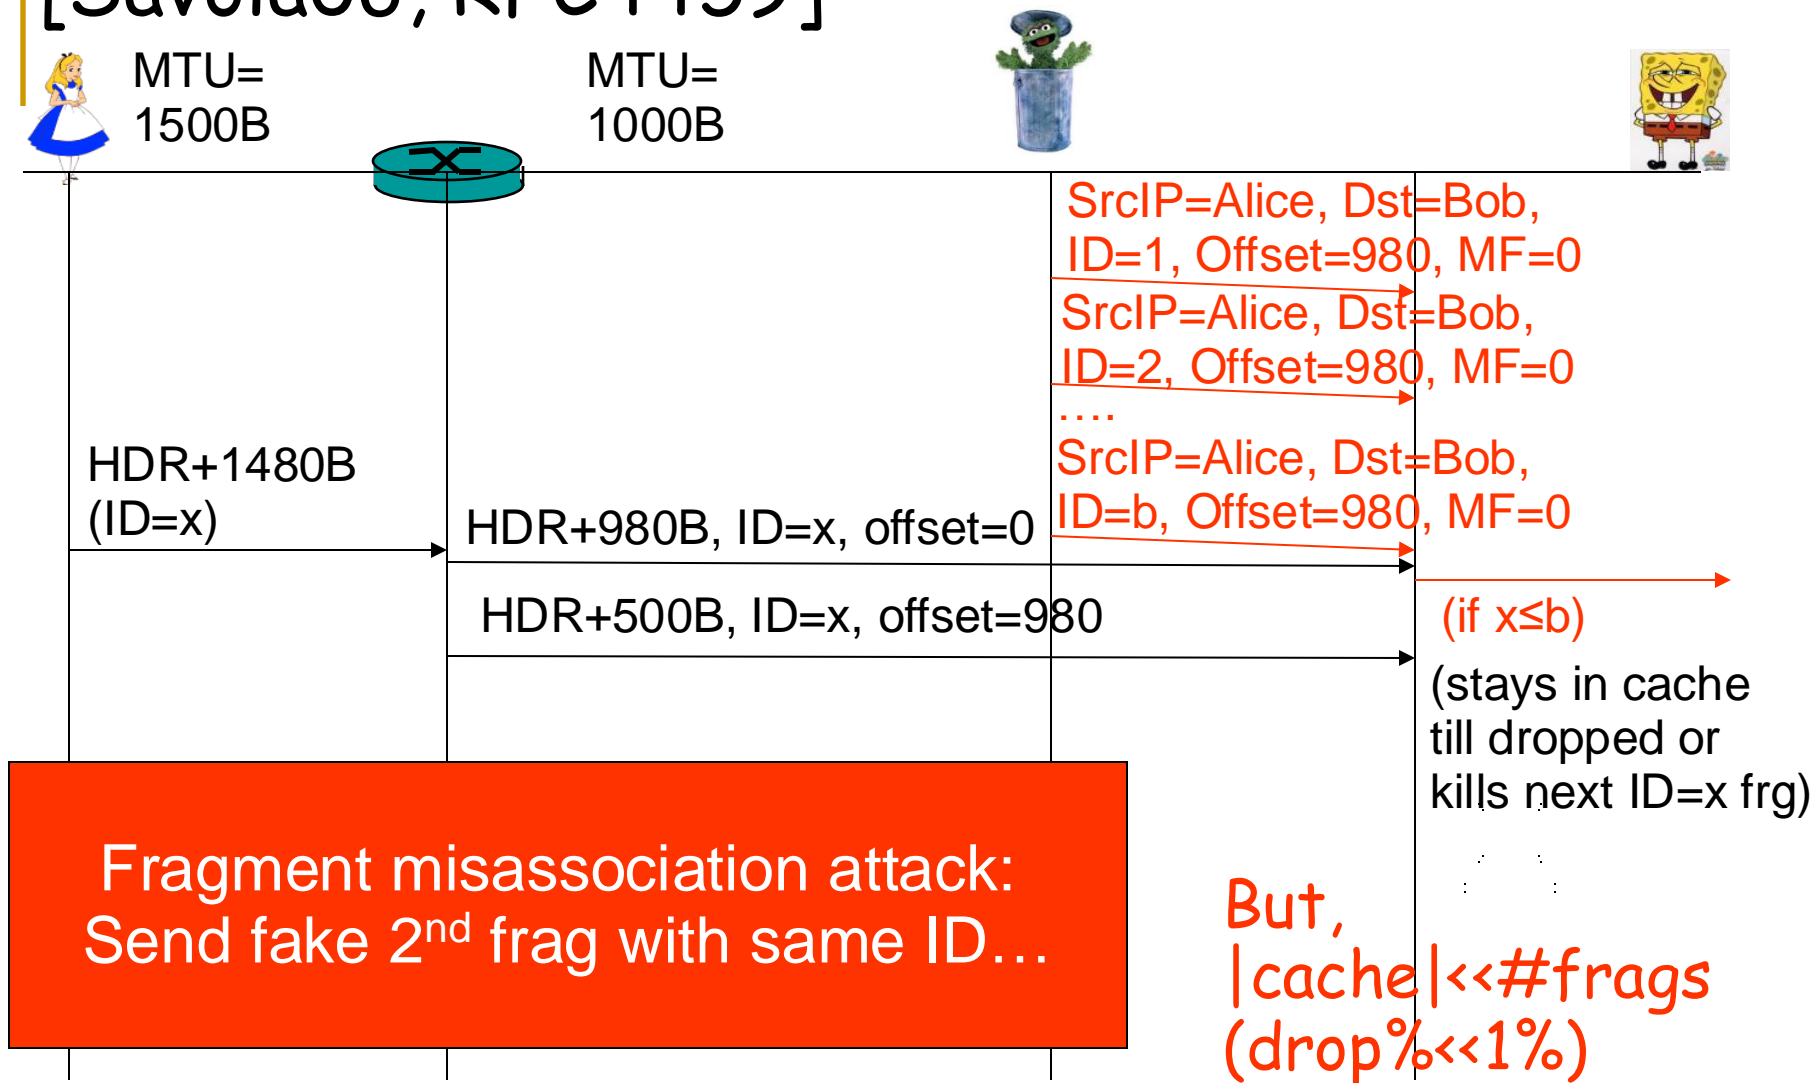
# Spoofed ICMP-Frag-Req Attack

- ❑ All links have same MTU (say, 1500B)
- ❑ Oscar sends ICMP Frag-Req to Alice
- ❑ Alice uses short MTU (TCP) or fragments at source (UDP - even in IPv6)
- ❑ So let's see these frag attacks...



# Fragment Misassociation Attack

[Savola06, RFC4459]

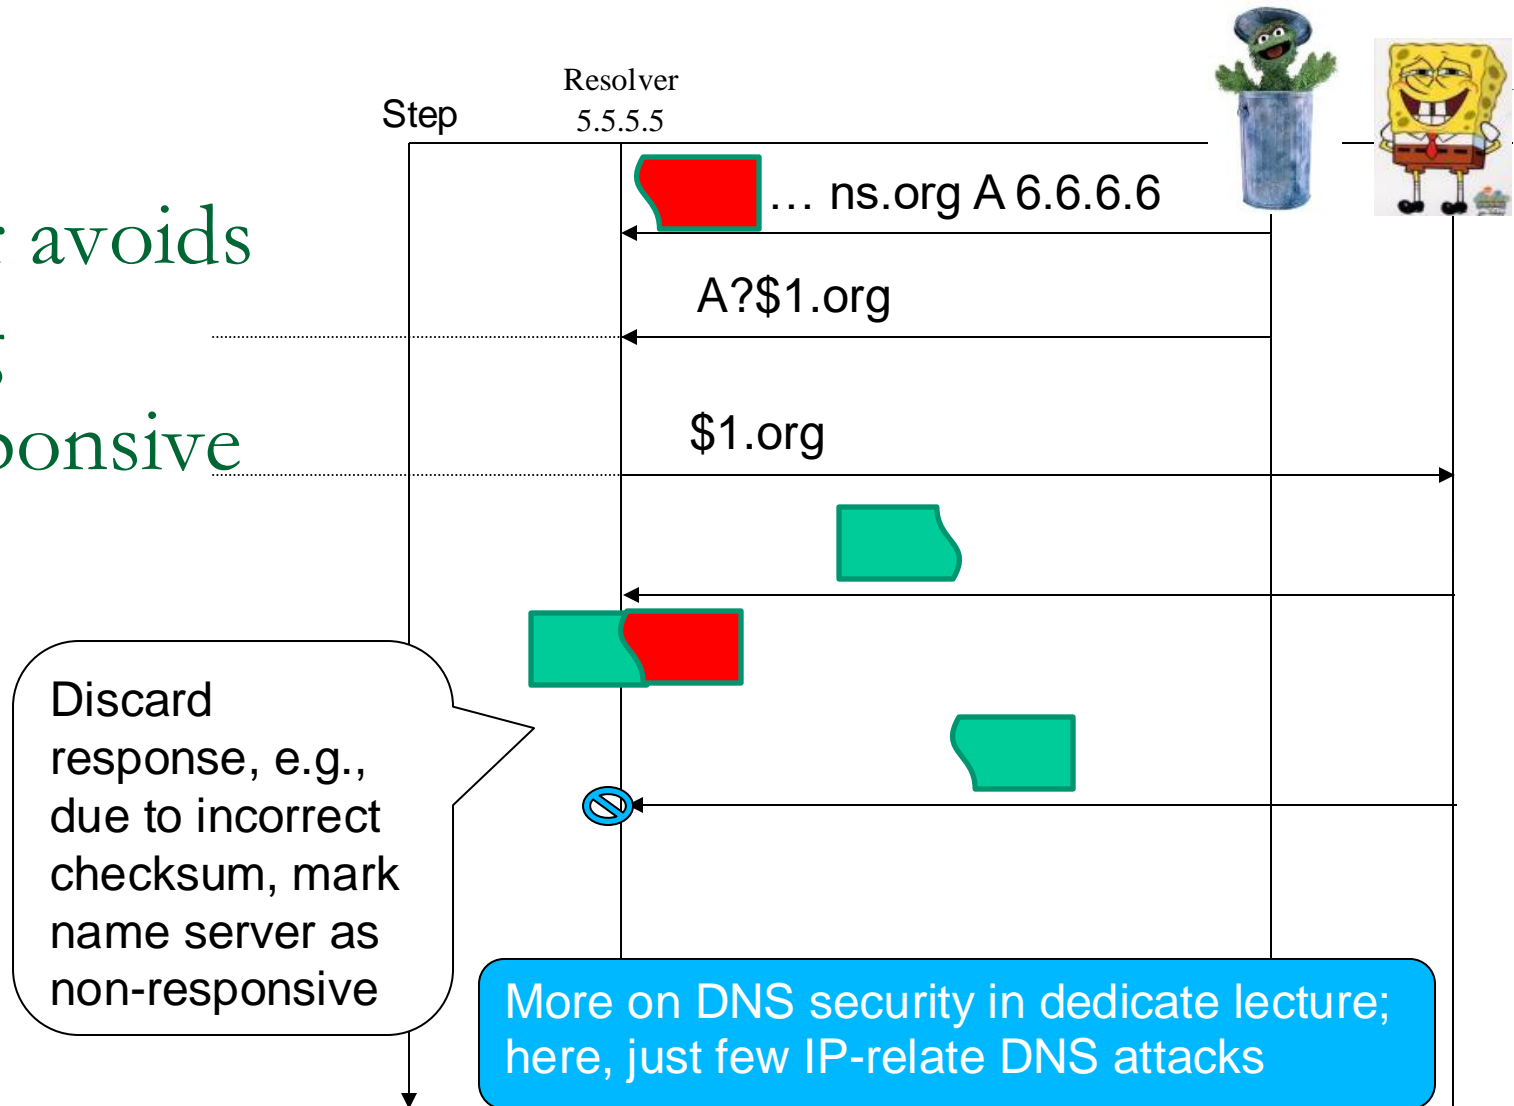


# What if Oscar can *find* the next IP-ID?

- Attacker can 'kill' packets (DoS)
- Assuming fragments are sent in-order:
  - Send 2<sup>nd</sup> fragment with next IP-ID
  - Cached by host (or GW, NAT)
  - When 1<sup>st</sup> fragment arrives (same srcIP, IP-ID):
    - Packet is reconstructed
    - UDP/TCP checksum error → packet discarded
    - Same as Savola's attack, but high success rates!
- Exercise: 'kill' packet when fragments are sent in reverse order (2<sup>nd</sup> fragments, then 1<sup>st</sup>)
- Application: block ('kill') DNS responses. Why?
  - So we have more time to poison
  - So resolver will use attacker-chosen name server: **transitive trust attack on DNS**


# DNS Response Blocking

Resolver avoids  
querying  
non-responsive  
name  
servers




# Can attacker *find* the IP-ID?

- How is the IP ID chosen (by the sender)?
  - Usually a counter – this is recommended in spec
    - To reduce probability of collisions (birthday paradox)
  - Three main approaches:
    - Per-destination incrementing (e.g., Linux)
    - Global incrementing (e.g., Windows, FreeBSD)
    - Mixed (Linux, from 4.1.8), see [1]
  - With all, it is often possible to find IP-ID
  - Can we find globally-inc
    - Send request to name server
    - Observe IP-ID in response
    - Measured at NS of ORG
    - And per-dest-inc IP-ID?



# Sometimes, finding per-dest IP-ID is Easy!

- Assume Oscar has an adversarial agent behind a NAT with the destination



- But why should Alice send a packet to the Zombie??
- Zombie may 'ping' Alice (in different ways), learn the current IP-ID Alice uses to send to NAT's IP from response
- Allows to intercept 2nd fragment by rewriting header(!)
- We'll return to this later...
- First: another way to find per-destination IP-ID

# Finding *per-destination* IP-ID

- ❑ Cause and detect loss to find per-dest IP-ID
- ❑ Example: find IP-ID of Name Server
  - Destination is open resolver
- ❑ Cause loss: send spoofed 2<sup>nd</sup> frags to resolver
- ❑ Detect loss: timeout of responses from resolver



6.6.6.6

OR (Open  
Resolver)

NS.foo.com  
(Name Server)

SrcIP=NS, DstIP=OR,  
**ID=i**, Offset=MTU, MF=0

SrcIP=6.6.6.6, DstIP=OR,  
**Req-j**: j.foo.com ANY, j=1,2,...,n

**Req-j**, j=1,2,...,n

Resp-j s.t. **ID<sub>j</sub>=i** discarded  
(wrong checksum)

**Resp-j s.t. ID<sub>j</sub>≠i**

**Resp-j**[1:MTU], **ID<sub>j</sub>**, MF=1,  
j.foo.com A 1.2.3.4,..., j=1,2,...,n

**Resp-j**[MTU:], **ID<sub>j</sub>**, MF=0,  
j.foo.com A 1.2.3.4,..., j=1,2,...,n



# Finding *per-destination* IP-ID

- ❑ Cause and detect loss to find per-destination IP-ID
- ❑ Example: find IP-ID of Name Server (sending to OR)



6.6.6.6

OR (Open  
Resolver)

NS.foo.com  
(Name Server)

SrcIP=NS, DstIP=OR,  
ID=i, Offset=MTU, MF=0

SrcIP=6.6.6.6, DstIP=OR,  
**Req-j**: j.foo.com ANY, j=1,2,...,n

**Req-j**, j=1,2,...,n

Resp-j s.t.  $ID_j = i$  discarded  
(wrong checksum)

**Resp-j**[1:MTU],  $ID_j$ , MF=1,  
j.foo.com A 1.2.3.4,..., j=1,2,...,n

Resp-j s.t.  $ID_j \neq i$

**Resp-j**[MTU:], MF=0,  
j.foo.com A 1.2.3.4,..., j=1,2,...,n

Detects timeout (loss) of Resp-j\*  
Concludes:  $ID_{j^*} = i$   
Hence: next IP-ID =  $i + (n - j) + 1$

Minor error, can you spot?

# Finding *per-destination* IP-ID

- ❑ Cause and detect loss to find per-destination IP-ID
- ❑ Example: find IP-ID of Name Server (sending to OR)



6.6.6.6

OR (Open  
Resolver)

NS.foo.com  
(Name Server)

SrcIP=NS, DstIP=OR,  
ID=i, Offset=MTU-20, MF=0

SrcIP=6.6.6.6, DstIP=OR,  
**Req-j**: j.foo.com ANY, j=1,2,...,n

**Req-j**, j=1,2,...,n

Resp-j s.t.  $ID_j = i$  discarded  
(wrong checksum)

**Resp-j**[1:MTU-20],  $ID_j$ , MF=1,  
j.foo.com A 1.2.3.4,..., j=1,2,...,n

Resp-j s.t.  $ID_j \neq i$

**Resp-j**[MTU-21:], MF=0,  
j.foo.com A 1.2.3.4,..., j=1,2,...,n

Detects timeout (loss) of Resp-j\*  
Concludes:  $ID_{j^*} = i$   
Hence: next IP-ID =  $i + (n - j) + 1$

## Challenges:

- ORs don't allow ANY query
- Ensuring fragmented response
- What is n? can it be smaller?

# Ensuring Fragmented Responses

---

- Method 1: Query to attacker's subdomains
  - For NS or .org, register subdomains:  
One-Domain-to-Rule-them-All.org
  - With many name servers with long names
  - ➔ Fragmented response
  - Or: apply attack to find IP-ID used by 'large' open resolver, responding to 'client OR'
- Method 2: Spoofed ICMP `fragmentation req' packet.

# Find *per-dest* IP-ID: battleship optimization

- As presented, we need  $n = 2^{16}$  queries to be sure we'll have Resp- $j$  such that  $ID_j = i$
- Goal: finding IP-ID of NS (to OR) with less queries
- Idea: send spread 'missiles' (like in Battleship game)



6.6.6

SrcIP= $i$   
ID= $i$ , Of

SrcIP= $i$   
Req- $j$ :  $j$



NS.foo.com  
(Name Serve

1,2,...,n

$D_j$ , MF=1,  
...,  $j=1,2,...,n$

, MF=0,  
...,  $j=1,2,...,n$

# Find *per-dest* IP-ID: battleship optimization

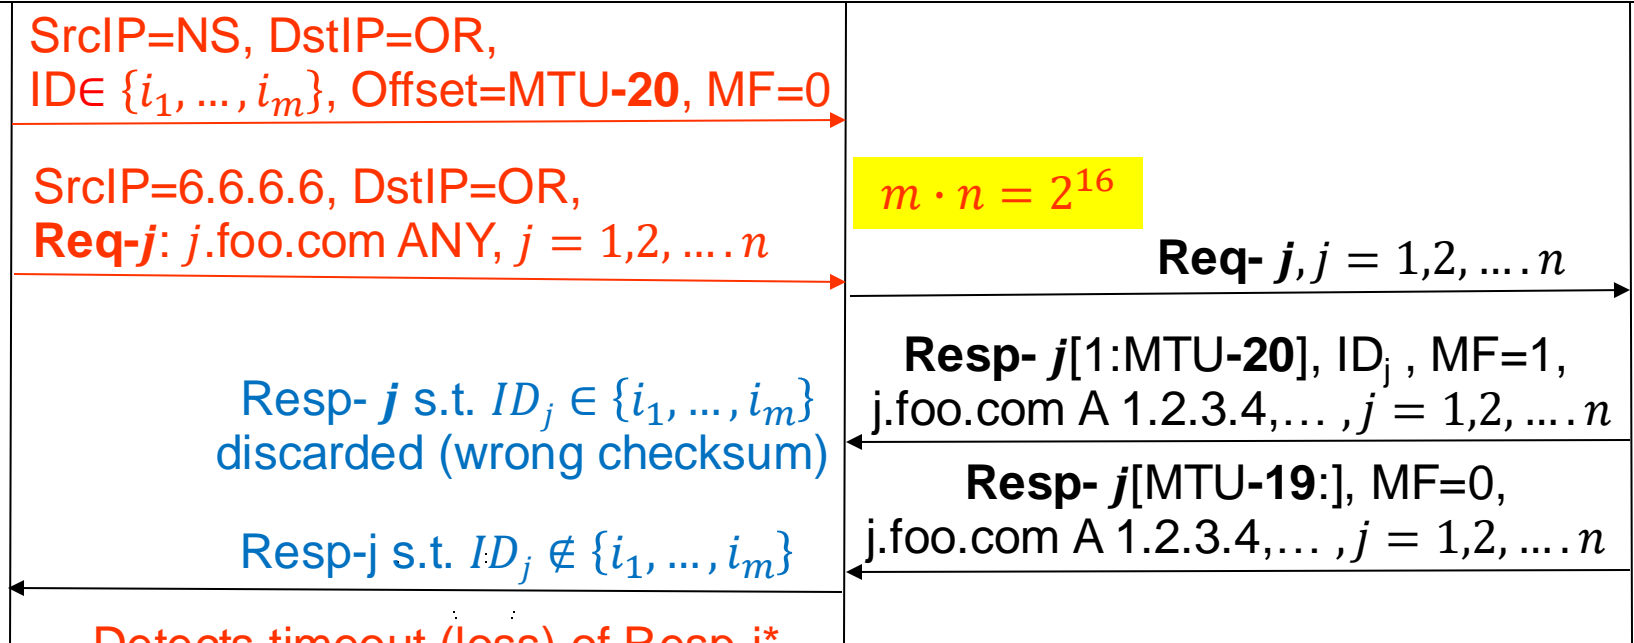
- Goal: finding IP-ID of NS (to OR) with less queries
- Idea: send multiple 'missiles' (like Battleship game)



6.6.6.6

OR (Open  
Resolver)

NS.foo.com  
(Name Server)



Detects timeout (loss) of Resp- $j^*$

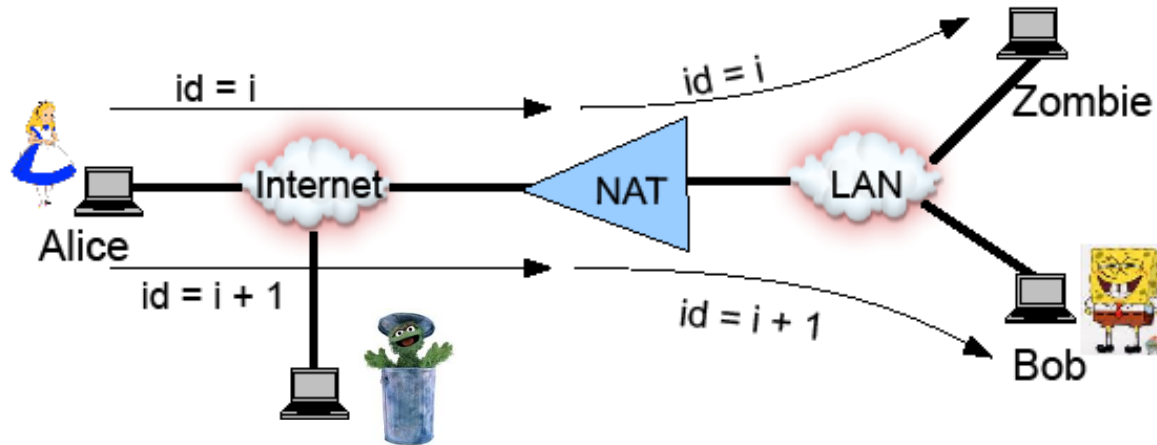
Concludes:  $ID_j \in \{i_1, \dots, i_m\}$ , hence: next IP-ID  $\in \{i_1 + n - j + 1, \dots\}$

How can attacker find exactly next IP-ID?

Hint:  $\log(m)$  'rounds' to divide-and-conquer the  $m$  possible values

# Recall: finding per-dest IP-ID behind NAT

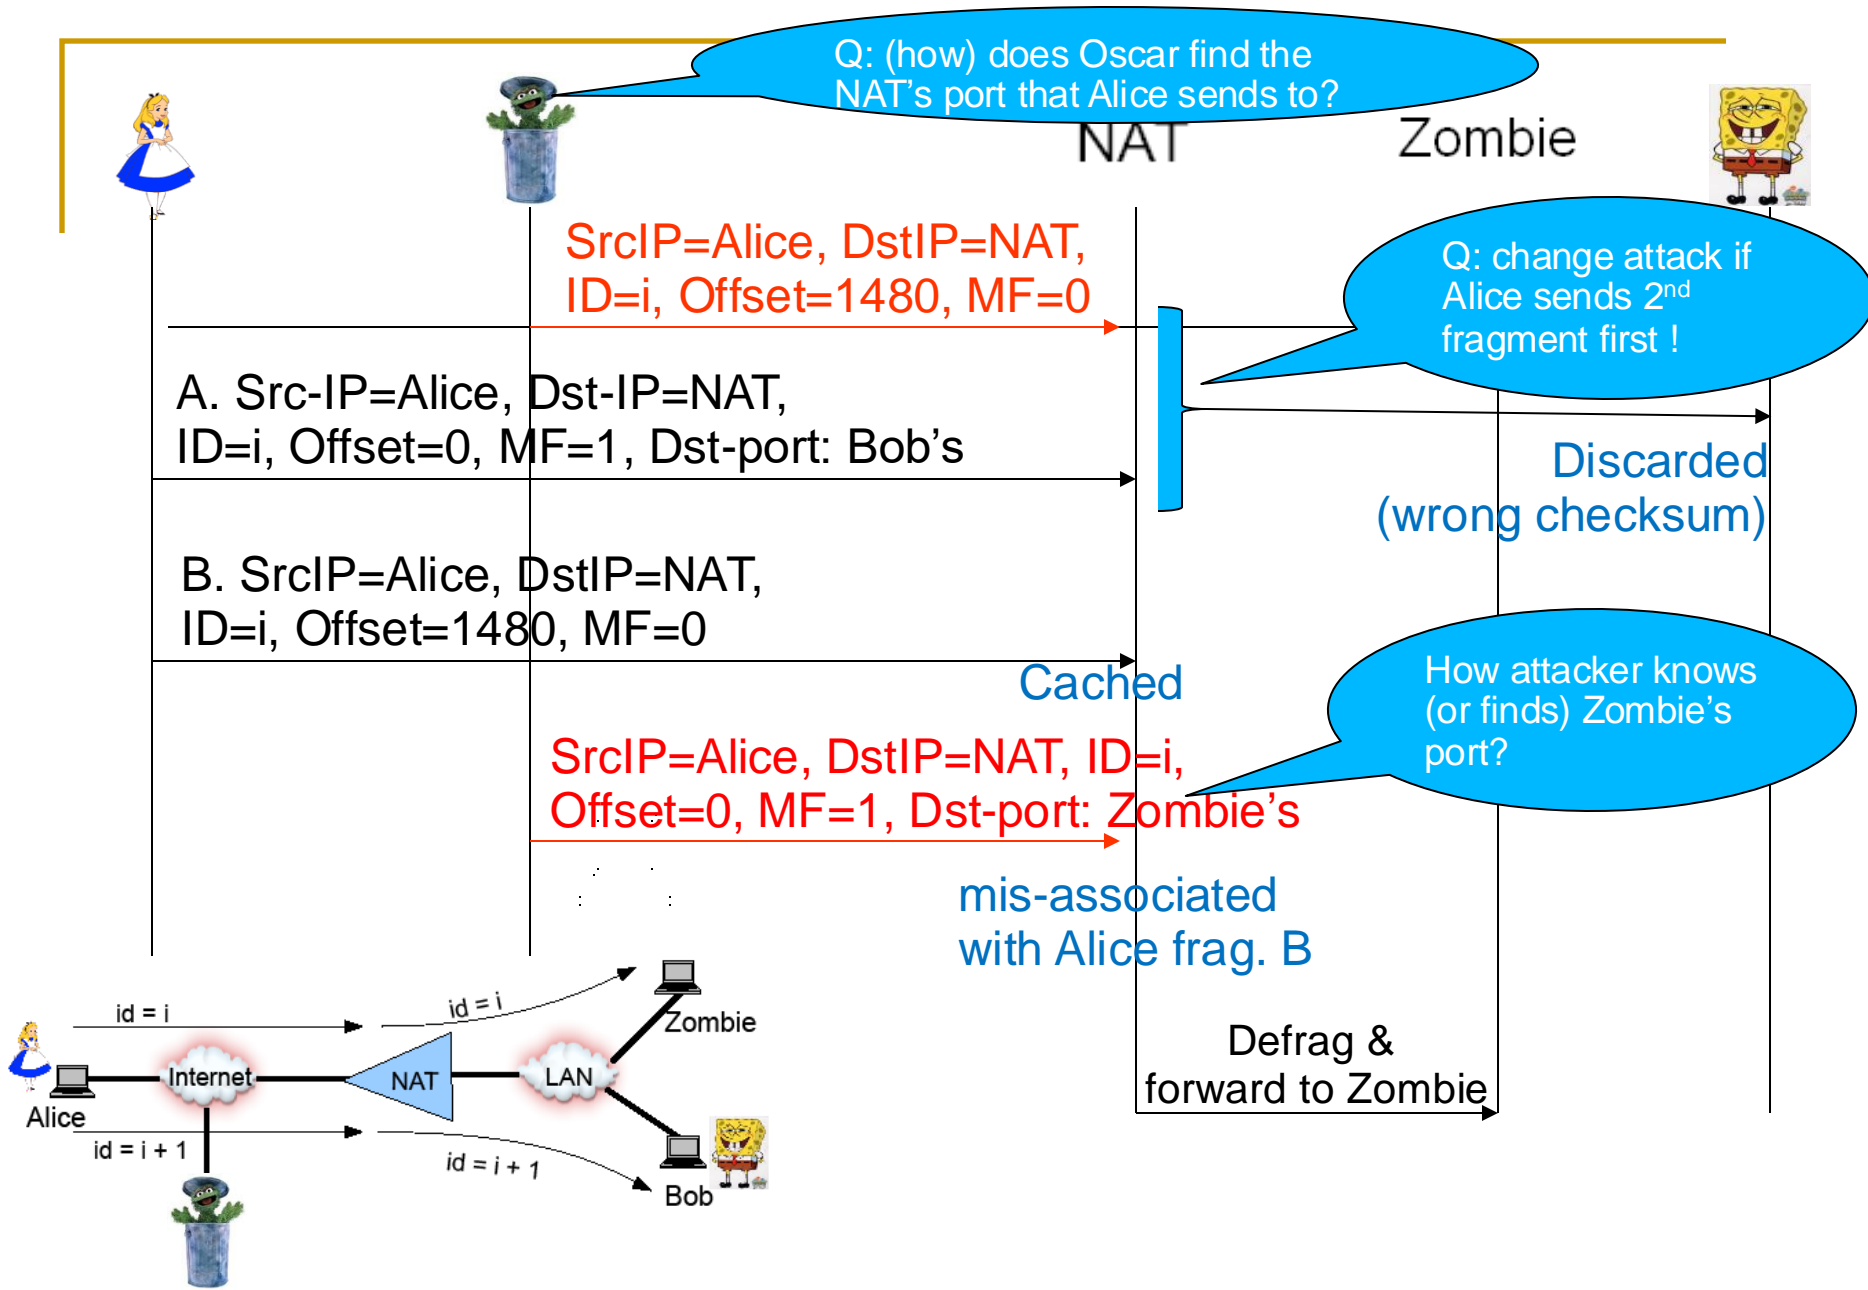
- Assume Oscar has an adversarial agent behind a NAT with the destination



# Fragmentation attacks on IP

Attacker models → Security goals	MitM	Eavesdropper	Off-Path
Confidentiality and privacy	None (without IPsec)	None (without IPsec)	1. 2 <sup>nd</sup> Frag interception attack
Integrity and authentication	None (without IPsec)	Trivial	2. 2 <sup>nd</sup> Frag spoofing attack modification
(ar	<ul style="list-style-type: none"> <li>- Next: 2<sup>nd</sup>-frag intercept / spoof attacks</li> <li>- Idea: replace 1<sup>st</sup> or 2<sup>nd</sup> fragment, respectively</li> <li>- Both: assume IP-ID known (and fragmentation)</li> </ul>		



# Off-path 2<sup>nd</sup> fragment intercepting





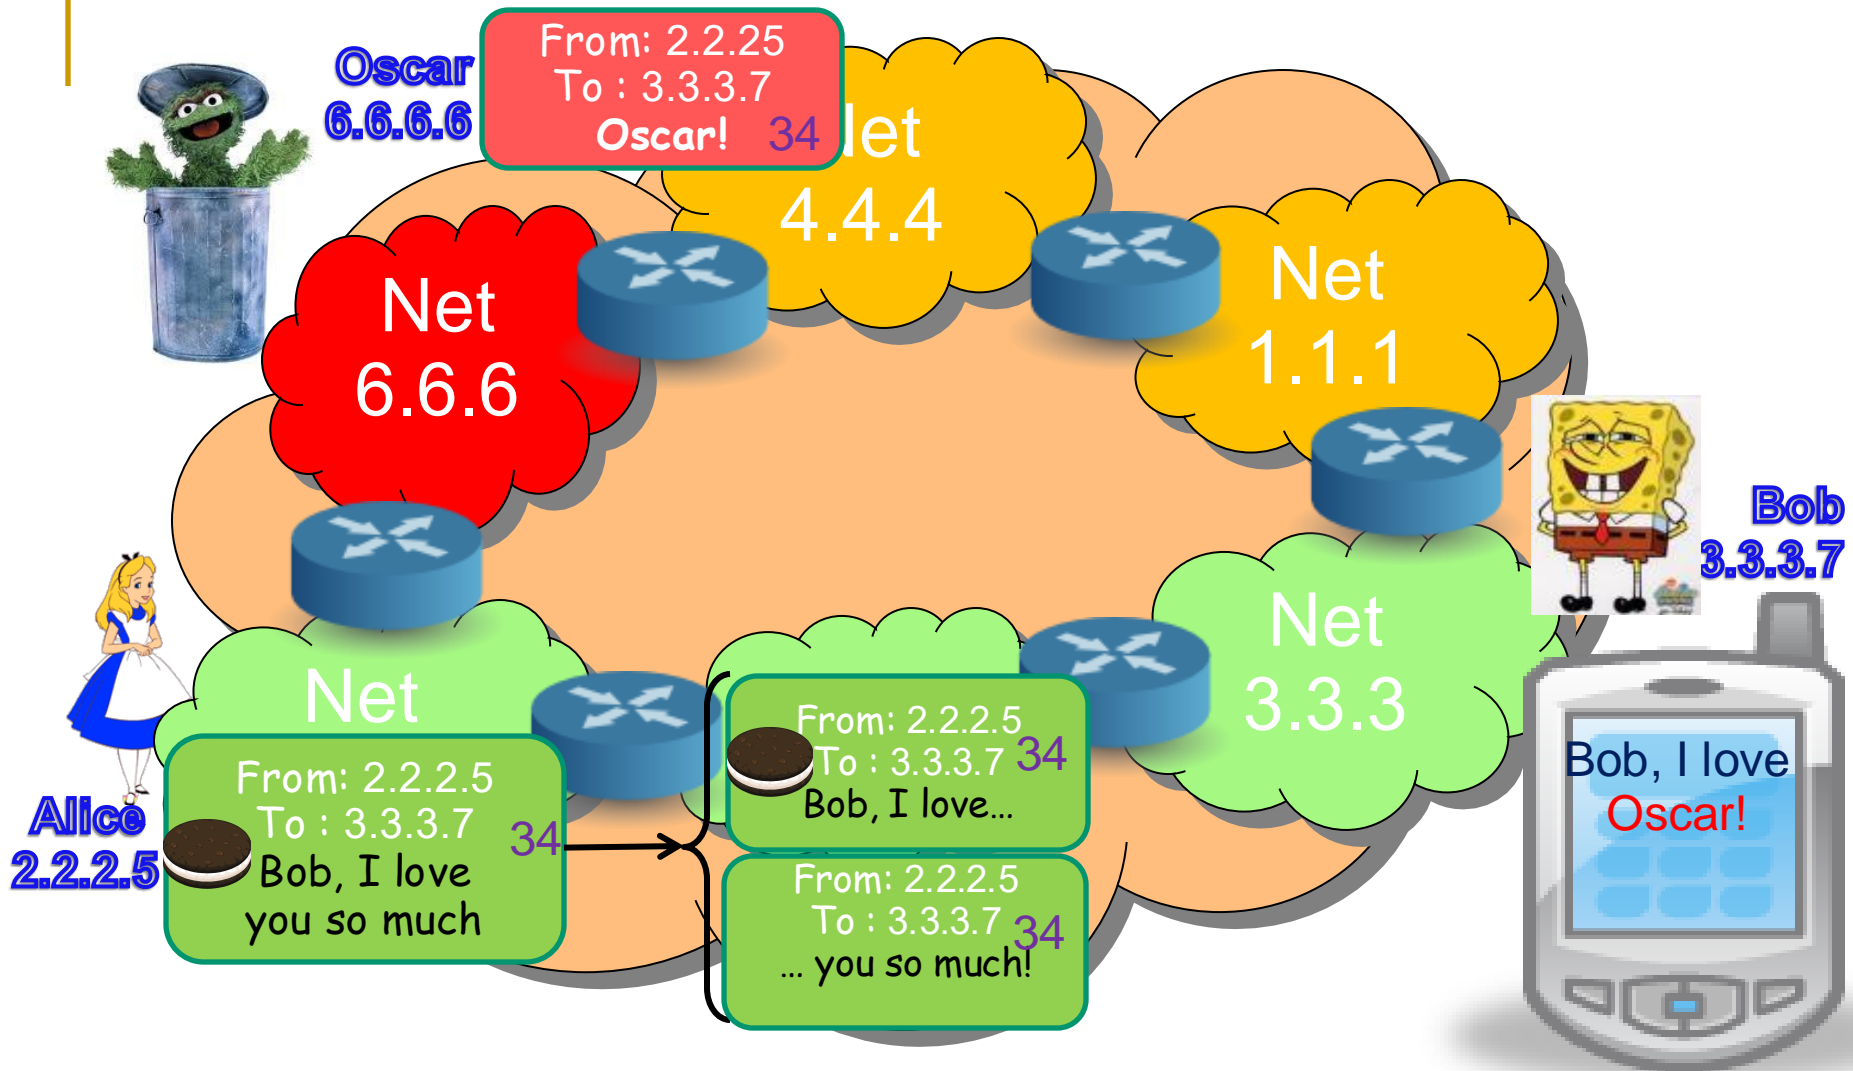
# 2<sup>nd</sup> Fragment Interception: Results

## ■ Results for IP tables based NAT



# Off-path 2<sup>nd</sup> fragment spoofing attack

(for *IP/DNS/TCP header cookie grabbing*)



# DNS Poisoning by spoofed 2<sup>nd</sup> fragment



Resolver  
5.5.5.5



NS.ORG


6.6.6.6

## Attack challenges:

1. Causing fragmented response
2. Predicting IP-ID (to spoof 2<sup>nd</sup> fragment)
3. Checksum: 2<sup>nd</sup> fragment must result in same checksum for packet

 ... ns.org A 6.6.6.6

A? \$1.org

A? \$1.org 

 ... \$1.org NXDOMAIN

A? vic.org

A? vic.org 

... vic.org A 6.6.6.6  
(message)

... vic.org A 6.6.6.6 

88423	199.249.120.1	IPv4	480	Fragmented IP protocol (proto=UDP 0x11, off=1480, ID=b063) [Reassembled in #207715]
207714	132.70.6.119	DNS	102	Standard query NS one-domain-to-rule-them-all.org
207715	199.249.120.1	DNS	1514	Standard query response
207716	199.249.120.1	IPv4	480	Fragmented IP protocol (proto=UDP 0x11, off=1480, ID=b063) [Reassembled in #207715]
▶ one-domain-to-rule-them-all.org: type NS, class IN, ns i23456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns j23456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns sns-pb.isc.org ▶ one-domain-to-rule-them-all.org: type NS, class IN, ns pdns3.ultradns.org ▶ h9p7u7tr2u9ld0v0ljs9llgidnp90u3h.org: type NSEC3, class IN ▶ h9p7u7tr2u9ld0v0ljs9llgidnp90u3h.org: type RRSIG, class IN ▶ o64vmqp2rn5ef3aou4g3hruir3ijhis4.org: type NSEC3, class IN ▶ o64vmqp2rn5ef3aou4g3hruir3ijhis4.org: type RRSIG, class IN				
▼ Additional records				
▶ a34353.123456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns i23456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns j23456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns sns-pb.isc.org ▶ b34353.123456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns pdns3.ultradns.org ▶ b34353.123456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NSEC3, class IN ▶ b34353.123456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type RRSIG, class IN ▶ b34353.123456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NSEC3, class IN ▶ a23456789101112131415161718192021222324252627282930313233343536.a234567891011121.one-domain-to-rule-them-all.org: type NS, class IN, ns i23456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns j23456789101112131415161718192021222324252627282930313233343536.123456789.one-domain-to-rule-them-all.org: type NS, class IN, ns sns-pb.isc.org ▶ c23456789101112131415161718192021222324252627282930313233343536.c234567891011121.one-domain-to-rule-them-all.org: type NS, class IN, ns pdns3.ultradns.org ▶ d23456789101112131415161718192021222324252627282930313233343536.d234567891011121.one-domain-to-rule-them-all.org: type NSEC3, class IN ▶ e23456789101112131415161718192021222324252627282930313233343536.e234567891011121.one-domain-to-rule-them-all.org: type RRSIG, class IN ▶ f23456789101112131415161718192021222324252627282930313233343536.f234567891011121.one-domain-to-rule-them-all.org: type NSEC3, class IN ▶ g23456789101112131415161718192021222324252627282930313233343536.g234567891011121.one-domain-to-rule-them-all.org: type RRSIG, class IN ▶ h23456789101112131415161718192021222324252627282930313233343536.h234567891011121.one-domain-to-rule-them-all.org: type NSEC3, class IN ▶ i23456789101112131415161718192021222324252627282930313233343536.i234567891011121.one-domain-to-rule-them-all.org: type RRSIG, class IN ▶ j23456789101112131415161718192021222324252627282930313233343536.j234567891011121.one-domain-to-rule-them-all.org: type NSEC3, class IN ▶ sns-pb.isc.org: type A, class IN, addr 132.70.6.244 ▶ pdns3.ultradns.org: type A, class IN, addr 132.70.6.202				

DNS query sent by resolver

Spoofed second fragment

DNS response: First authentic fragment reassembled with spoofed second fragment

Authentic second fragment (discarded after timeout)

0020	ad c4 72 60 e0 ed fd
0030	8f 85 9f 7f cb 7a b8
0040	a5 28 7e 29 a9 08 9f
0050	d1 92 86 22 4e 13 ca
0060	80 00 04 84 46 06 c8
0070	80 00 04 84 46 06 c8
0080	80 00 04 84 46 06 c9
0090	80 00 04 84 46 06 ca
00a0	80 00 04 84 46 06 f4
00b0	80 00 04 84 46 06 ca
00c0	80 00 04 84 46 06 ca
00d0	80 00 04 84 46 06 f4
00e0	80 00 04 84 46 06 77
00f0	80 00 04 84 46 06 f4
0100	80 00 04 84 46 06 f4
0110	80 00 04 84 46 06 f4
0120	80 00 04 84 46 06 ca
0130	80 00 10 20 01 0d b8
0140	70 73 34 c2 eb 00 1c
0150	01 0d b8 85 a3 00 42
0160	21 00 01 00 01 00 01
0170	00 00 00 00 00 00 00

# TCP/IP Security


---

- **Internet Protocol (IP) Security**
  - Fragmentation attacks
  - IPsec
- **Secure Transport**
  - TCP injections and other attacks
  - Quic
- **Conclusion**



# Secure Virtual Private Networks

- **Private Network:** owned by single organization
- **Secure VPN:** secure networking over insecure Net
  - ❑ Prevent eavesdropping, spoofing, injecting, modifying
  - ❑ MitM attacker
  - ❑ Main (standard) tool: **IP-Sec**



# Fragmentation attacks on IP

Attacker models → Security goals	MitM	Eavesdropper	Off-Path
Confidentiality and privacy	None (without IPsec)	None (without IPsec)	2 <sup>nd</sup> Frag interception attack
Integrity and authentication	None (without IPsec)	Trivial	2 <sup>nd</sup> Frag spoofing attack modification
Availability (and efficiency)	None	Trivial	Frag-based packet drop and overhead attacks

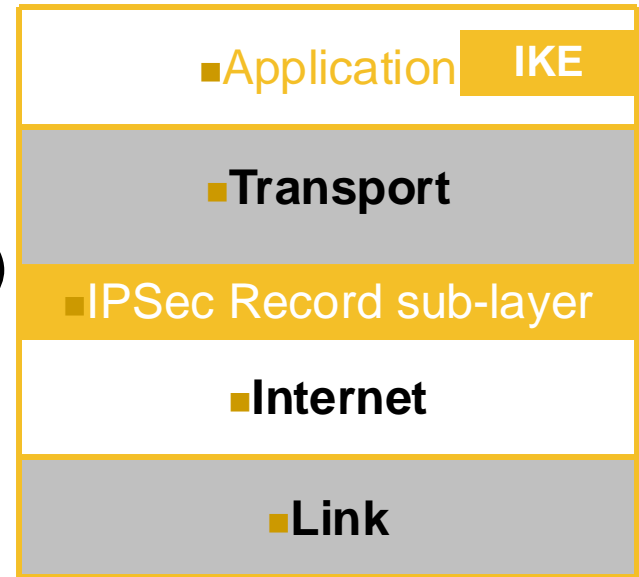
# IP Security: Impact of IPsec

Attacker models → Security goals	MitM	Eavesdropper	Off-Path
Confidentiality and privacy	Fixed by IPsec ICMP-redirect and feedback attacks on encryption-only ESP	None (without IPsec)	Expected
Integrity and authentication	Fixed by IPsec Modification attack on encryption-only ESP	Trivial	Expected: spoofing, but no modification
Availability (and efficiency)	None Stealthy, efficient attacks in spite of IPsec	Trivial	Frag-based packet drop and overhead attack (even for TCP due to tunnel?)



# IP-Sec Layers

- Two separate layers
- **IKE – Internet Key Exchange**
  - ❑ Establish shared key (application layer)
- **IP-Sec Record Sub-Layer: traffic encapsulation & protection**
  - ❑ **ESP protocol – Encapsulating Security Payload**
  - ❑ **AH protocol – Authentication Header** (only authentication, no encapsulation)
  - ❑ Signal to IKE when detecting traffic that requires IP-Sec but without established IP-Sec connection



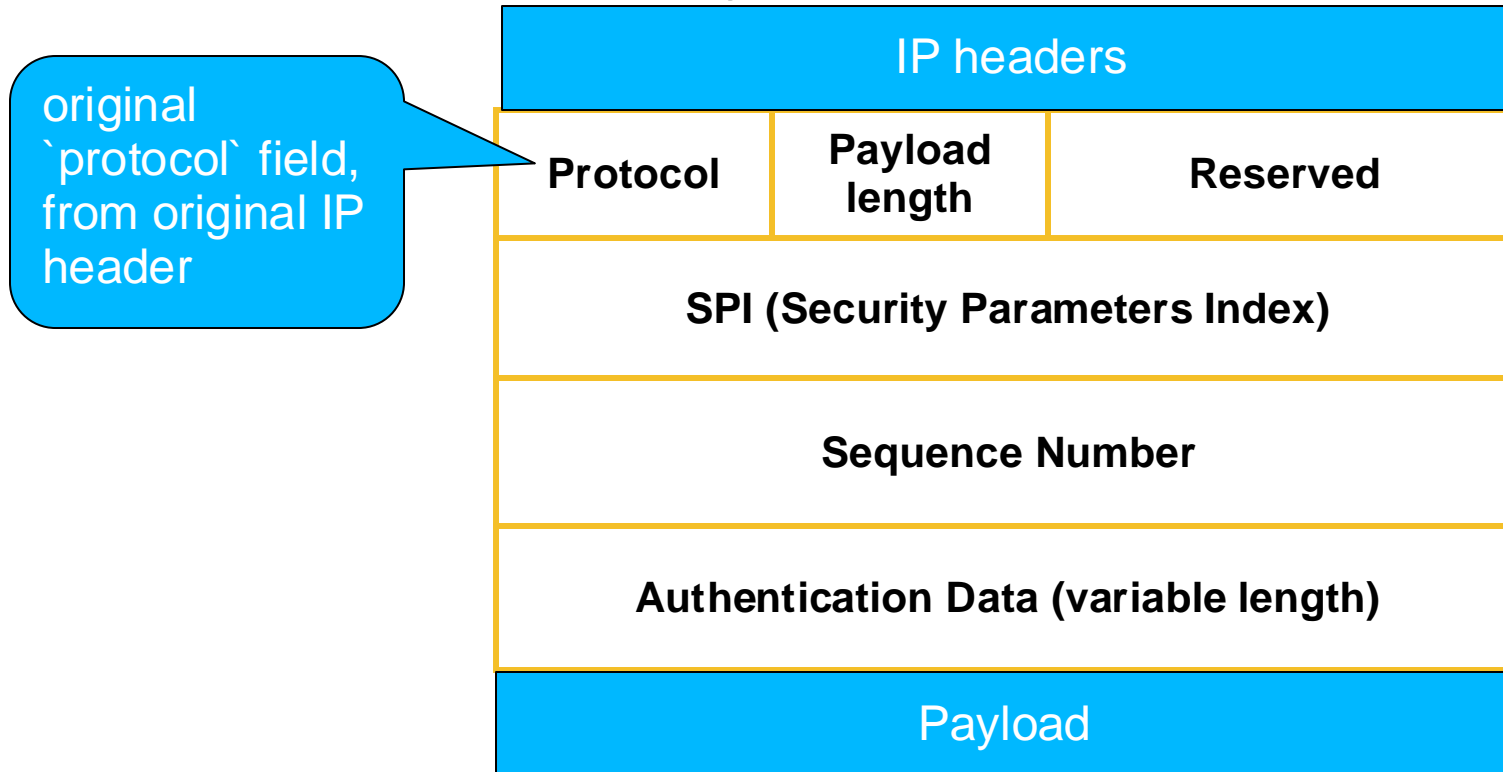
# IP-Sec Record Sub-Layer

---

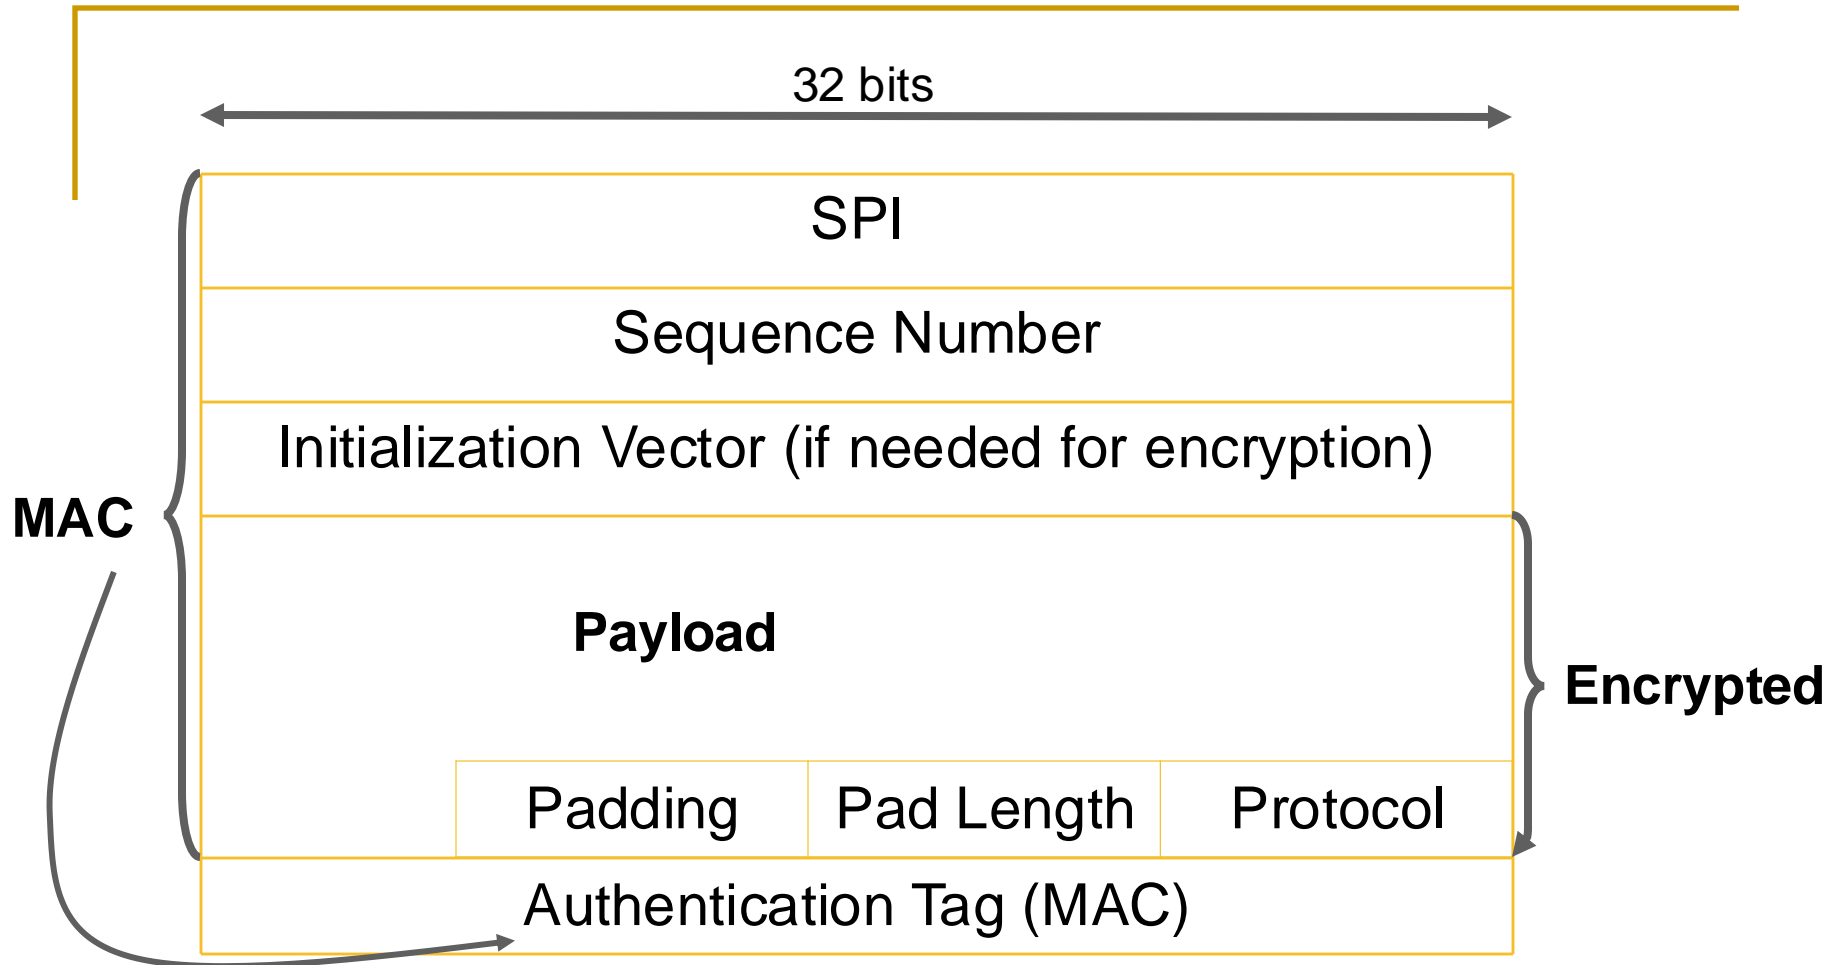
- Consists of two header protocols
  - ❑ Sent as protocols over IP
  - ❑ **ESP: Encapsulating Security Payload** (IP protocol 50): supports encryption and/or authentication (MAC)
  - ❑ **AH: Authentication Header** (IP protocol ID 51): only MAC
- Both AH and ESP:
  - ❑ Add sequence numbers to packets (to prevent duplication)
  - ❑ `Synchronized` keys, algorithms and counters
    - **Security Association (SA)**
    - Identify SA for packet via its **Security Parameter Index (SPI)**

# AH - Authentication Header [RFC4302]

- Authenticates entire IP packet: IP header, AH header, and payload
  - ❑ Except router-modified fields in IP-header, e.g. TTL (hop count)
  - ❑ UDP/TCP payload begins with ports, contains checksum



# ESP – Encapsulating Security Payload



# *Encryption-Only ESP can be Vulnerable!*

---

- Authentication is optional in ESP
  - Advised, but not always done
  - Q: does the (encrypted) checksum ensure integrity?
- No! ESP without authentication has no integrity
  - How can attacker change without the key?
  - Easier to see assuming stream-cipher encryption (e.g. OTP)...

# *Encryption-Only ESP can be Vulnerable!*

- Authentication is optional in ESP
  - Advised, but not always done (often not even default)
- ESP without authentication may not provide integrity
  - In spite of TCP/UDP checksum, etc.
  - Easier to see for stream-cipher encryption (OTP, OFB)
    - Some changes are easy also for CBC, CTR, CFB
- Confidentiality is vulnerable too!
  - Changed ciphertext attacks, exploiting feedback (side channel)
    - [Bellovin 96] TCP silently discards wrong checksum (no ack)
    - [Vaudenay 02, Paterson 06] response to incorrect padding
  - Modify header to ICMP echo – redirect to attacker – expose contents!
  - ➔ IPsec should always also authenticate
  - Attacks may depend on mode of operation of cipher, IPSec

# IP-Sec: Transport vs. Tunnel Modes

---

## ■ Transport Mode (end to end only)


- ❑ Change protocol field to AH (51) or ESP (50)
  - ❑ Protocol field kept in (& restored from) IPsec header
- ❑ Except protocol field: use existing IP header
- ❑ End-to-end – encapsulation by source host, decapsulation by destination host (receiver)

## ■ Tunnel Mode (gateway to gateway or end to end)

- ❑ Entire original IP packet (including header) is payload
- ❑ IP-Sec puts a new IP header (protocol AH (51) / ESP (50))
- ❑ Allows **Secure Virtual Private Network (VPN)**

# IP-ID Exposing Attacks on Tunnel


- Main difference from NAT scenario:
  - Packets 'on the Internet' have a different IP header
    - Zombie, cannot see the 'Internet IP-ID'
- Improved motivation: fragmentation is common in tunnels, due to extra header
  - Or caused by spoofed ICMP frag-required - possibly even for TCP, since embedded in ESP in tunnel





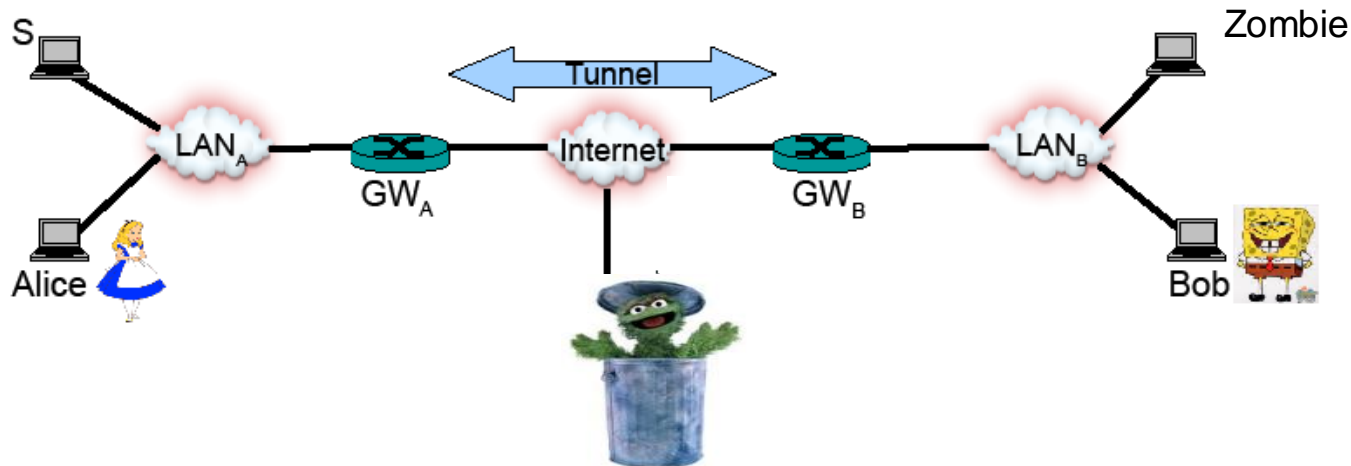
# IP-ID Exposing Attack

- Use **packet loss as a side channel** to identify the current IP-ID between two hosts
  - Example: btw two gateways (end of `tunnel`)
  - Similarly: btw (DNS) proxy and server
- Simplifying assumption: no benign traffic or packet loss

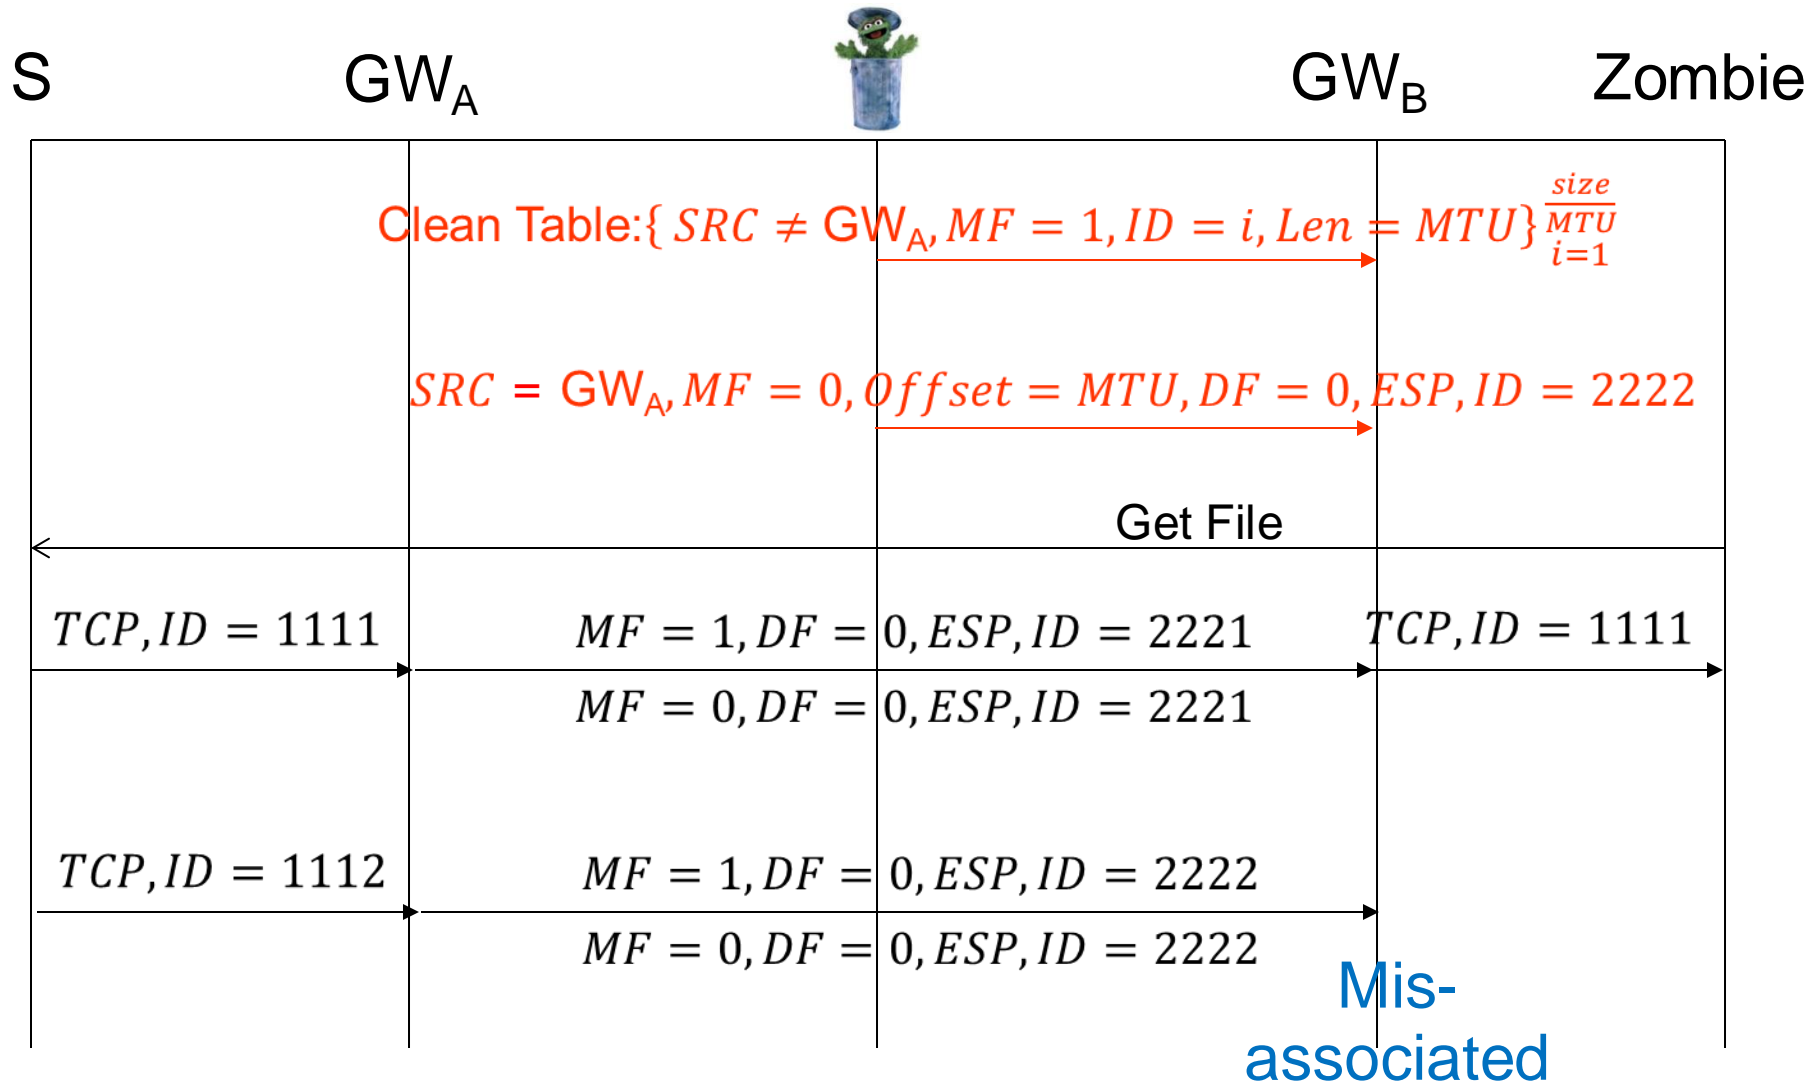


# IP-ID Exposing: review

- Goal: find IP-ID of next pkt from  $GW_A$  to  $GW_B$ 
  - Most efficiently (minimal # of packets)
- Assume:
  - Zombie can 'get' packets from  $S$  (e.g., files)
  - Packets encapsulated btw gateways, and fragmented
  - Incremental IP-ID
  - No other (benign) traffic, no packet loss



# ID Exposing Attack: tunnel scenario



# IP-ID Exposing Attack

S

GW<sub>A</sub>




GW<sub>B</sub>

Zombie

TCP, ID = 1113	MF = 1, DF = 0, ESP, ID = 2223	TCP, ID = 1113
▪	MF = 0, DF = 0, ESP, ID = 2223	▪
⋮	▪	⋮
▪	▪	▪
TCP, ID = 1110	MF = 1, DF = 0, ESP, ID = 2220	TCP, ID = 1110
	MF = 0, DF = 0, ESP, ID = 2220	
<p style="color: red;">Feedback: <math>1110 - 1112 = -2 \pmod{2^{16}}</math></p> <p style="color: red;">←</p> <p style="color: red;">compute:</p> <p style="color: red;">next = <math>2222 - 2 + 1 = 2221 \pmod{2^{16}}</math></p>		


# ID Exposing Attack - Meet in the Middle

- But... if  $n$  is the number of possible identifiers, this attack requires to send  $O(n)$  packets.
  - $2^{16}$  for IPv4, for  $2^{32}$  IPv6
- More efficient: use Battleship optimization




# Exploit: Continual Deny & Expose

- Off-path attacker has the current IP-ID
  - Goal: deny fragmented traffic
- Difficulty: maintain synchronization with current IP-ID
  - Incremented for every packet (regardless of arrival/loss)




# Continual Deny & Expose

- Idea: use zombie to `monitor' IP-ID progress
  - Send two sequences of spoofed fragments with consecutive IDs
  - Small `gap' of unsent IDs between them
  - Zombie makes a periodic request for data
  - Response arrives → ID within the gap
  - Send the next sequence
- Causing over 95% loss rate



# Continual Deny & Expose - Results

- Success depends on the number of forged fragment attacker can 'cache in'
  - Usually 64 or no limitation (except cache size, 6500+)





# TCP/IP Security

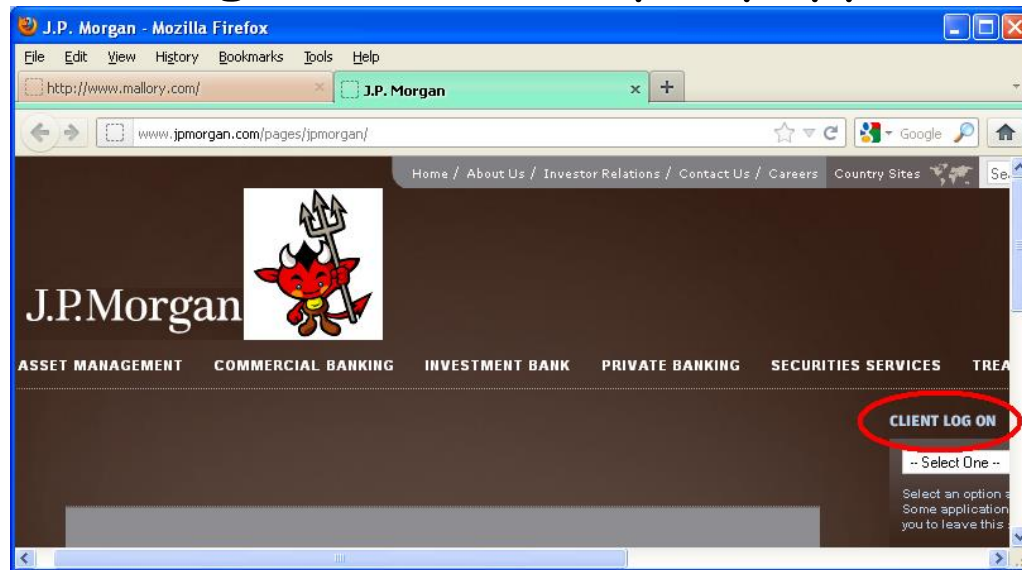
---

- Internet Protocol (IP) Security
  - Fragmentation attacks
  - IPsec
- **Secure Transport**
  - TCP injections and other attacks
  - Quic
- Conclusion

# Off-Path Attacks on TCP

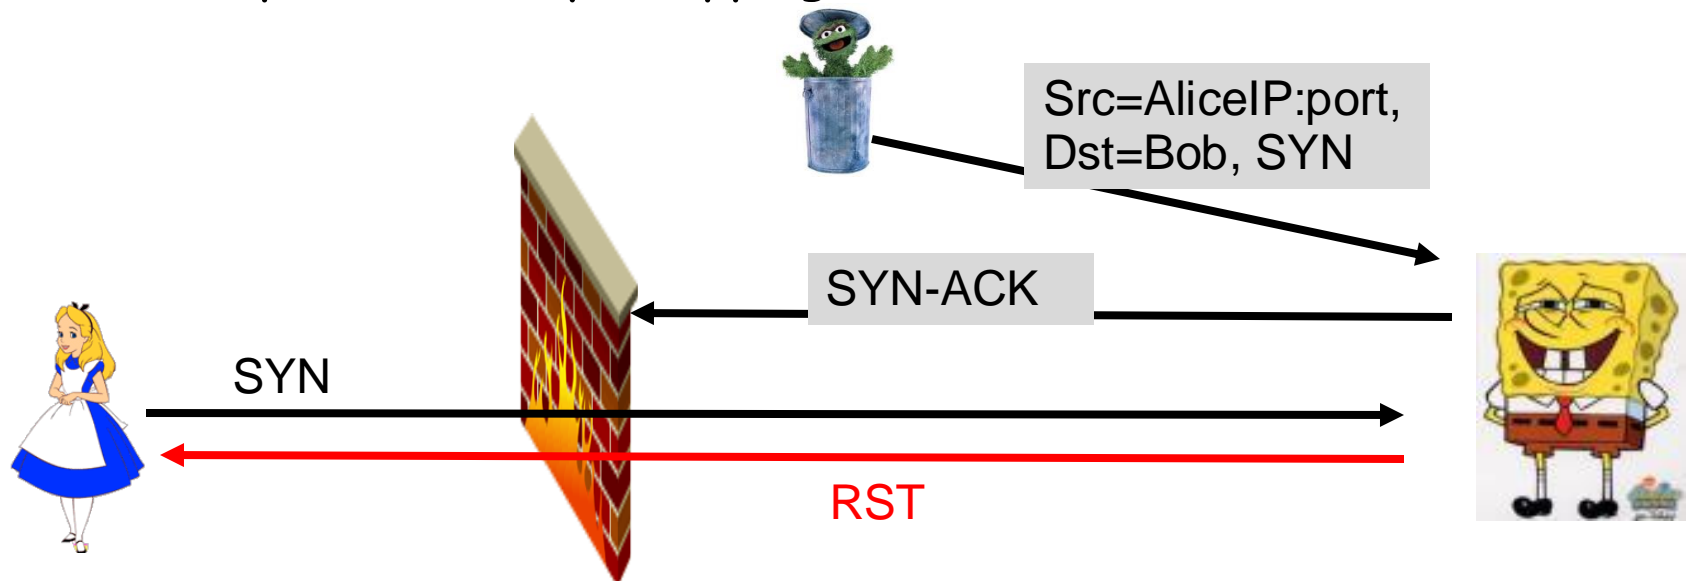
- Denial/degradation of Service (DoS)
  - ❑ SYN-flood: DoS server (crash / no new connections)
  - ❑ BW/performance degradation attacks
  - ❑ TCP Amplification attacks
  - ❑ Foil connection from client (by IP or also port)
  - ❑ Break ongoing TCP connection
- TCP Injection, e.g., of mal-script (puppet)

DoS  
lectures




# More Off-Path TCP DoS Attacks

- Foil new connection from given client
  - Method 1: send many SYNs → server blocks client IP
    - Due to (some) rate limiting defenses (against SYN-Flood)
  - Method 2: Allocate client's 4-tuple (aka *4-tuple blocking*)
    - 4-tuple: (clientIP:port, serverIP:port)
    - Send SYN using 4-tuple; server allocates socket, rejects legit-client connection using same 4-tuple
    - Exploits silently-dropping FW/NAT to avoid RST to Bob



# More Off-Path TCP DoS Attacks

- Foil new connection from given client
  - Method 1: send many SYNs → server blocks client IP
    - Due to (some) rate limiting defenses (against SYN-Flood)
  - Method 2: Allocate client's 4-tuple (aka *4-tuple blocking*)
    - 4-tuple: (clientIP:port, serverIP:port)
    - Send SYN using 4-tuple; server allocates socket, rejects legit-client connection using same 4-tuple
    - Exploits silently-dropping FW/NAT or **RST-dropping FW** to avoid RST to Bob

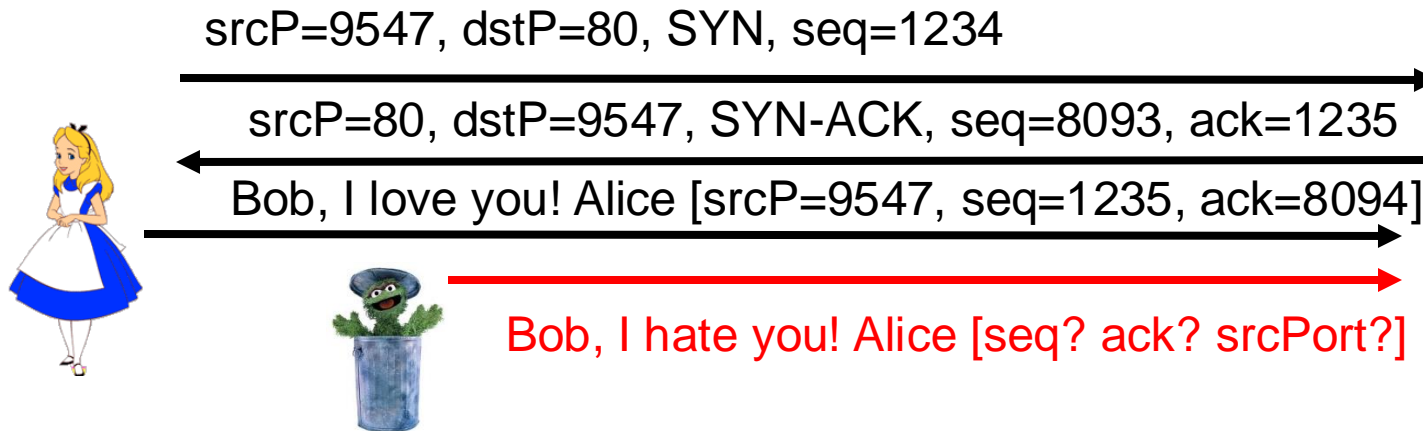


# More Off-Path TCP DoS Attacks

- Foil new connection from given client
  - Method 1: send many SYNs → server blocks client IP
    - Due to (some) rate limiting defenses (against SYN-Flood)
  - Method 2: Allocate client's 4-tuple (aka *4-tuple blocking*)
    - 4-tuple: (clientIP:port, serverIP:port)
    - Send SYN using 4-tuple; server allocates socket, rejects legit-client connection using same 4-tuple
    - IPs and server port are often known; Client port = ??
    - Note: may use 4-tuple-blocking also to find client's next port
      - For incrementing-source-port clients (most OSs)
- Break ongoing TCP connection
  - Method 1: spoofed RST
    - Must have correct IPs+ports and seq-num 'in window'
  - Method 2: Spoofed ICMP error message (same 4-tuple)

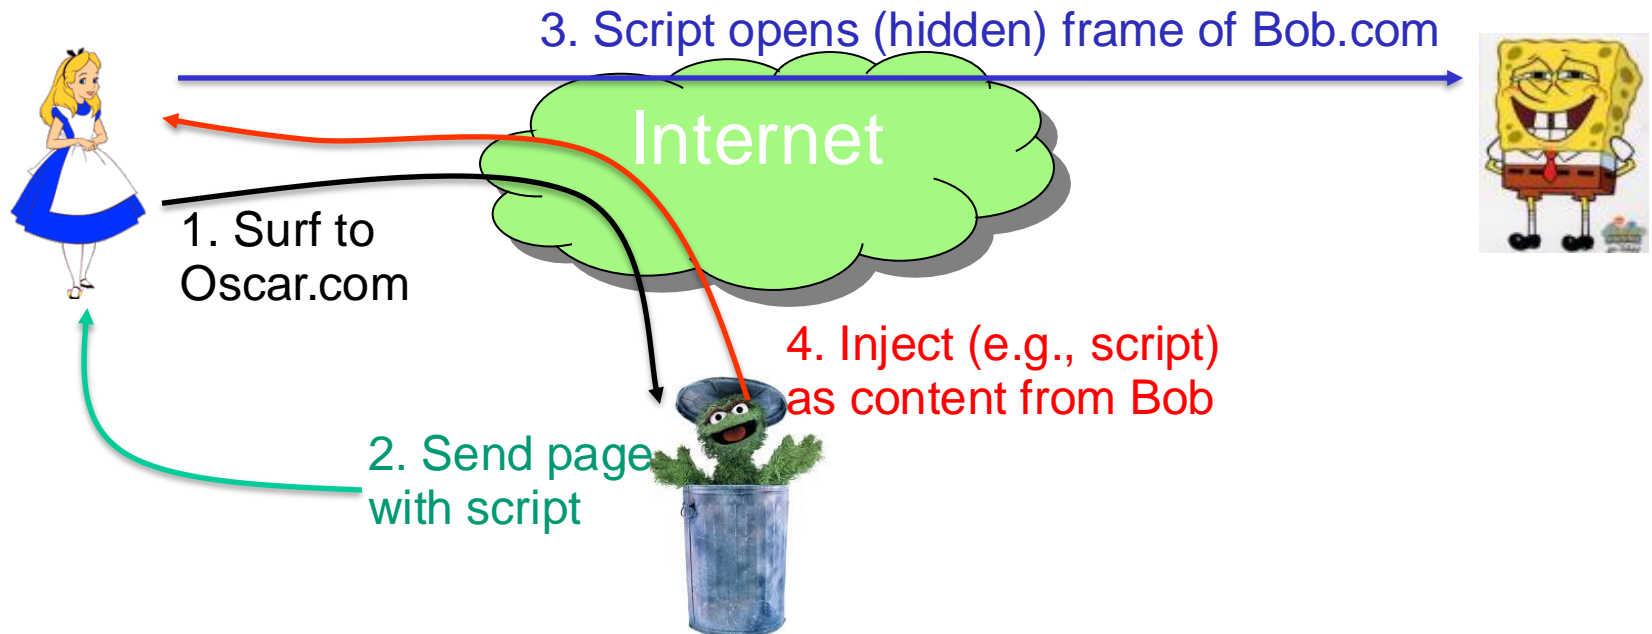
# Off-path TCP inject challenges

- No explicit off-path defenses in TCP
- But... injection requires:
  - 4-tuple: (clientIP:port, serverIP:port)
    - IPs and server port are often known
  - And sequence/ack numbers
  - Initialized randomly (since the 1990s)



# Off-Path TCP Injection: Scenario

1. Alice surfs to Oscar's site
2. Alice's browser runs Oscar's script (puppet)
3. Puppet sends HTTP requests to Bob
4. Oscar injects response (e.g. mal-script) into connection
  - ❑ Alice's browser assigns mal-script with origin of `Bob`
  - ❑ Cached → long term attack
  - ❑ Use for phishing, credentials/info-theft, malware



# Off-path TCP Injection: Challenges


---

- First: identify `victim-connection`
  - I.e., find 4-tuple: srcIP+port, dstIP+port
  - Ongoing or (easier) initiated by attacker (puppet)
  - Only part we show in this presentation
- Then: learn sequence numbers (seq+ack)
  - TCP discards packets with invalid seq #
  - We will not show in this presentation
- Finally: inject and exploit
  - Send (spoofed) data in correct (HTTP) context
  - Carefully manage TCP counters to avoid Ack storm from packets of legit destination (victim)
  - Browser assigns data the credentials of server (Bob)
    - Defeating `Same Origin Policy`
  - Typically: send `cross site script` (XSS)



# Finding 4-tuple: <ServerIP:port, ClientIP:port>

- Trivial - for MitM or eavesdropper
- Easy for globally-seq client ports:
  - Puppet (script in browser) opens connection to Bob (server)
  - ServerIP:port known - or selected by puppet (attacker)
  - Client IP: known from client connection to Oscar



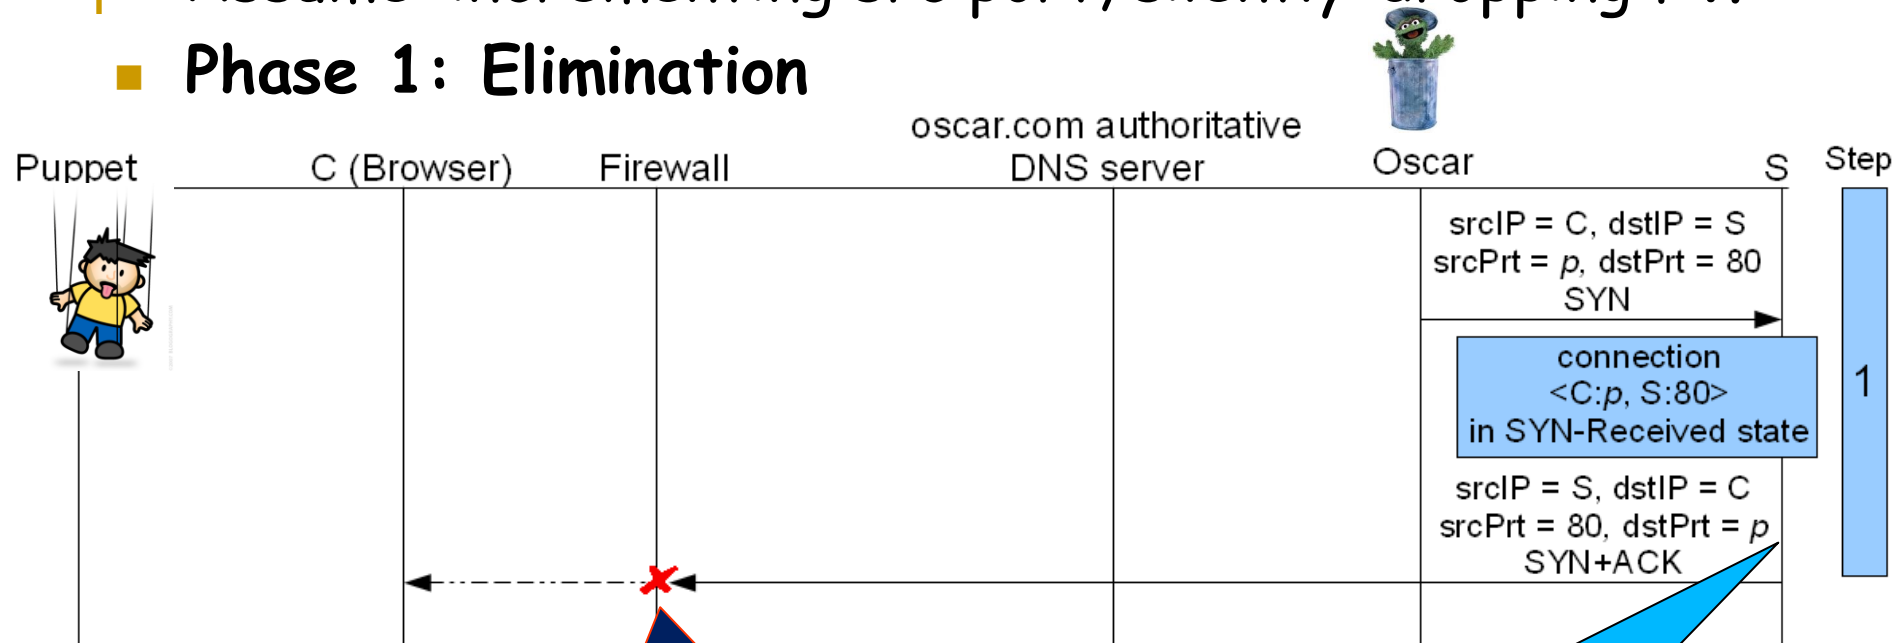
- Challenge: most clients use per destination incrementing (hash-based) ports

# Per-dest incrementing (Hash-Based) Ports

- Algorithm:
  - On first connection to S:  $\text{port} = \text{hash}(S, \text{key})$
  - Increment port for new connections with S
- A per-destination 'randomized counter'
- Common
  - Linux (since 2006), Android, iOS
  - DNS resolvers (since 2008)
  - ...
- Find next port by *Eliminate-then-Test* attack:
  - Find next port to be used by per-destination incrementing client
  - By using *4-tuple blocking* to block specific ports, then detect if current port is blocked
  - Assume: silently-dropping of unsolicited SYN-ACK packets

# Eliminate then Test Attack

- Goal: predict next  $C \rightarrow S$  src port
- Assume: incrementing src port, silently-dropping FW
- Phase 1: Elimination

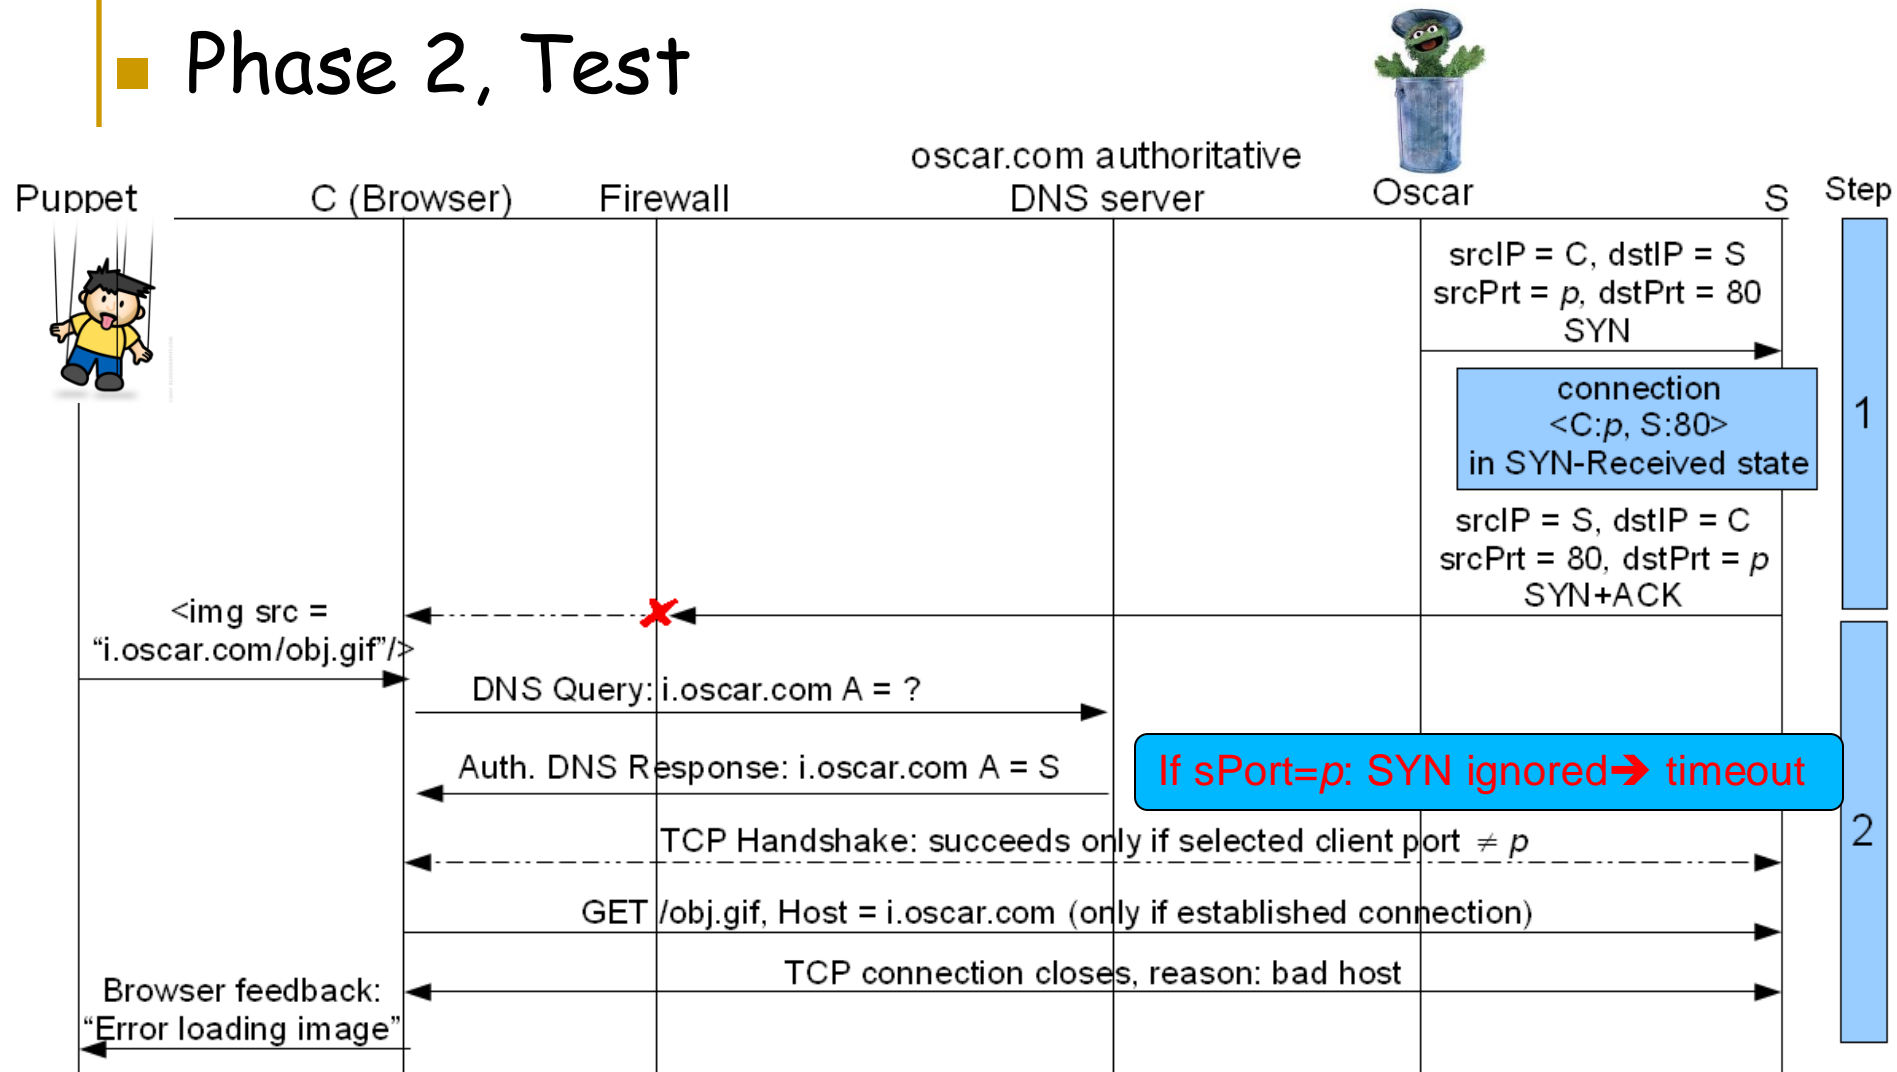


silently-dropping of unsolicited SYN-ACK packets

Move to SYN-RCVD state; ignore any pkt with incorrect seq# (incl. SYN)...

# Eliminate then Test

## ■ Phase 2, Test

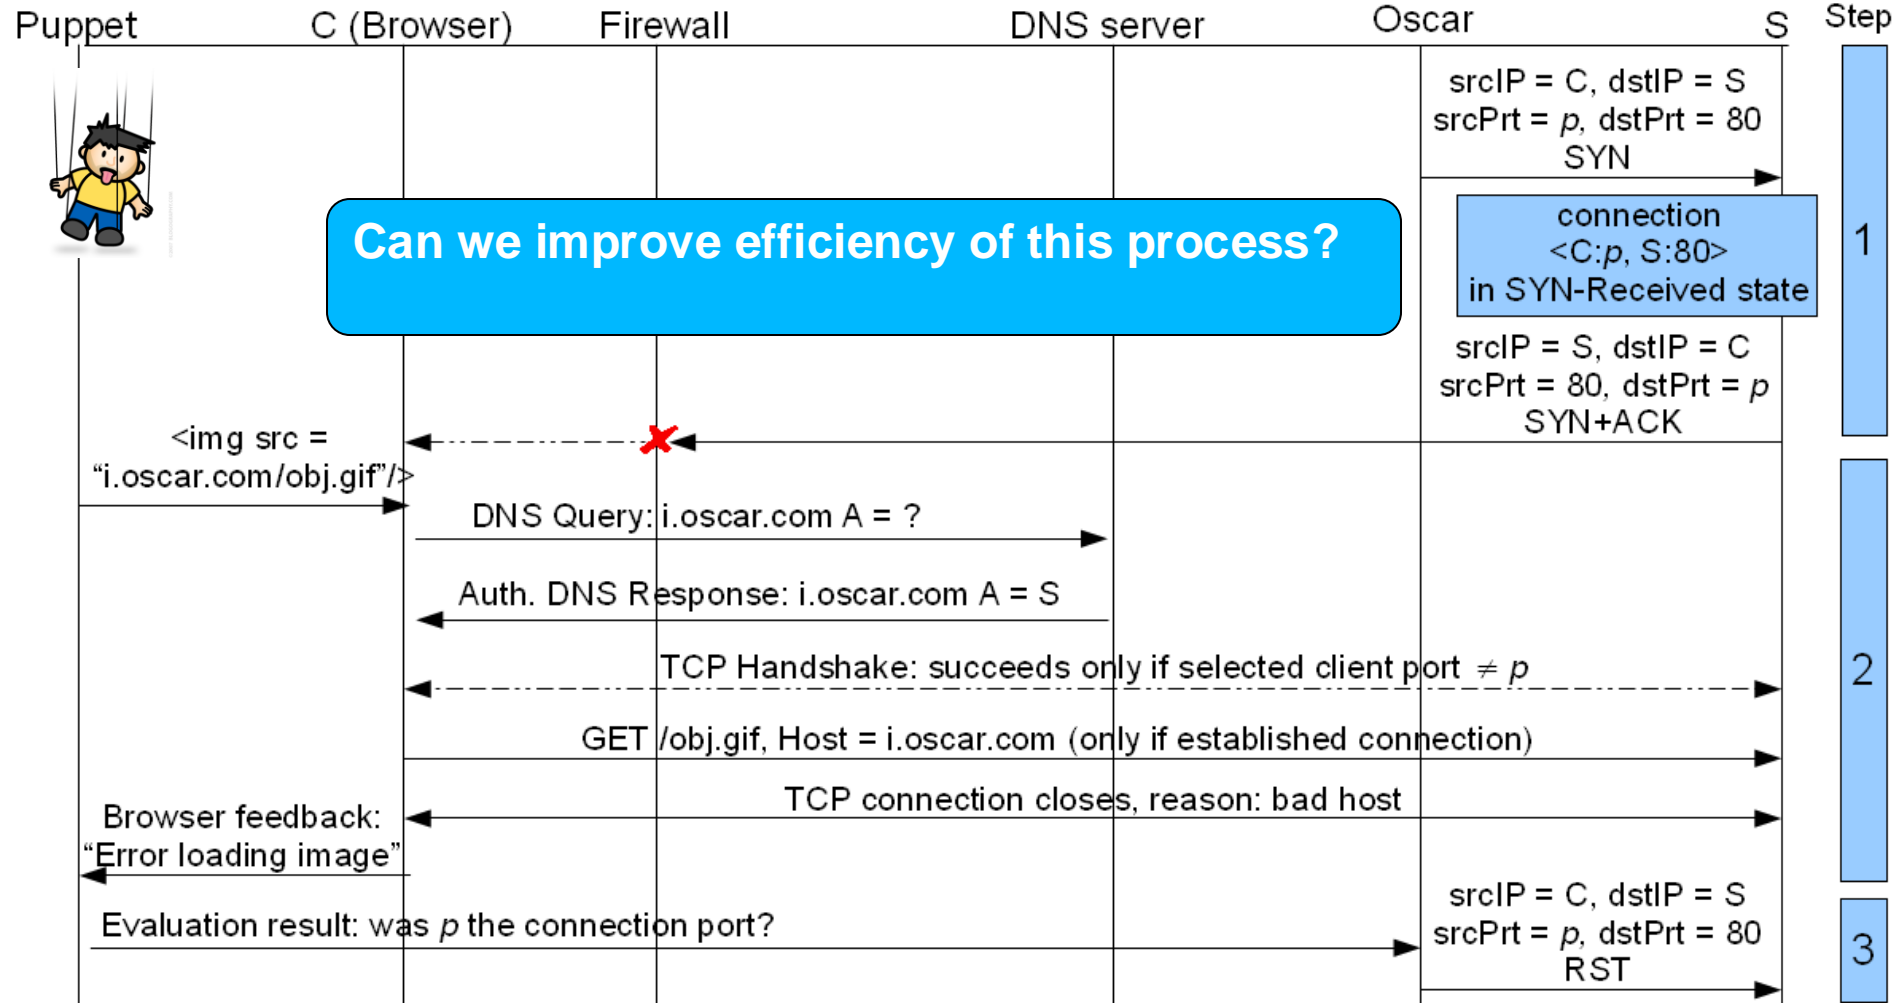


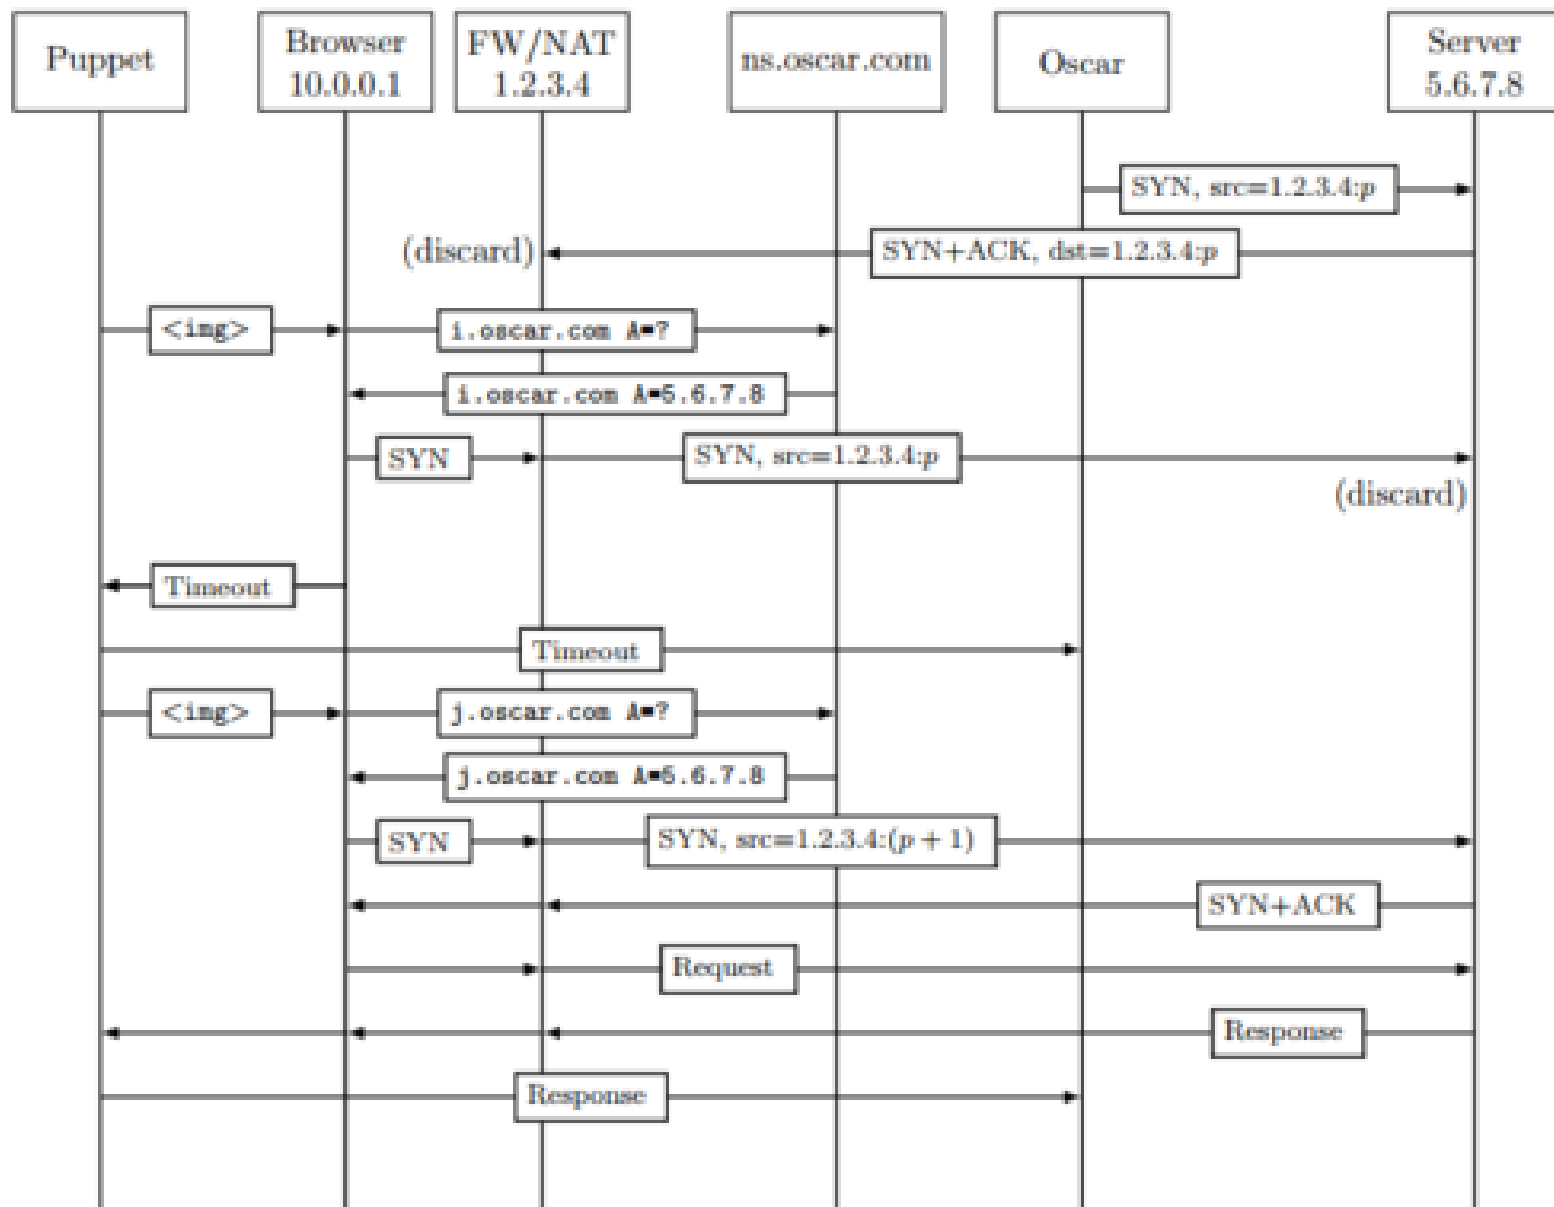
# Eliminate then Test

## ■ Also Phase 3, Cleanup...



oscar.com authoritative  
DNS server





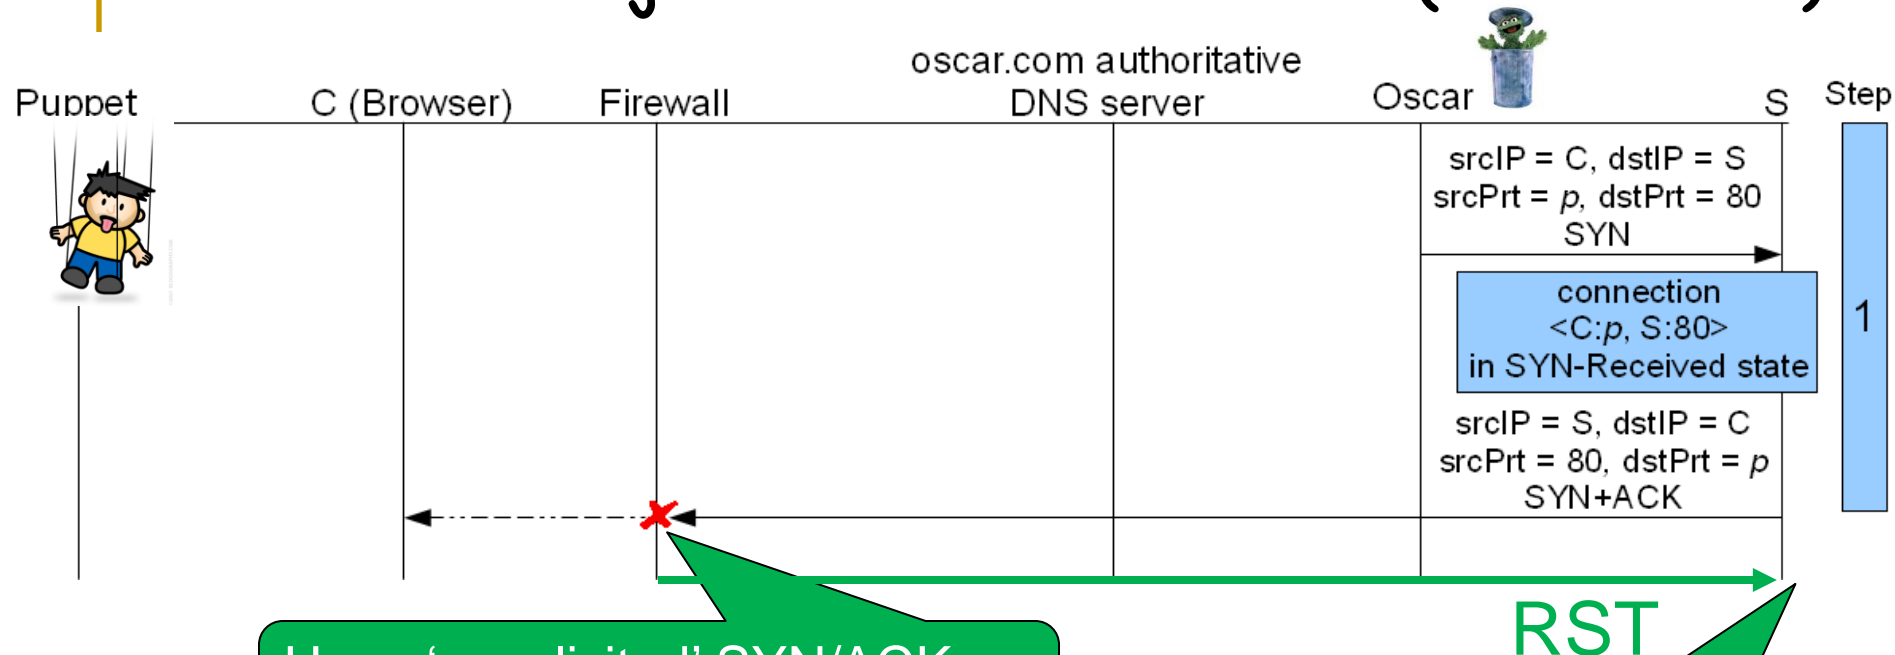
# Finding client port: battleship optimization

---

- `Eliminate and test' can test all ports
  - Works, but slow: 65536 ports!
- Optimize with battleship optimization:
  - Eliminate ports  $256*i$ , for  $i=1\dots256$
  - Open 256 connections to server (as before)
  - ... complete details - exercise!
- Over 90% success rate!
- Question: how could FW foil the attack?

# Firewall Could Foils Elimination Phase

- Firewall 'rejects' bad SYN/ACK (send RST)



Upon 'unsolicited' SYN/ACK:  
drop, but also send back RST !

Note: sending RST also foils SYN-Flood and  
SYN-ACK amplification  
Alternative attack exploits IP-ID; see [1]

Connection aborted,  
back to 'Listen' state...  
Will not ignore SYN !



# Finding Sequence Numbers

---

- Injection requires server, client seq#
- To inject data to client:
  - Server seq# should be 'in window'
    - If not exact: cached
  - Client seq#: depends on OS
    - Exact, in window, or arbitrary
- Different methods work (or fail) for different systems...
  - Cat and mouse game
  - See papers... not in course

# TCP/IP Security

---

- Internet Protocol (IP) Security
  - Fragmentation attacks
  - IPsec
- **Secure Transport**
  - TCP injections and other attacks
  - Quick introduction to Quic and its security
- Conclusion

# Quick UDP Internet Connections

---

- QUIC: (Google's) new transport protocol
  - TCP-features: connections, reliability, controls...
  - Over UDP: clean-slate, user-space
    - Some UDP concerns (fragmentation, identify client)
  - Significantly (QUICkly) deployed: esp. Google ☺
  - Improve latency, perceived performance:
    - Multiplexed requests
    - Prioritized requests
    - Compression
  - Many of the functions/mechanisms of HTTP2
- Security?

# QUIC Security

---

- Built-in secure connections
  - Combines TCP and TLS handshake and features
  - Always encrypted, authenticated
- Reuse state across connections:
  - Source-address token: detect spoofing
  - Cookie: server's public DH values
    - Allows stateless server to resume session
- Main goal: 0-RTT (reduce latency) exchange
  - Typical case: key already established, has cookie
    - Notice: 0-RTT request may be replayed
  - Related to 0-RTT exchange of TLS 1.3

# Quic vs. Firewalls & other tools

---

- Easy to allow Quic to flow (open UDP:443)
- NAT: how to identify end of Quic session?
  - General UDP problem
  - May become harder as Quic gets popular
- Hard to validate, intercept Quic traffic
  - Adding the 'fake CA' to the OS certificate store is not sufficient: which fake cert to send??!
    - Good or bad? Depends whom you ask
    - E.g., prevents inspection by firewall, IDS/IPS,...
    - Possible response: block Quic !

# Conclusions: TCP/IP security

---

- TCP/IP stack is insecure against MitM
  - Unless using crypto: IPsec, TLS, DNSSEC
- But also: many off-path attacks
  - Some we learned, many we didn't
  - Some trivial, many not so trivial
  - Most work only in some scenarios
    - Typically: with insecure versions, configurations
- Few of the many topics not covered:
  - IPv6 Security
  - Failure localization and accountability